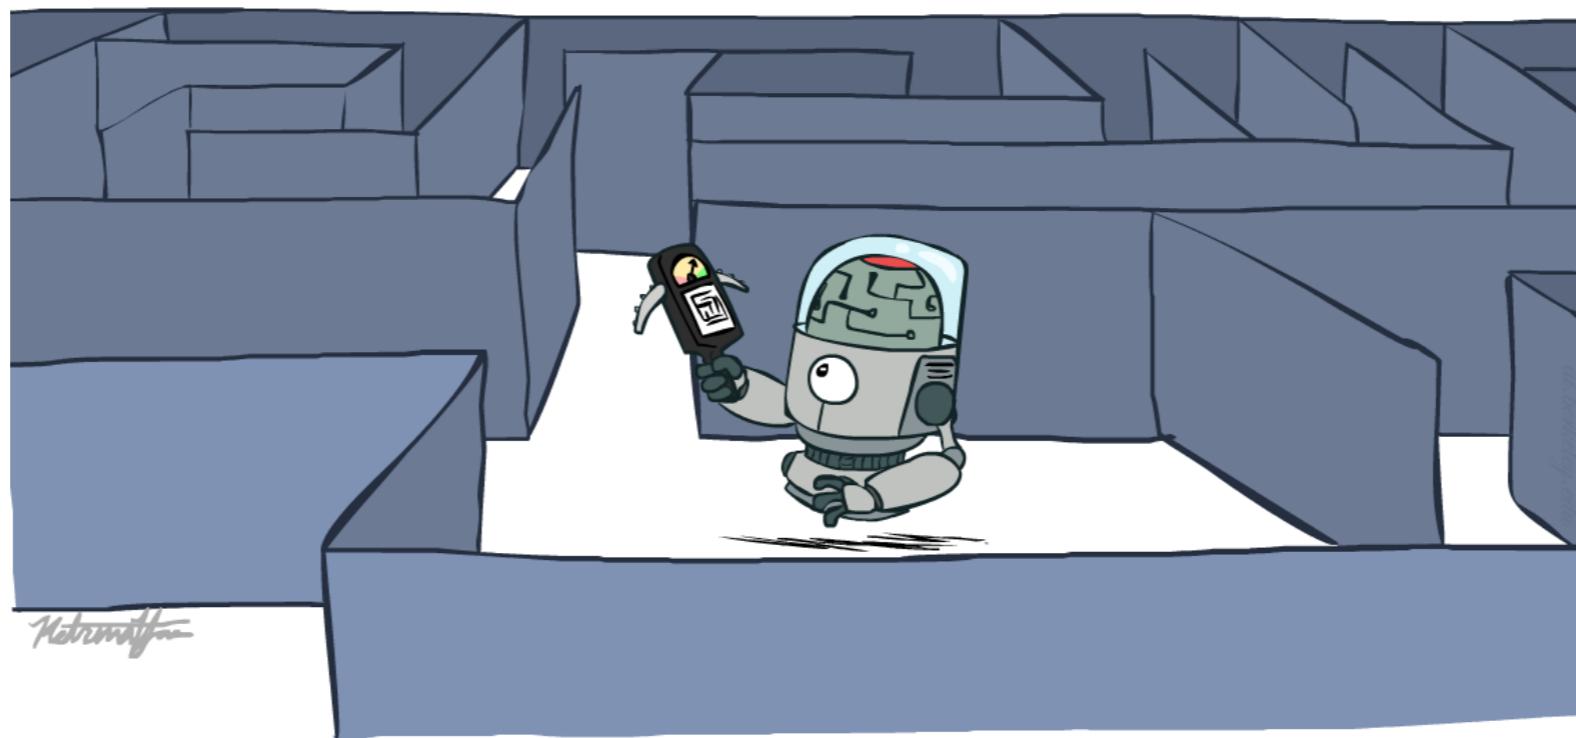

Announcements

- ❖ Homework 1: Search
 - ❖ Part I AND Part II due Wed. May 27 at 11:59pm.
 - ❖ Part I through Online Judge.
 - ❖ Part II through Canvas -- submit pdf
- ❖ Project 1: Search
 - ❖ Released today. Due Wed June 3 at 11:59pm
 - ❖ Start early and ask questions. It's longer than most!
- ❖ Grading policy for HW and Project
 - ❖ 20% deduction per day of late submission
 - ❖ Drop 2 lowest grades for electronic HW part and 2 lowest grades for written HW part
 - ❖ 5 slip days for projects; maximum 2 slip days for a given project

Ve492: Introduction to Artificial Intelligence

Informed Search and Graph Search



Paul Weng

UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

Outline

- ❖ Informed Search
 - ❖ Heuristics
 - ❖ Greedy Search
 - ❖ A* Search
- ❖ Graph Search



Recap: Search

❖ Search problem:

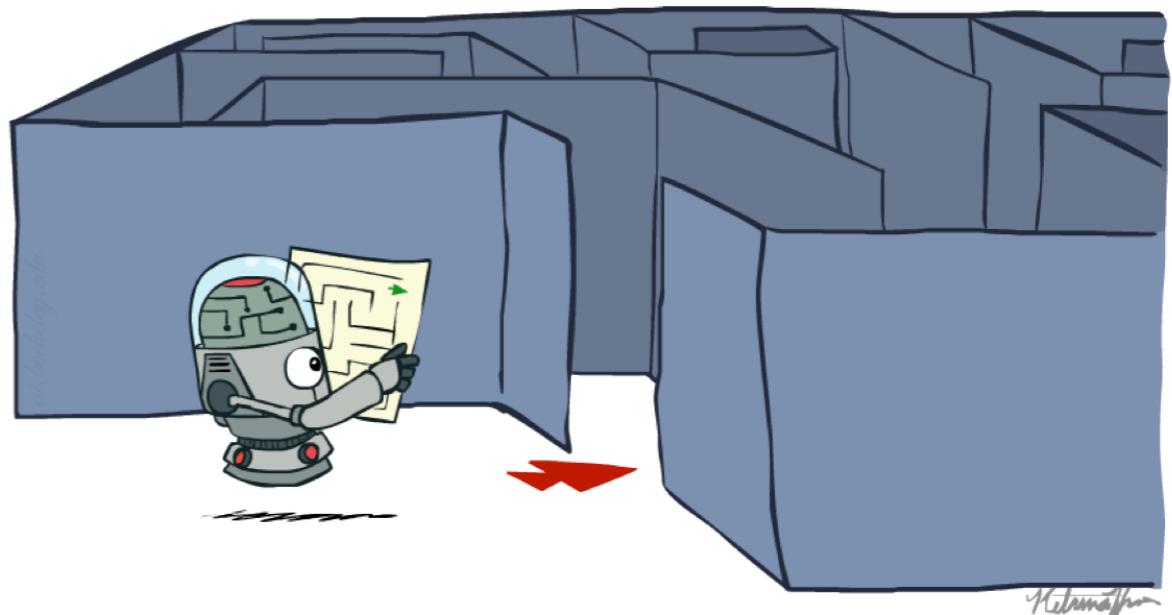
- ❖ States (configurations of the world)
- ❖ Actions and costs
- ❖ Successor function (world dynamics)
- ❖ Start state and goal test

❖ Search tree:

- ❖ Nodes: represent plans for reaching states
- ❖ Plans have costs (sum of action costs)

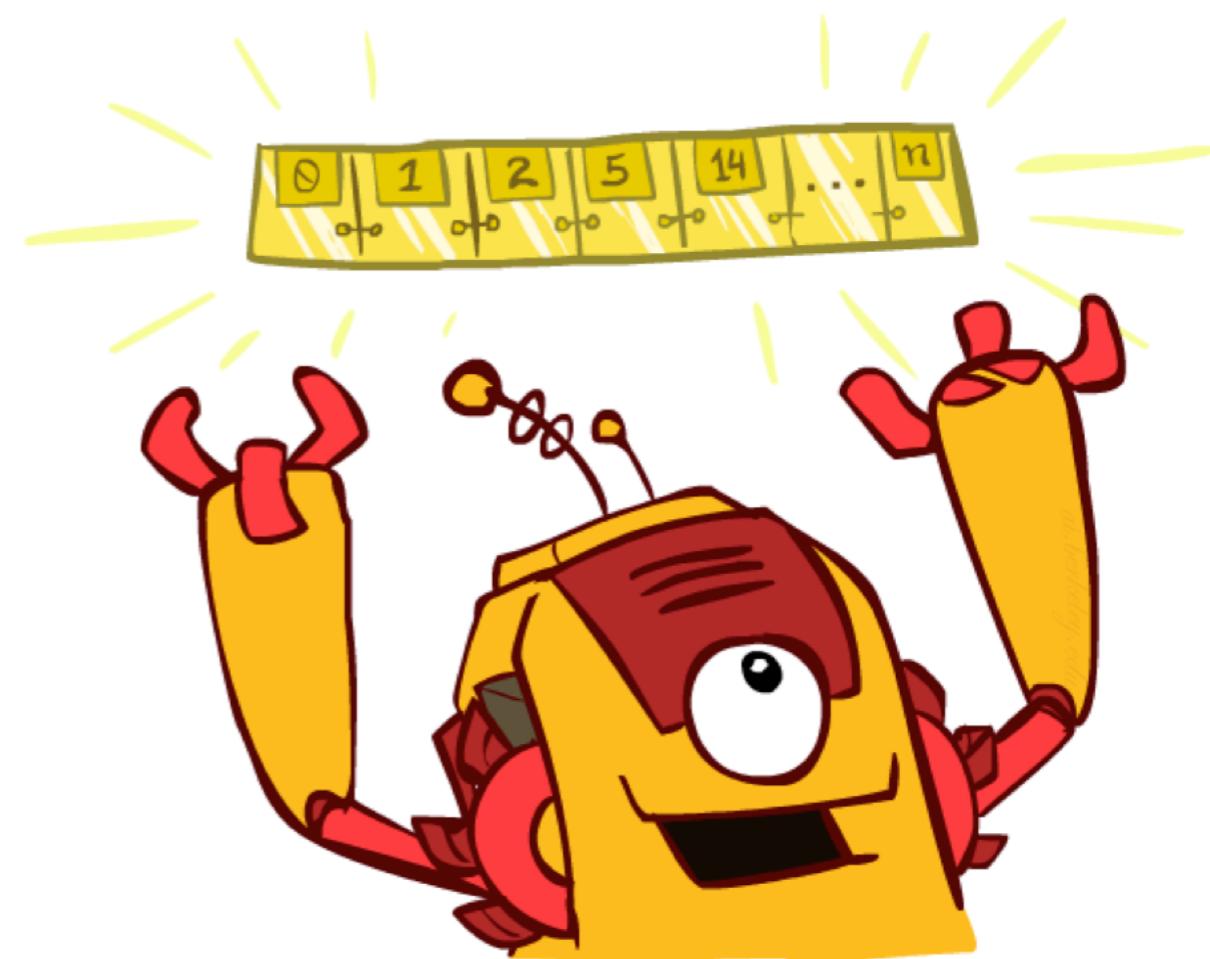
❖ Search algorithm:

- ❖ Systematically builds a search tree
- ❖ Chooses an ordering of the fringe (unexplored nodes)
- ❖ Optimal: finds least-cost plans



The One Queue

- ❖ All these search algorithms are the same except for fringe strategies
 - ❖ Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
 - ❖ Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stacks and queues
 - ❖ Can even code one implementation that takes a variable queuing object



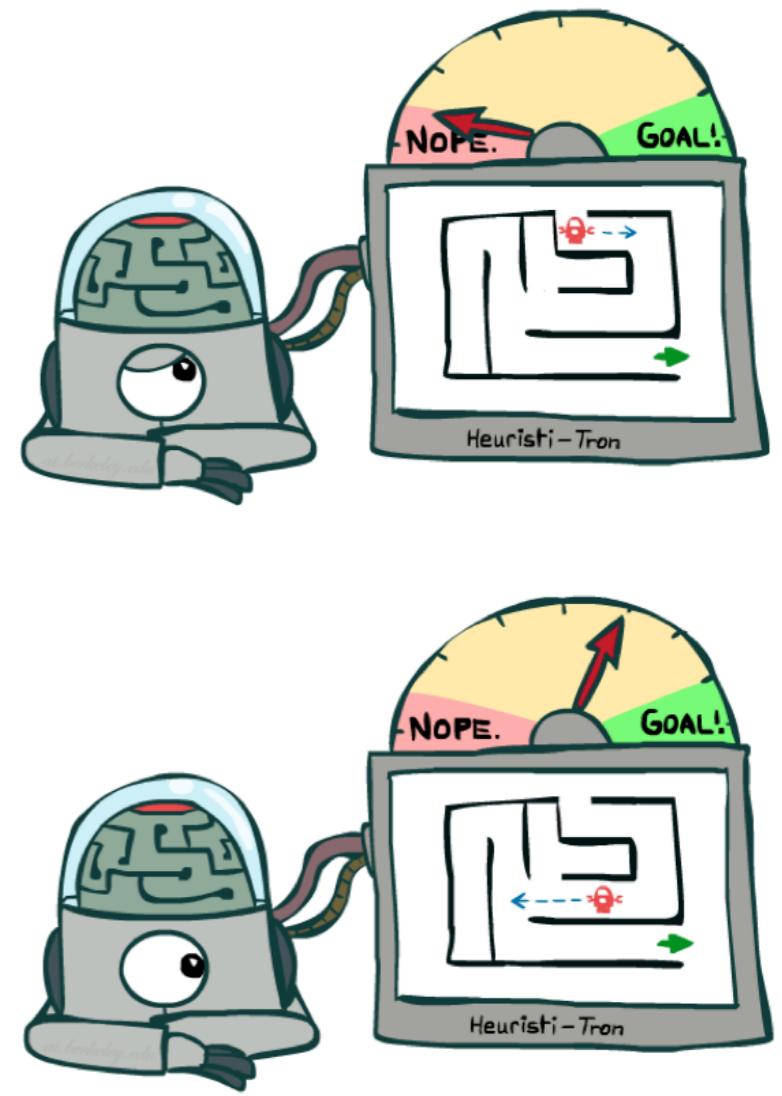
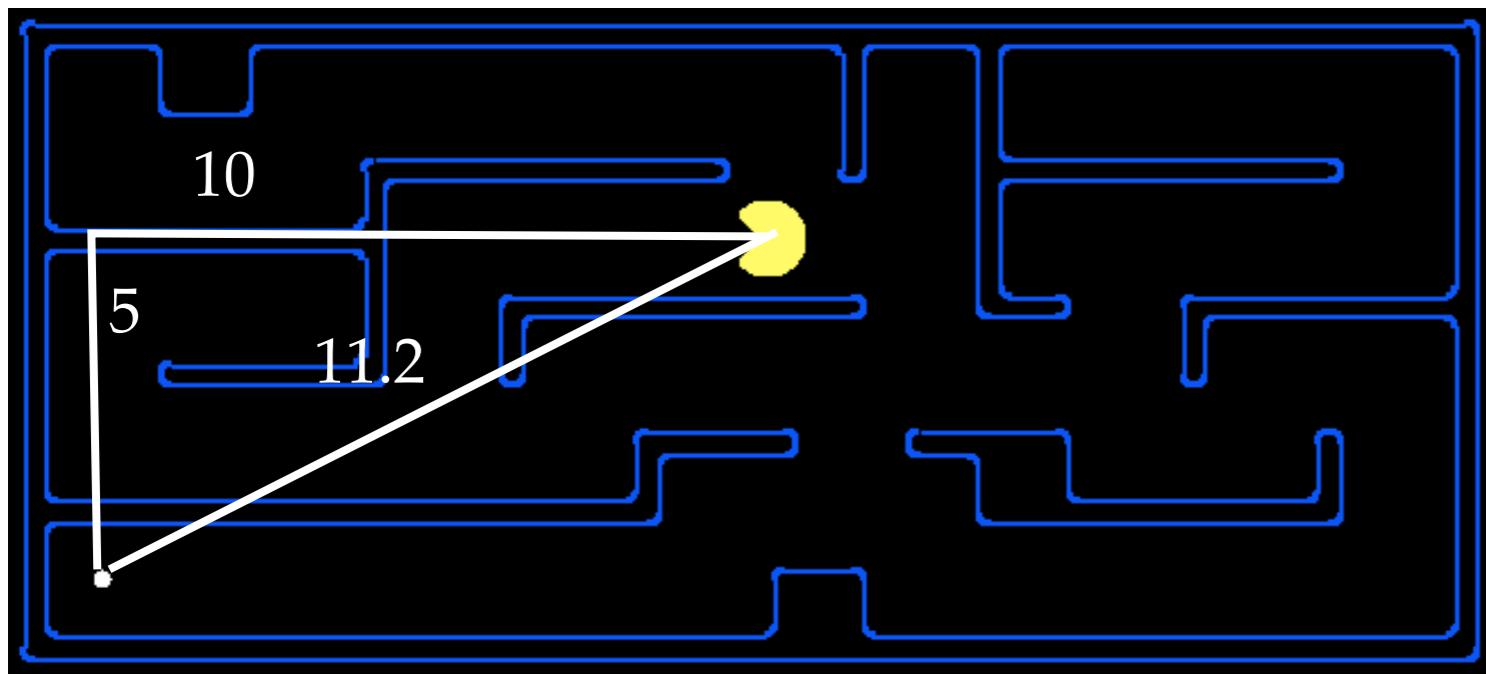
Informed Search



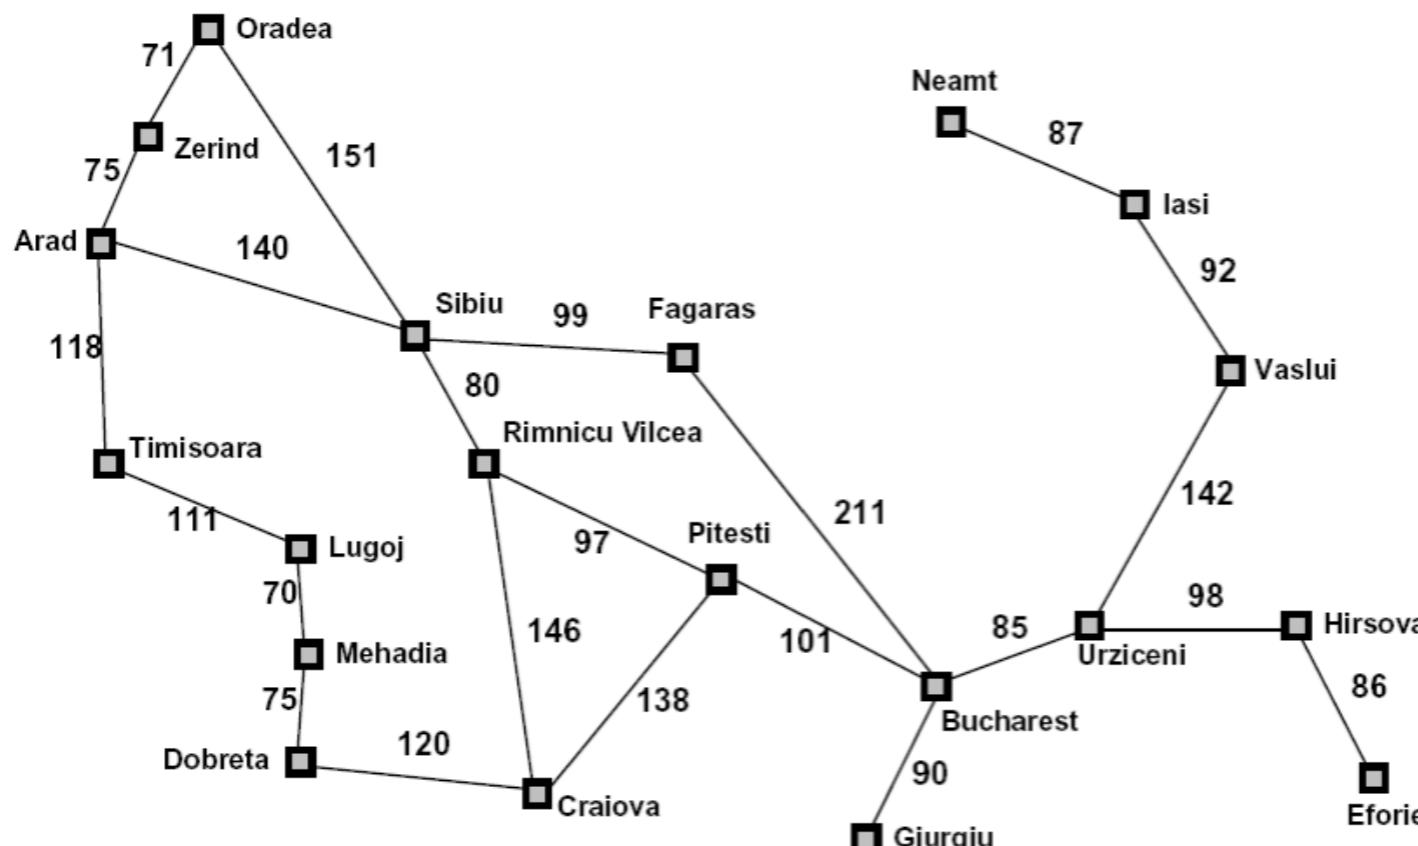
Search Heuristics

- ❖ A heuristic (function) is:

- ❖ A function that estimates how close a state is to a goal
- ❖ Designed for a particular search problem
- ❖ Examples: Manhattan distance, Euclidean distance for navigation



Example: Heuristic Function



| City | Straight-line distance to Bucharest |
|----------------|-------------------------------------|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 178 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |

$h(x)$

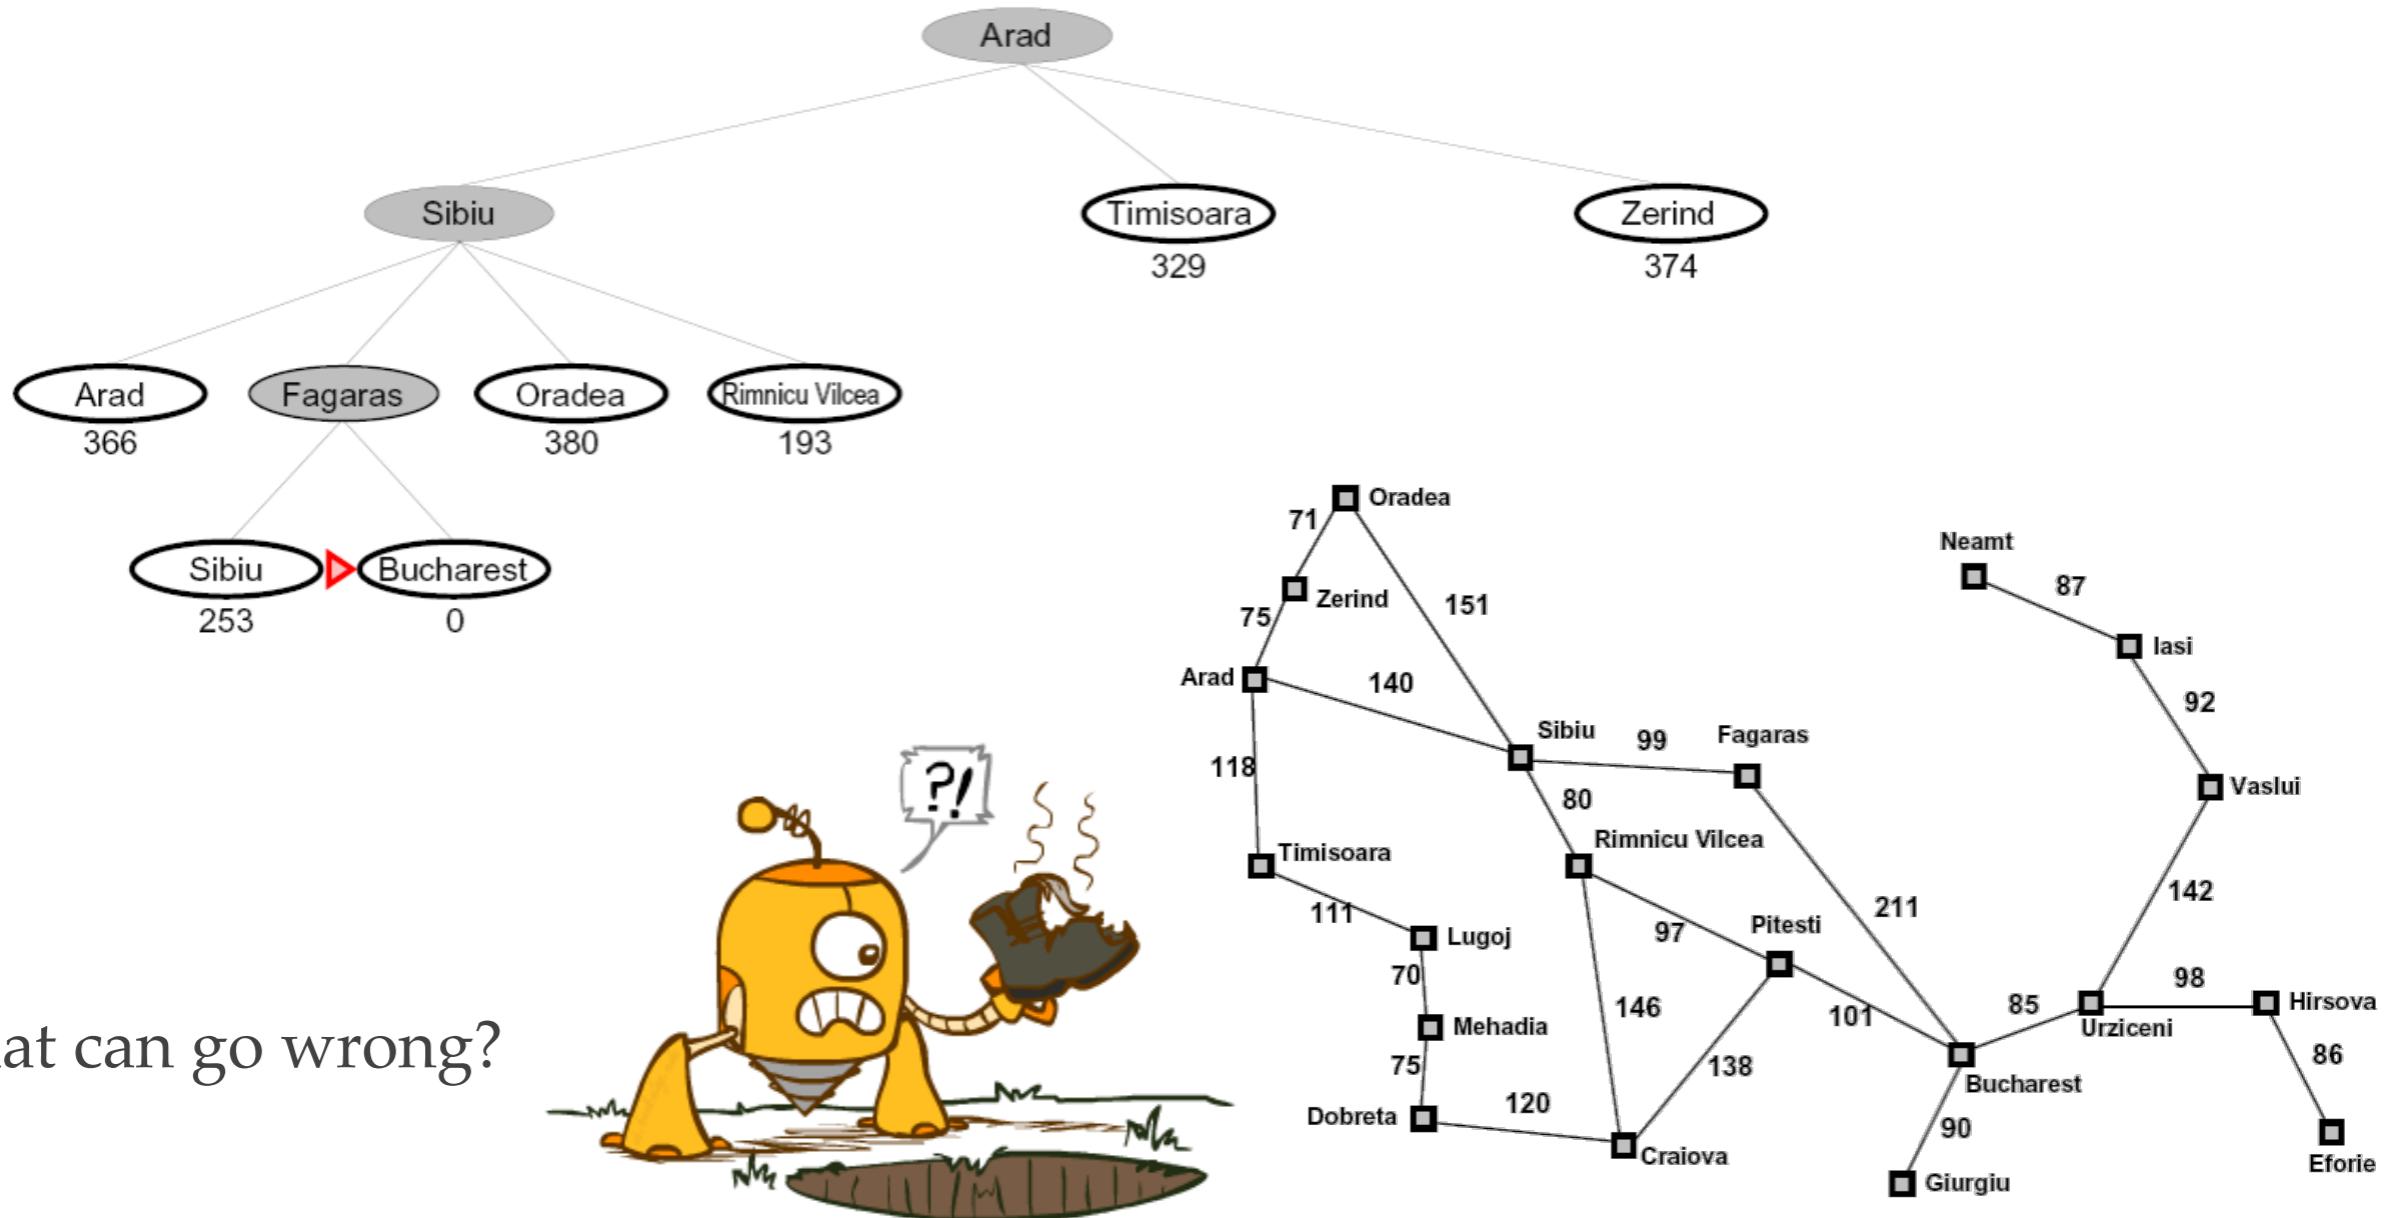
Informed Search

Greedy Search



Greedy Search

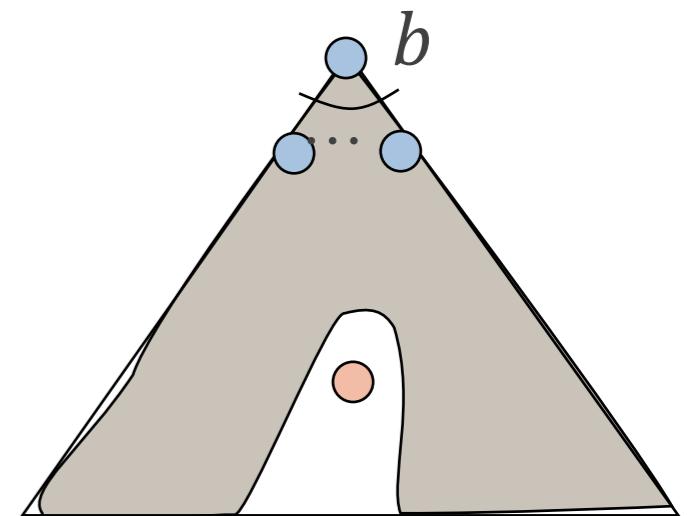
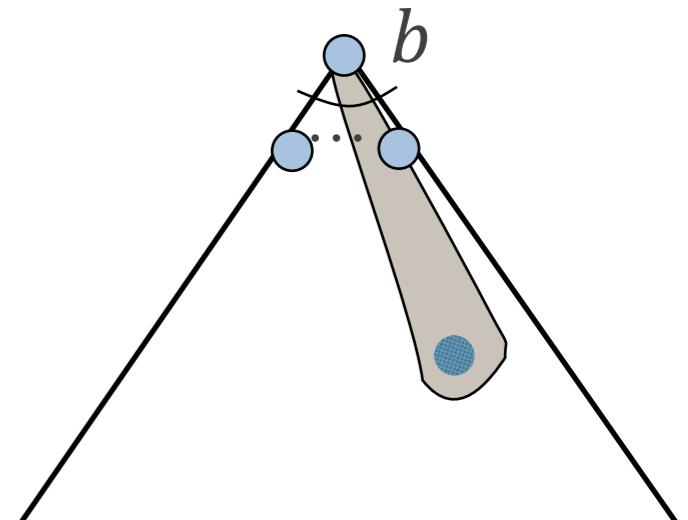
- ❖ Expand the node that seems closest...



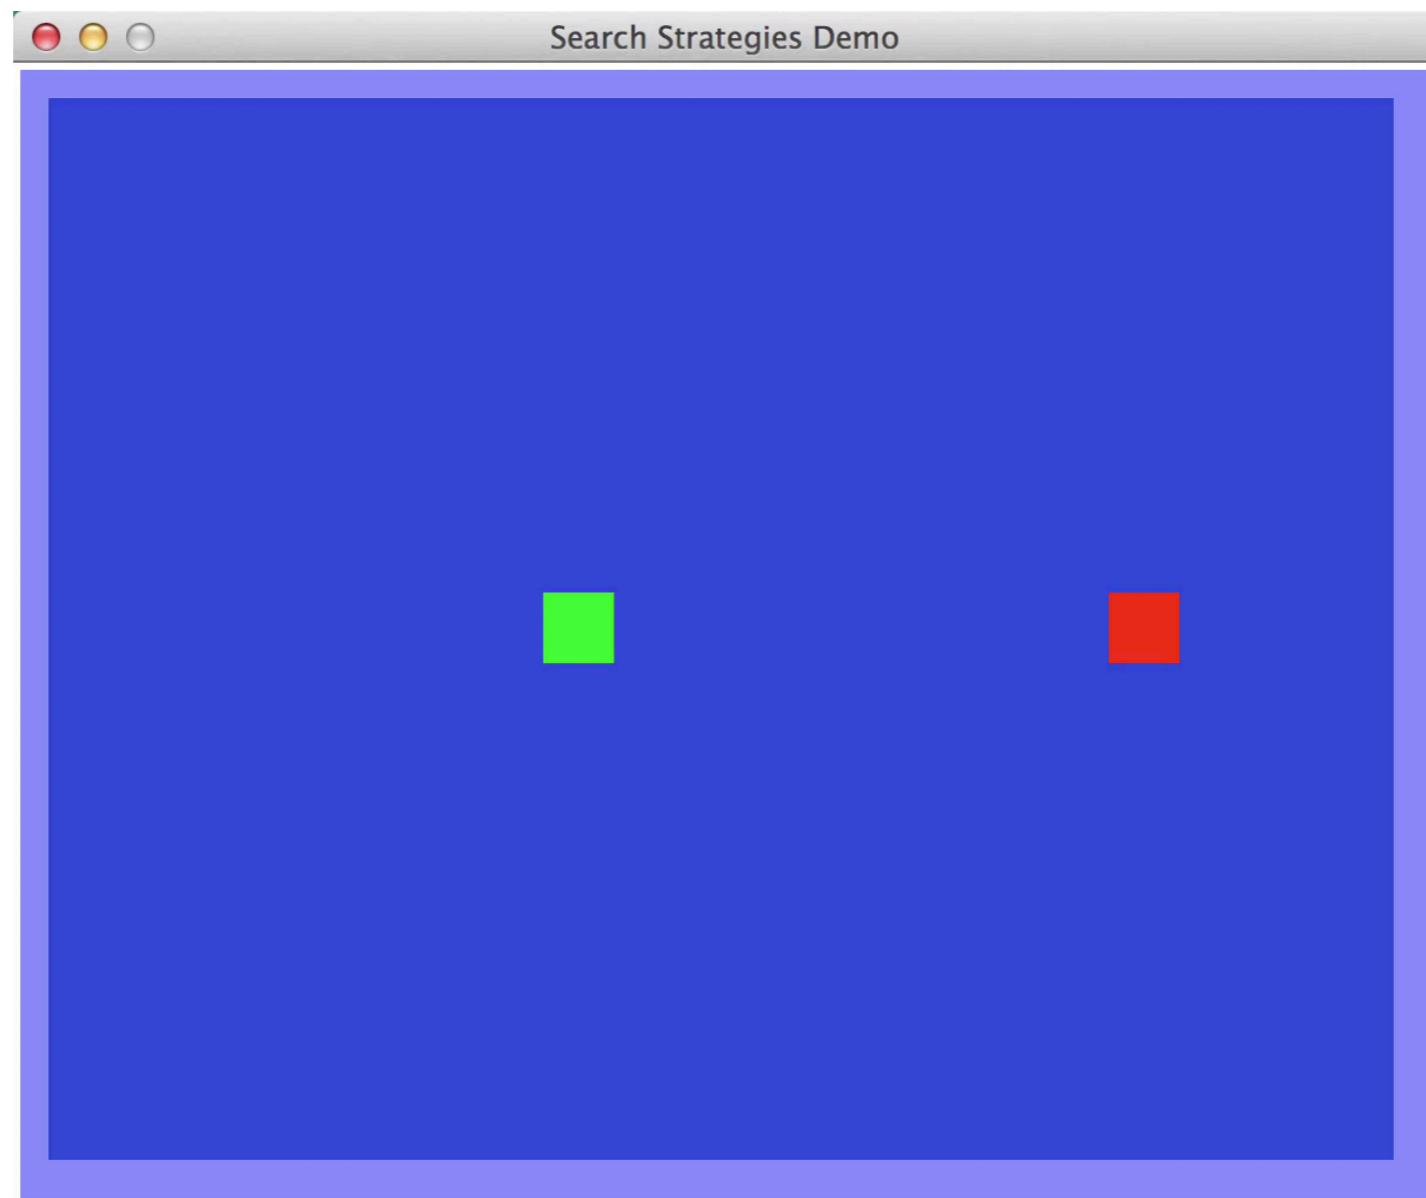
- ❖ What can go wrong?

Greedy Search

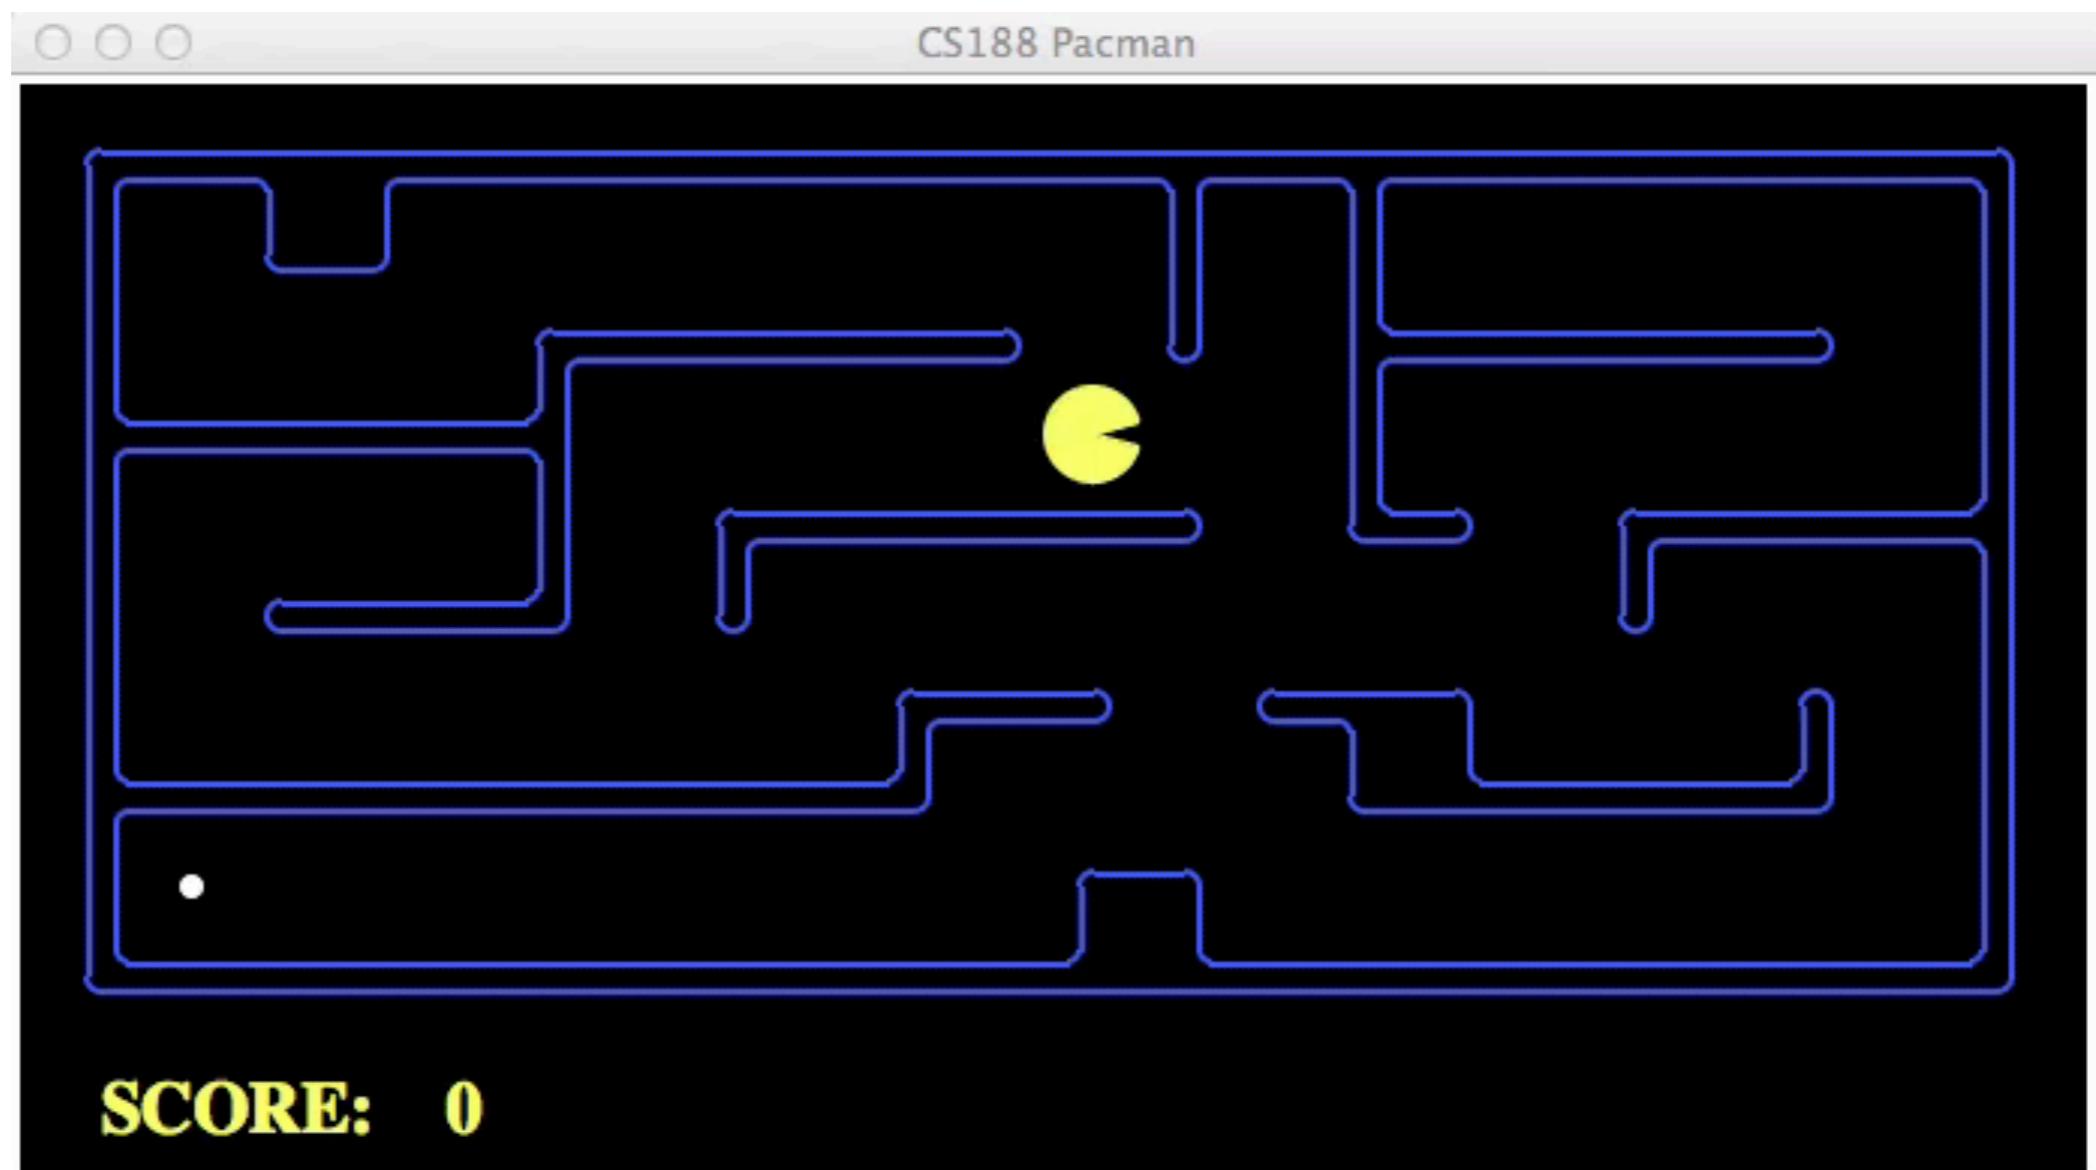
- ❖ **Strategy:** expand a node that you think is closest to a goal state
 - ❖ Heuristic: estimate of distance to nearest goal for each state
- ❖ **A common case:**
 - ❖ Best-first takes you straight to the (wrong) goal
- ❖ **Worst-case:** like a badly-guided DFS



Video of Demo Contours Greedy (Empty)



Video of Demo Contours Greedy (Pacman Small Maze)



Informed Search

A* Search



A* Search



UCS



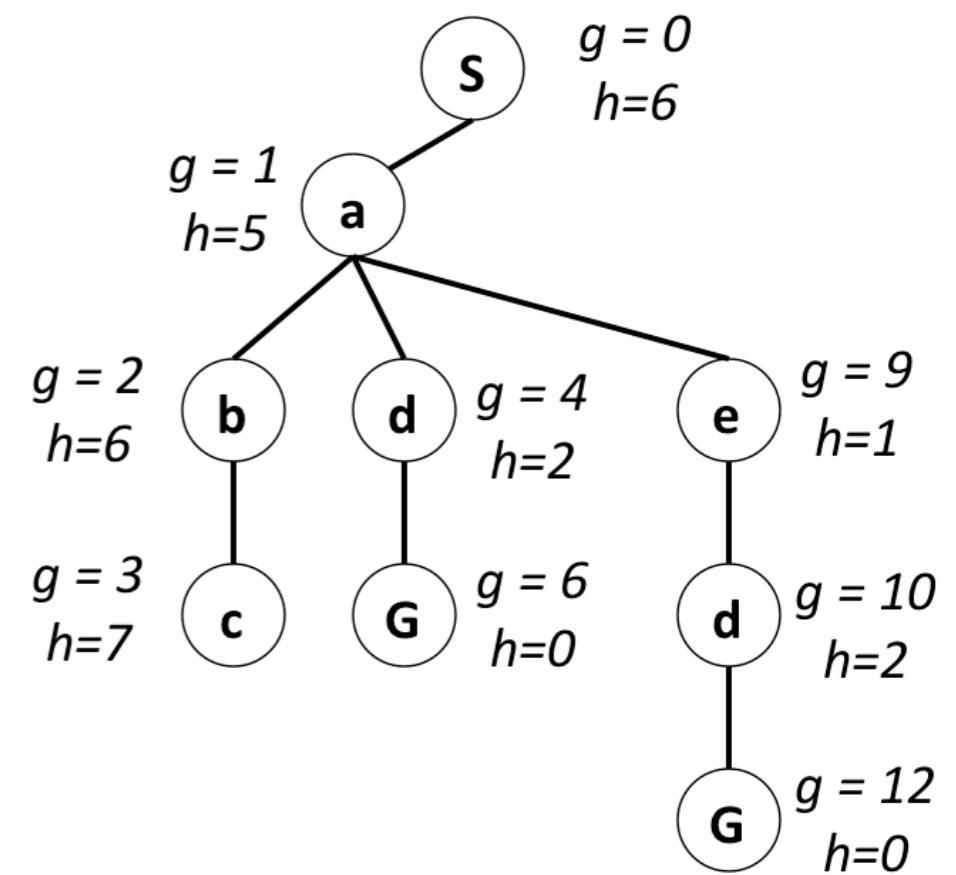
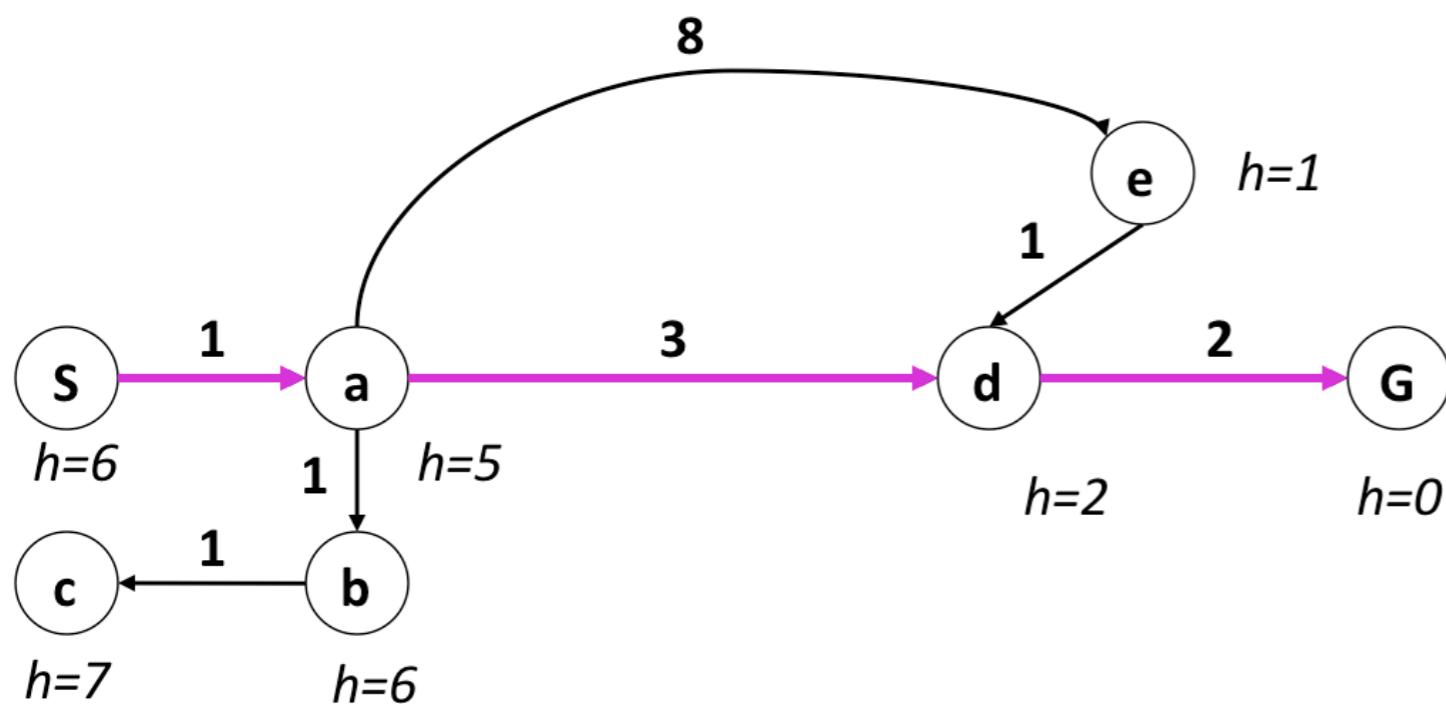
Greedy



A*

Combining UCS and Greedy

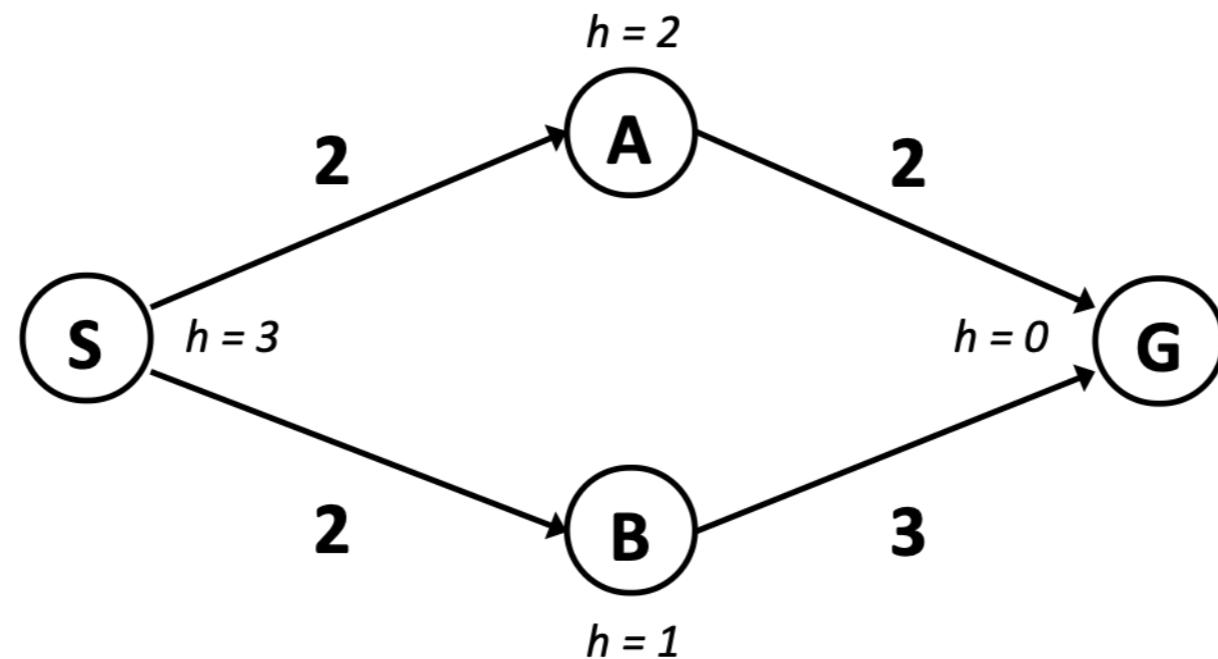
- ❖ Uniform-cost orders by path cost, or backward cost $g(n)$
- ❖ Greedy orders by goal proximity, or forward cost $h(n)$
- ❖ A* Search orders by the sum: $f(n) = g(n) + h(n)$



Example: Teg Grenager

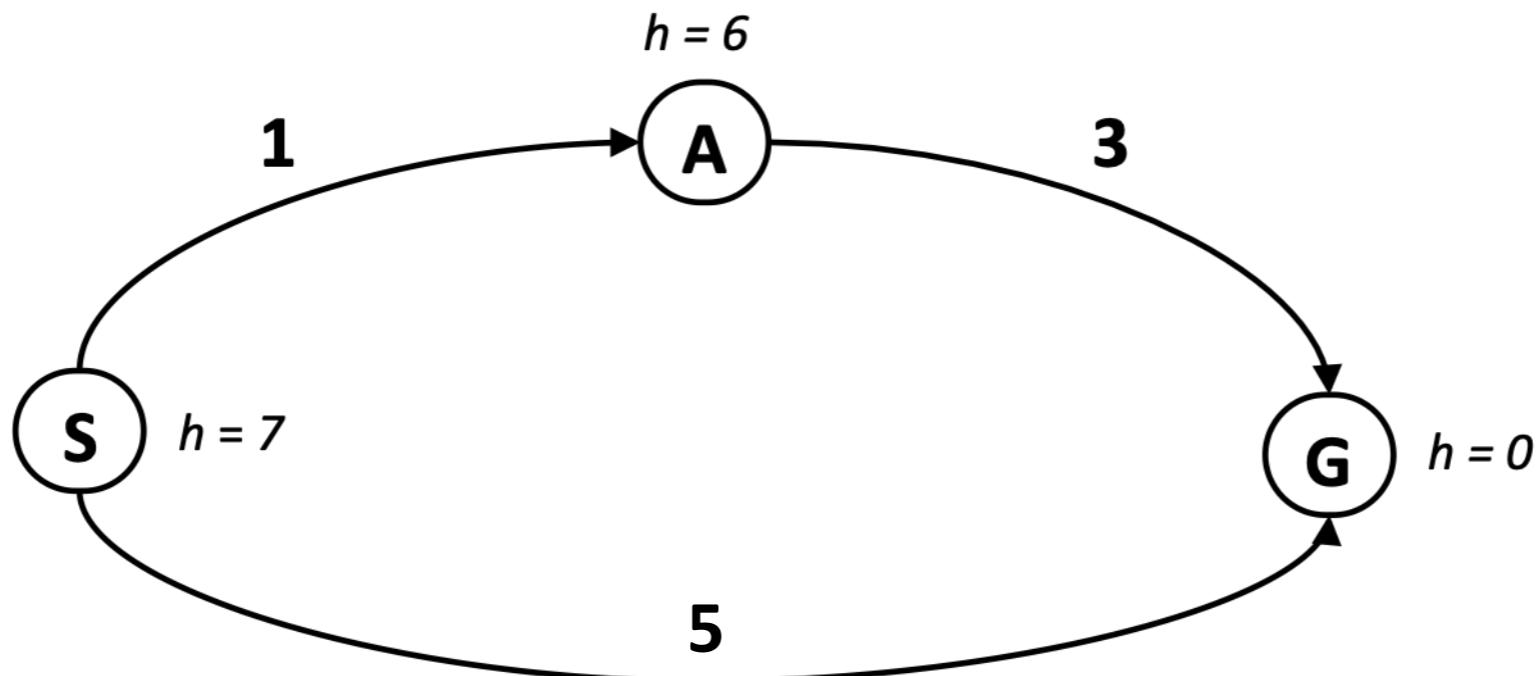
When should A* terminate?

- ❖ Should we stop when we enqueue a goal?



- ❖ No: only stop when we dequeue a goal

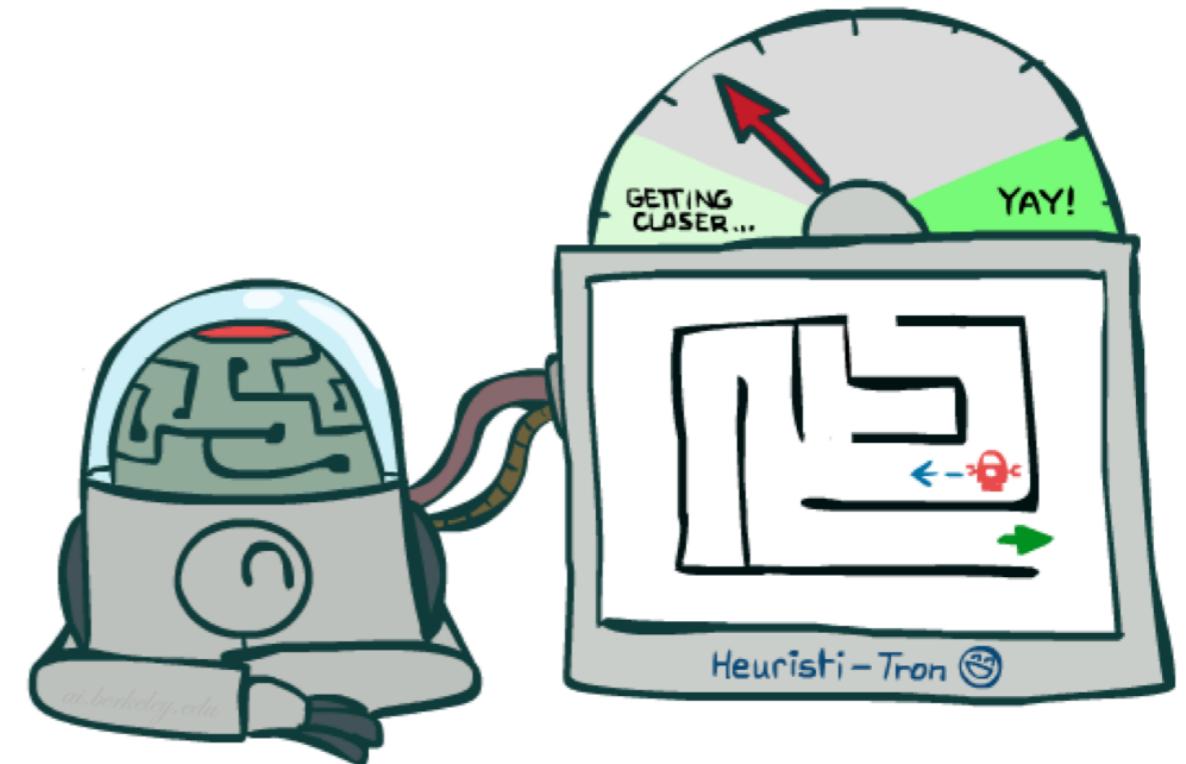
Is A* Optimal?



- ❖ What went wrong?
- ❖ Actual bad goal cost < estimated good goal cost
- ❖ We need estimates to be less than actual costs!

*Informed Search: A**

Admissible Heuristics



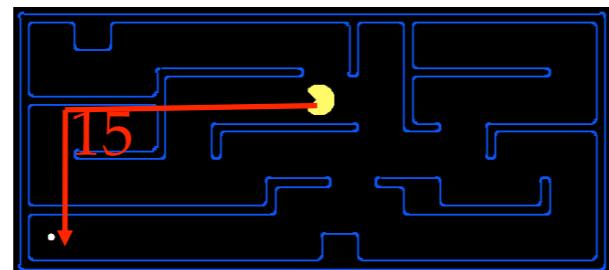
Admissible Heuristics

- ❖ A heuristic h is admissible (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

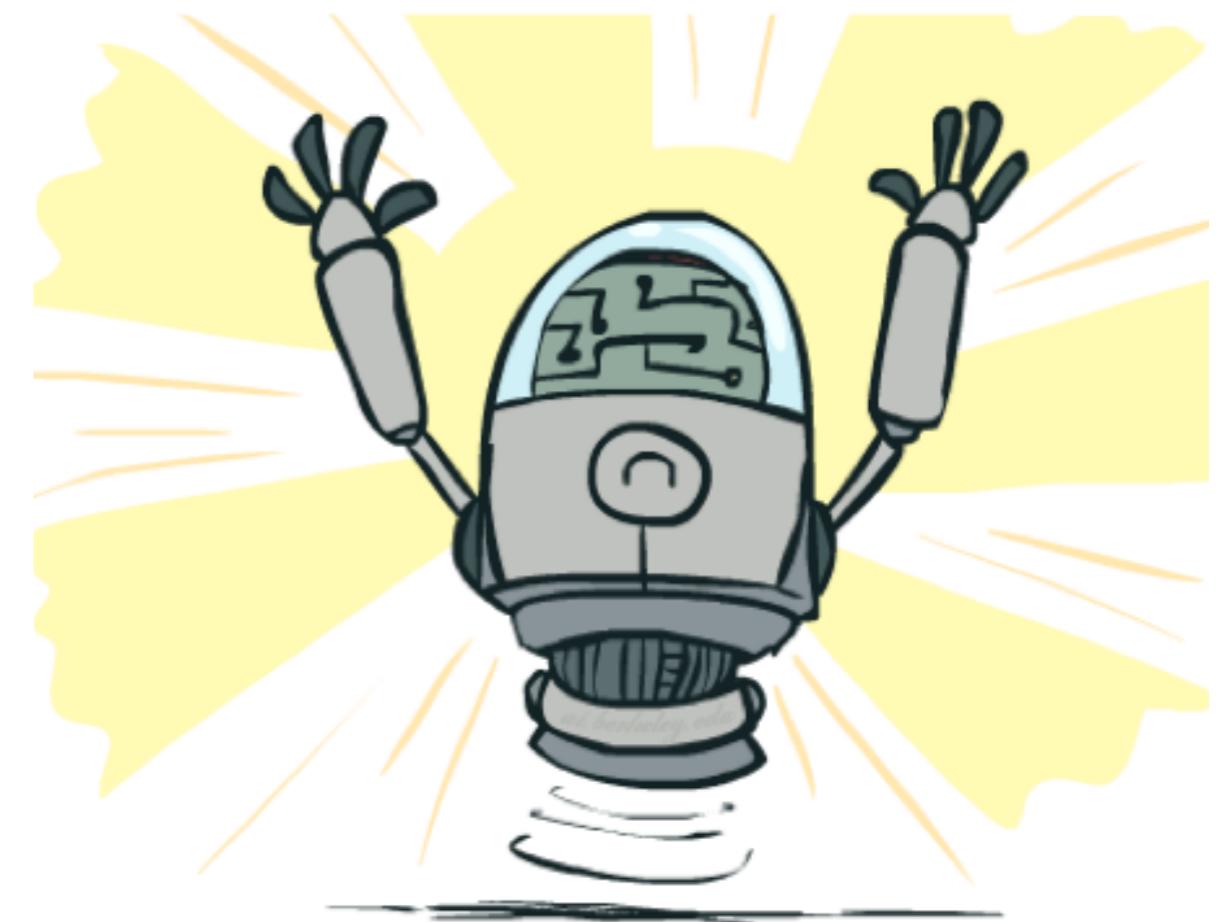
- ❖ Examples:



- ❖ Coming up with admissible heuristics is most of what's hard in using A* in practice.

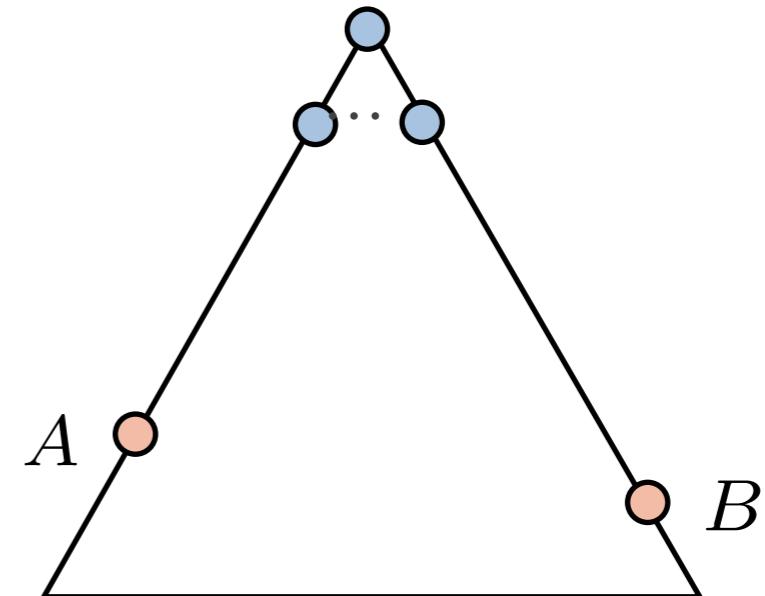
*Informed Search: A**

Optimality of A* Tree Search



Optimality of A* Tree Search

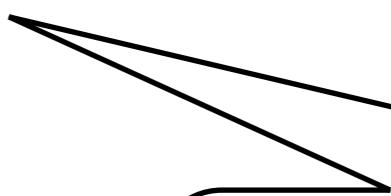
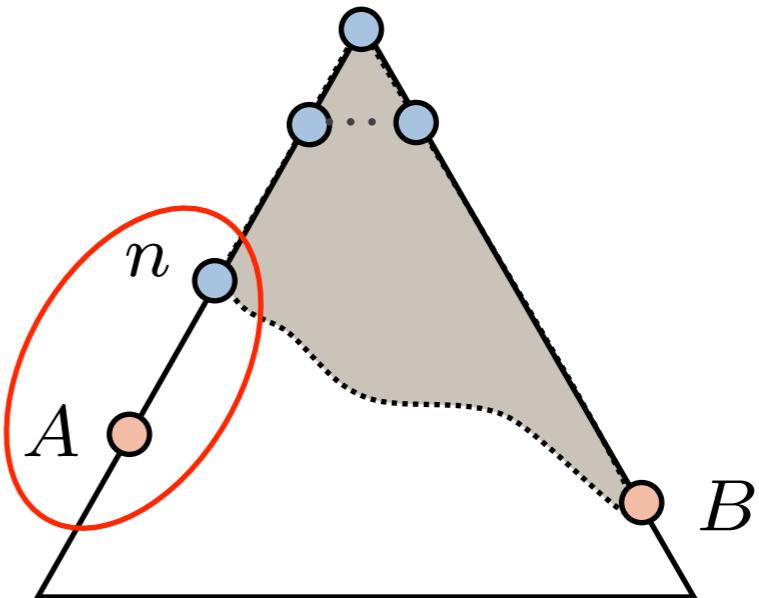
- ❖ **Assume:**
 - ❖ A is an optimal goal node
 - ❖ B is a suboptimal goal node
 - ❖ h is admissible
- ❖ **Claim:**
 - ❖ A will exit the fringe before B



Optimality of A* Tree Search: Blocking

Proof:

- ❖ Imagine B is on the fringe
- ❖ Some ancestor n of A is on the fringe, too (maybe A!)
- ❖ Claim: n will be expanded before B
 - ❖ $f(n)$ is less or equal to $f(A)$



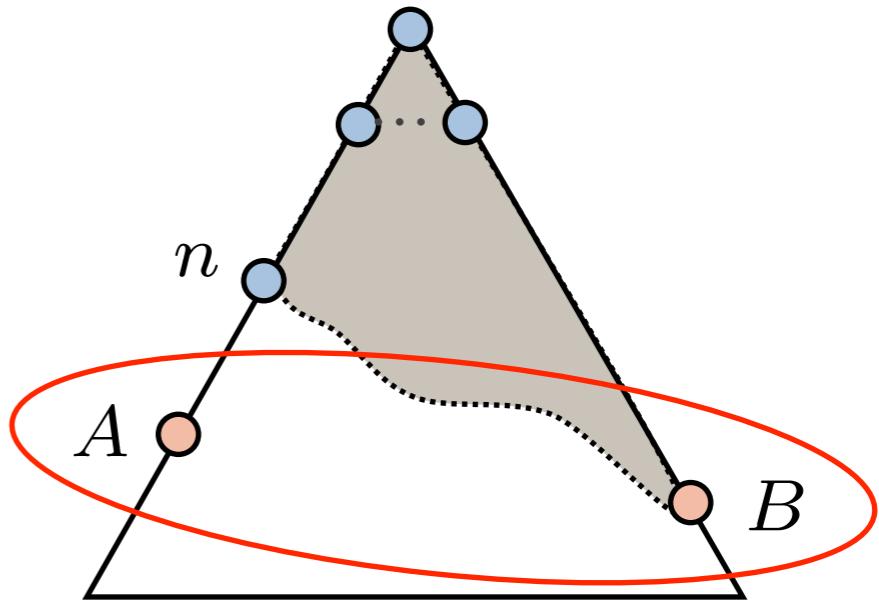
$$\begin{aligned}f(n) &= g(n) + h(n) \\f(n) &\leq g(A) \\g(A) &= f(A)\end{aligned}$$

Definition of f-cost
Admissibility of h
 $h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:

- ❖ Imagine B is on the fringe
- ❖ Some ancestor n of A is on the fringe, too (maybe A!)
- ❖ Claim: n will be expanded before B
 - ❖ $f(n)$ is less or equal to $f(A)$
 - ❖ $f(A)$ is less than $f(B)$



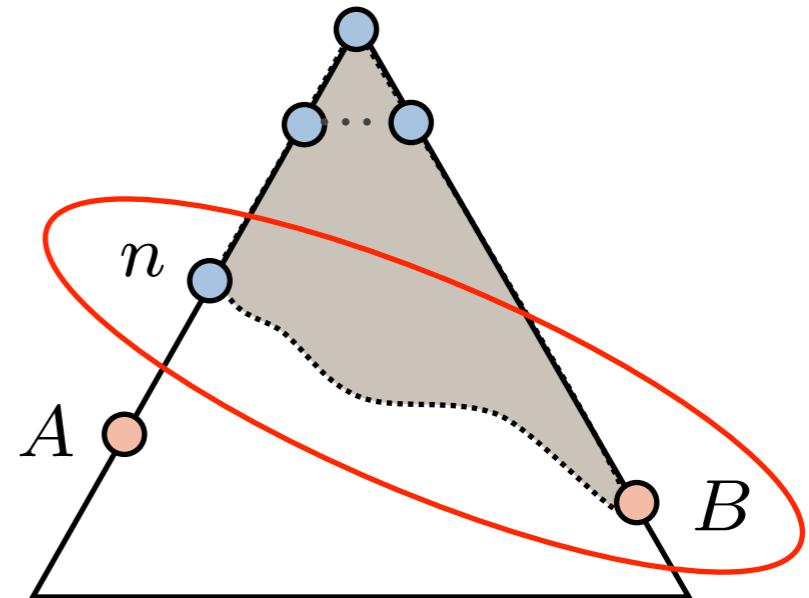
$$\begin{aligned} g(A) &< g(B) \\ f(A) &< f(B) \end{aligned}$$

B is suboptimal
 $h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:

- ❖ Imagine B is on the fringe
- ❖ Some ancestor n of A is on the fringe, too (maybe A!)
- ❖ Claim: n will be expanded before B
 - ❖ $f(n)$ is less or equal to $f(A)$
 - ❖ $f(A)$ is less than $f(B)$
 - ❖ n expands before B
 - ❖ All ancestors of A expand before B
 - ❖ A expands before B
 - ❖ A* search is optimal

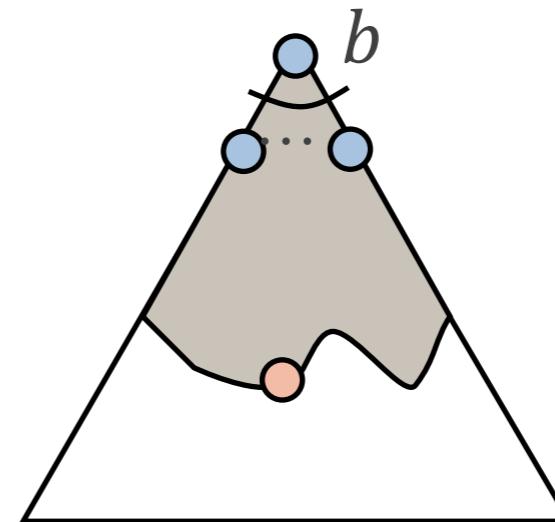


$$f(n) \leq f(A) < f(B)$$

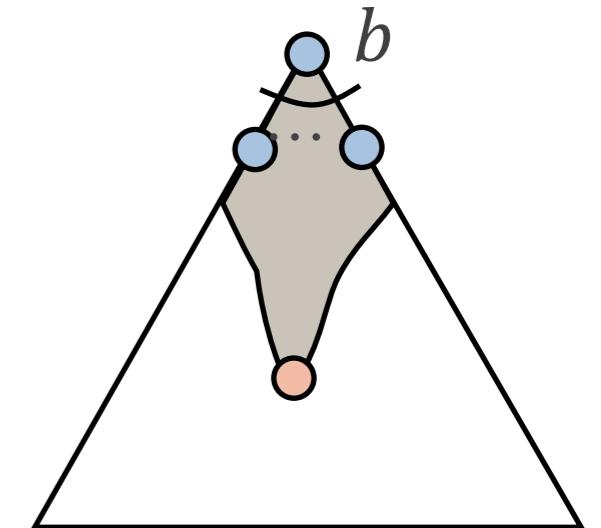
*Informed Search: A**

Properties of A*

Uniform-Cost

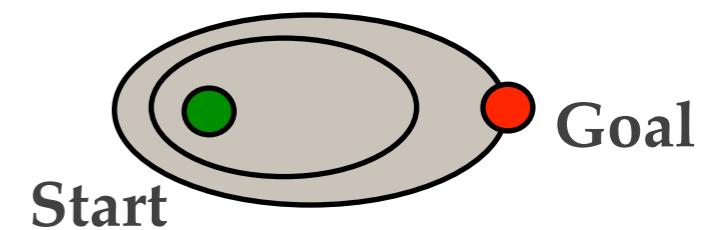
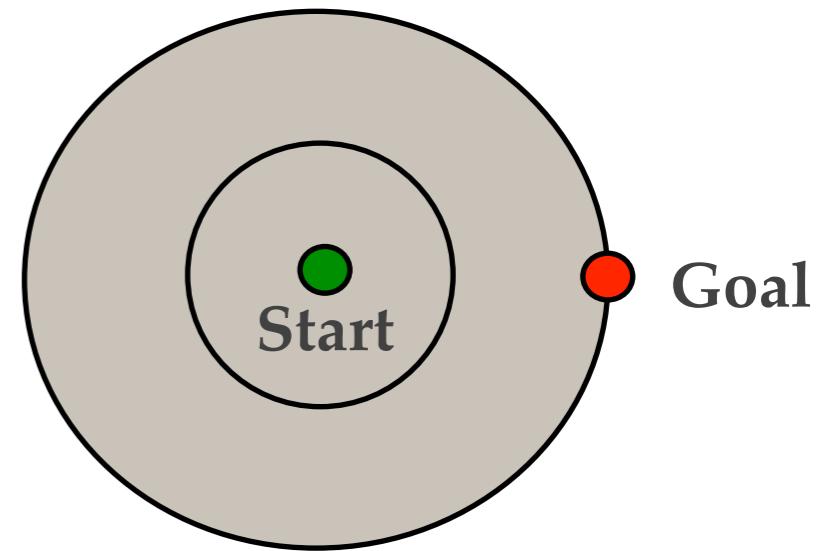


A*

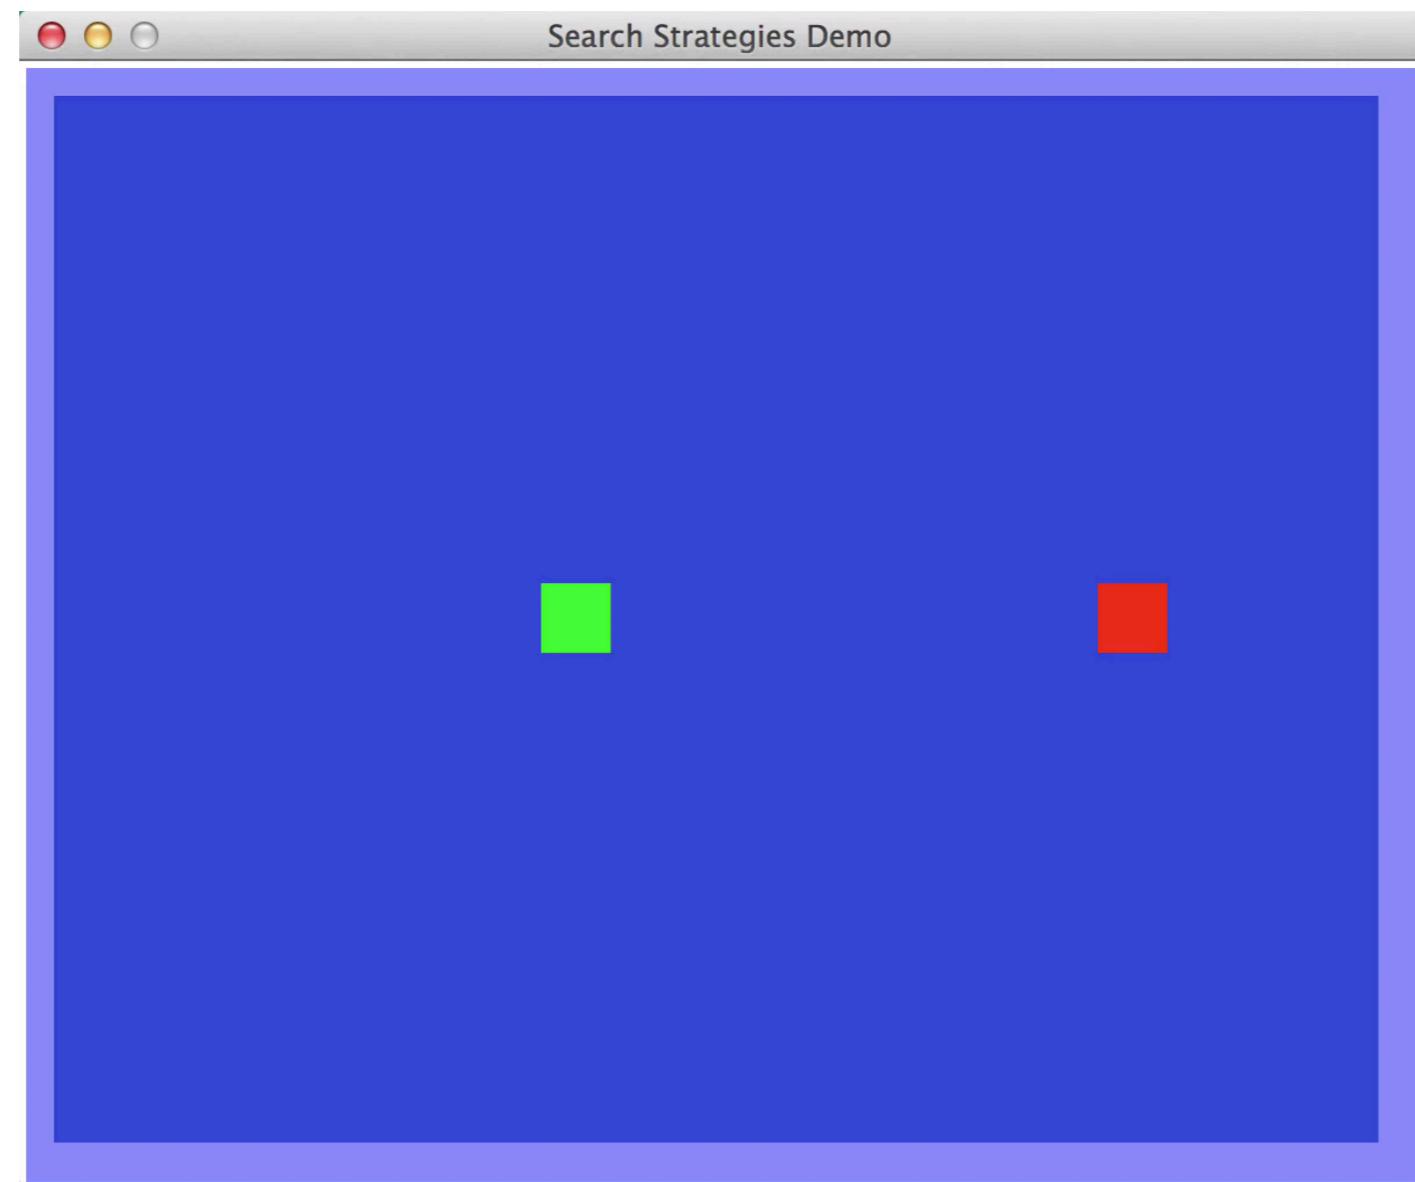


UCS vs A* Contours

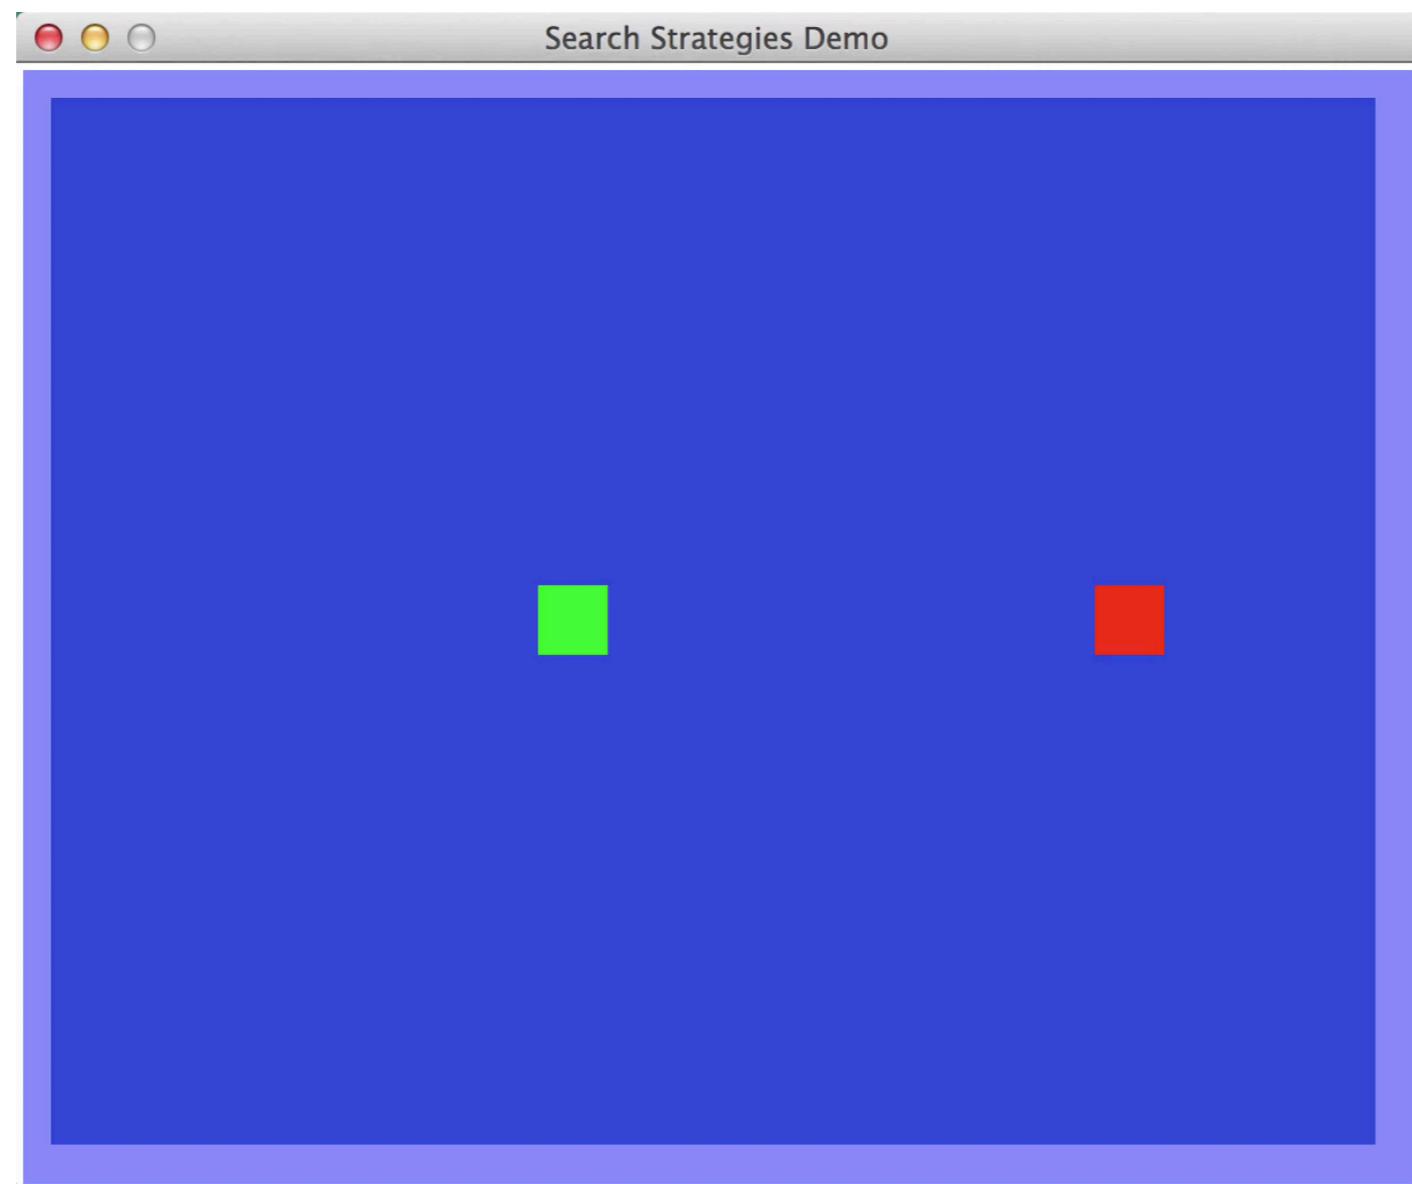
- ❖ Uniform-cost expands equally in all “directions”
- ❖ A* expands mainly toward the goal, but does hedge its bets to ensure optimality



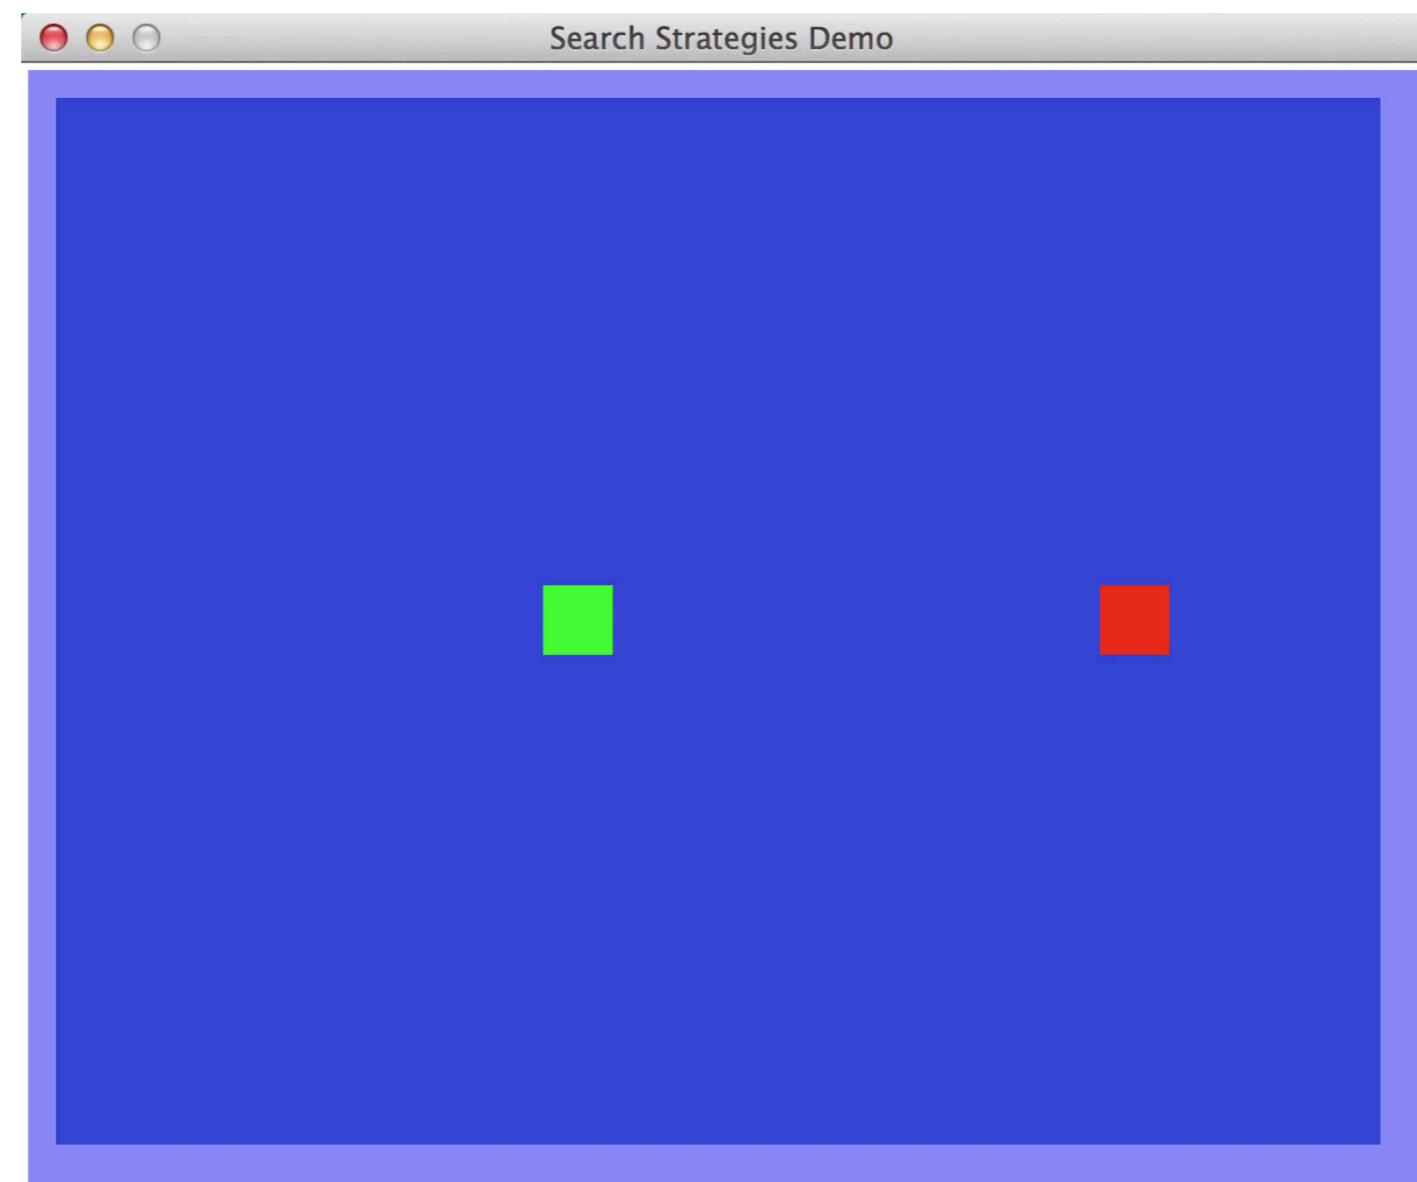
Video of Demo Contours (Empty) -- UCS



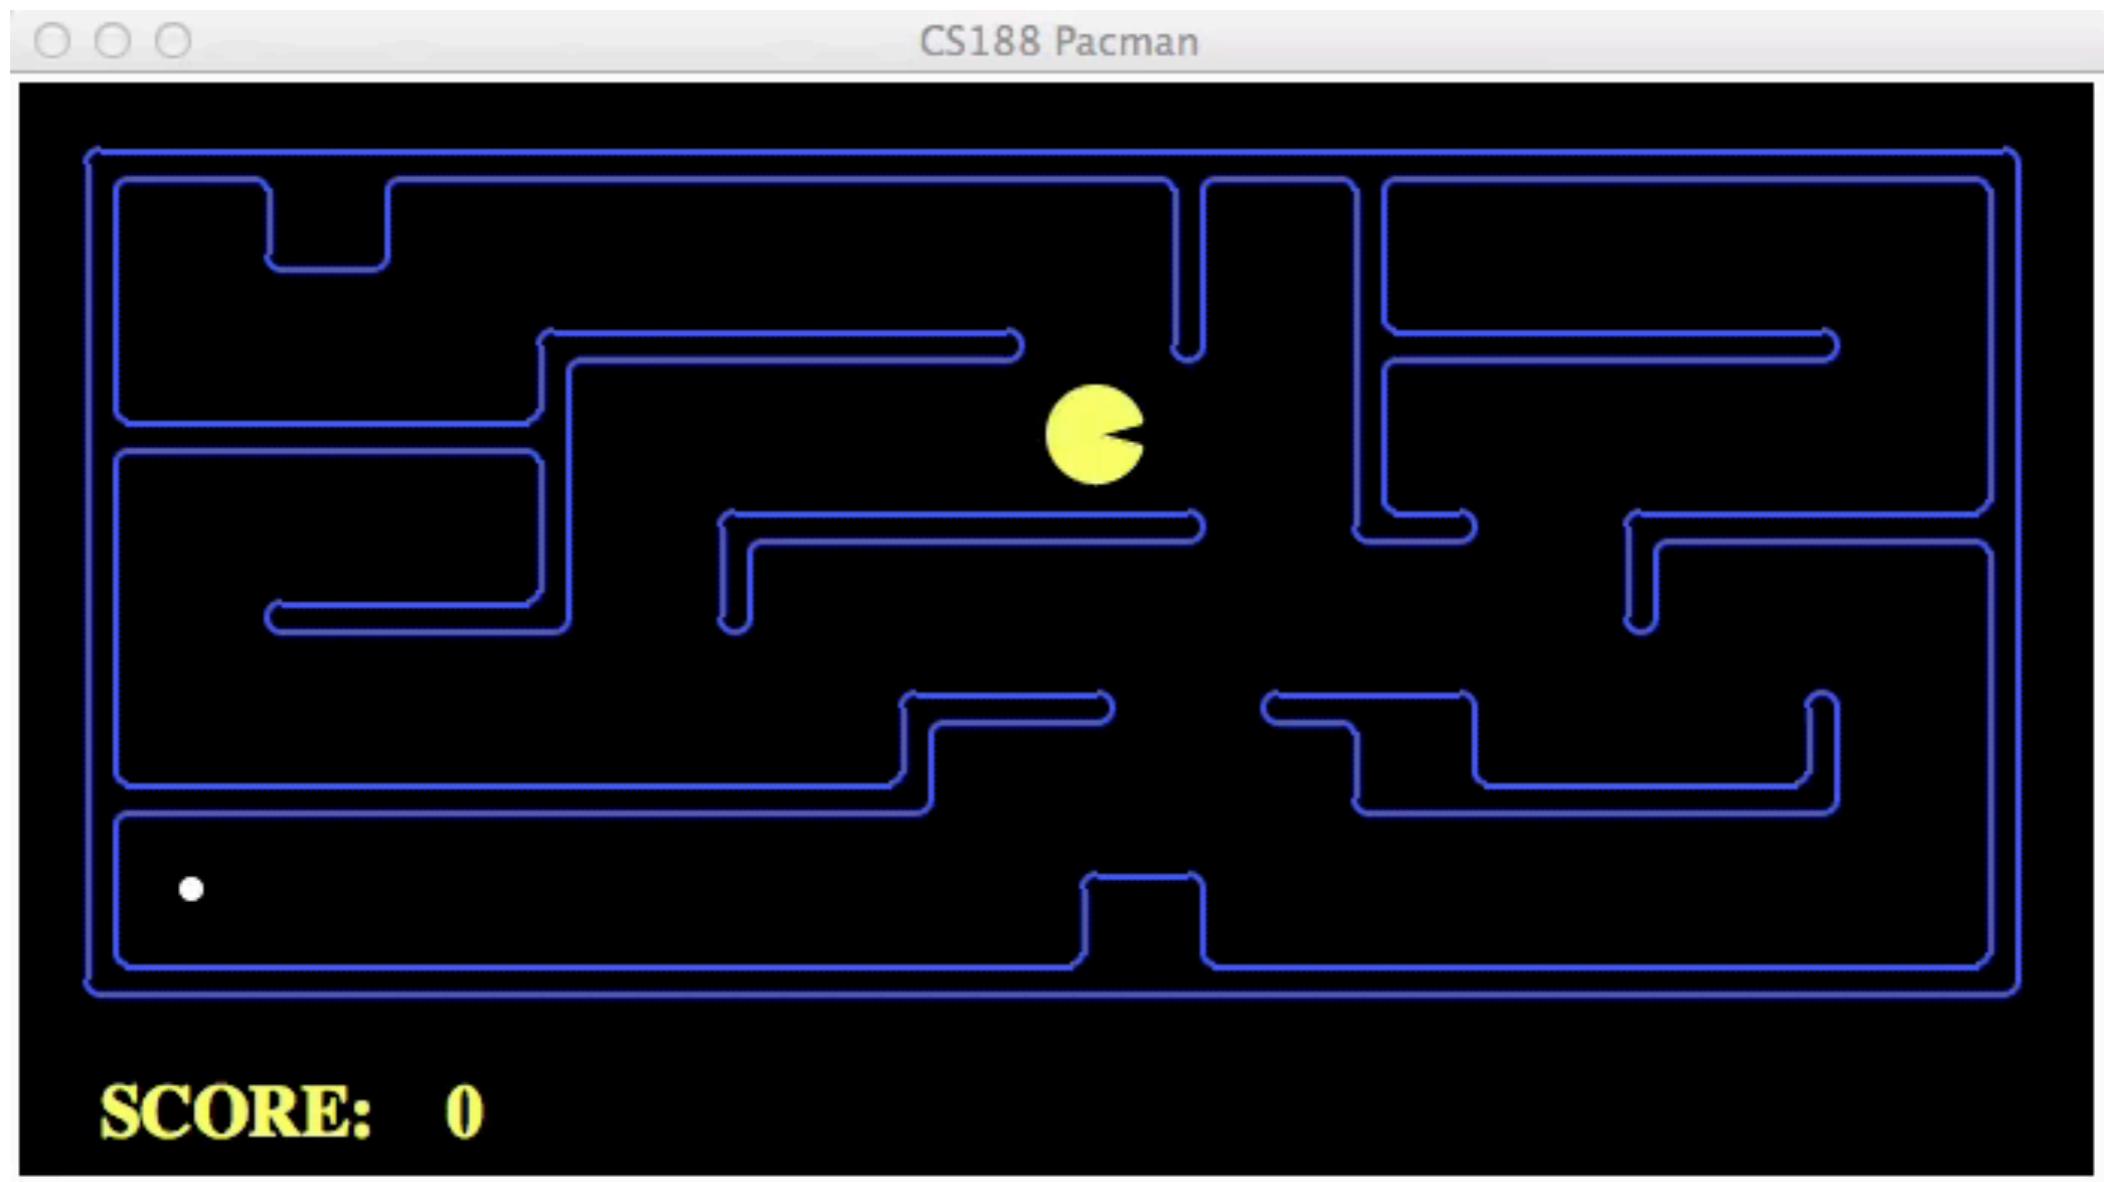
Video of Demo Contours (Empty) -- Greedy



Video of Demo Contours (Empty) – A*



Video of Demo Contours (Pacman Small Maze) – A*



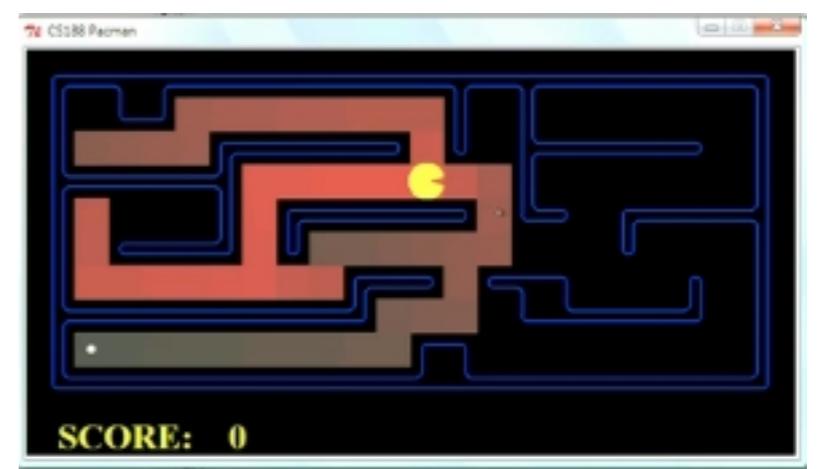
Comparison



Greedy



Uniform Cost



A*

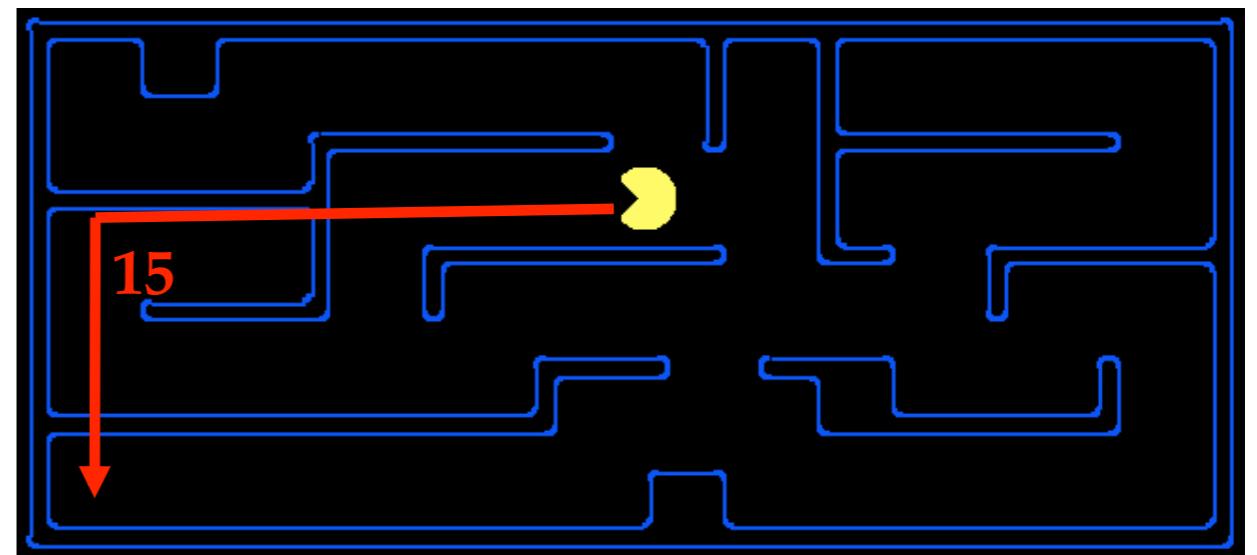
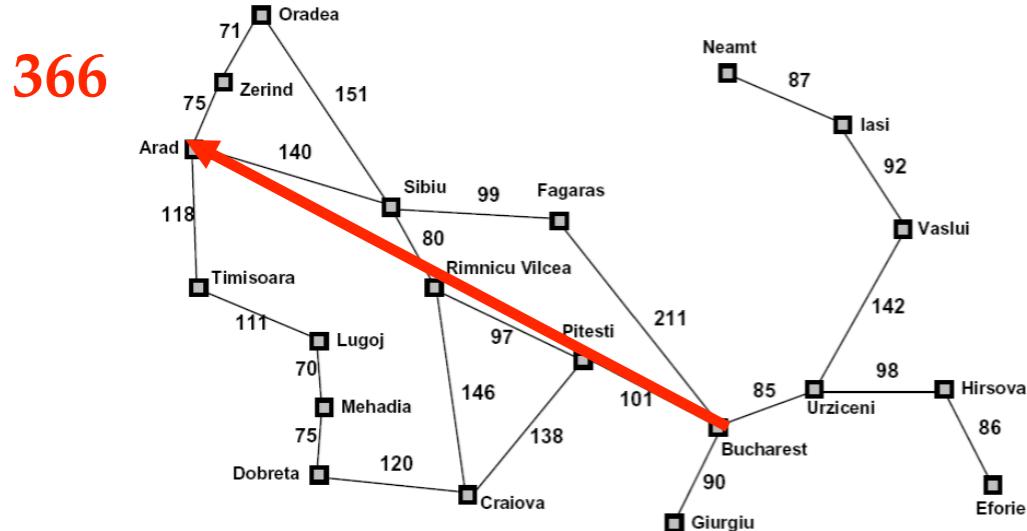
*Informed Search: A**

Creating Heuristics



Creating Admissible Heuristics

- ❖ Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- ❖ Often, admissible heuristics are solutions to relaxed problems, where new actions are available

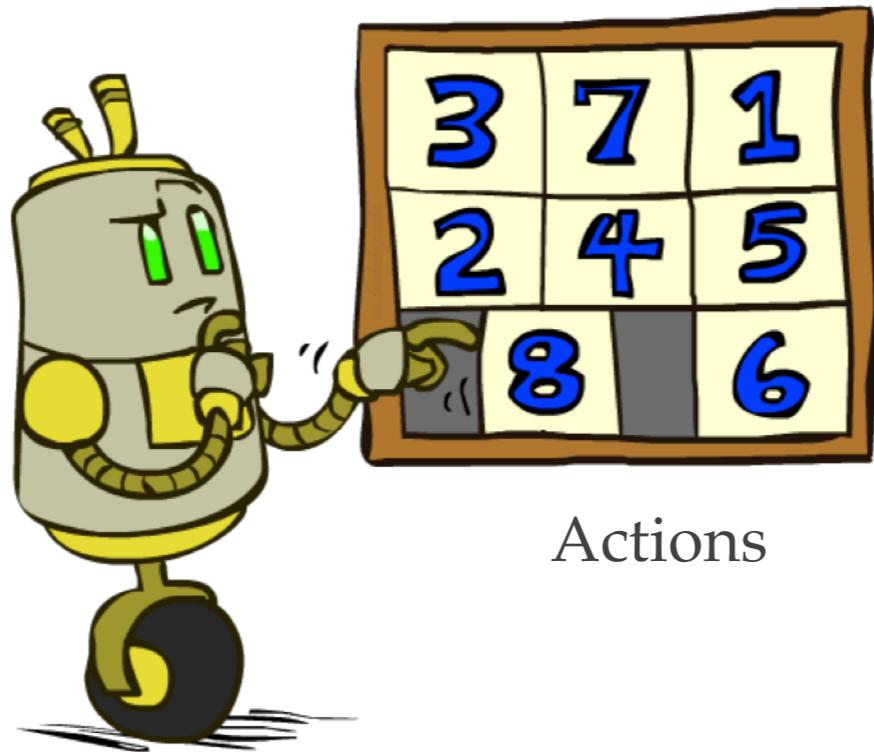


- ❖ Inadmissible heuristics are often useful too

Example: 8-Puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State



Actions

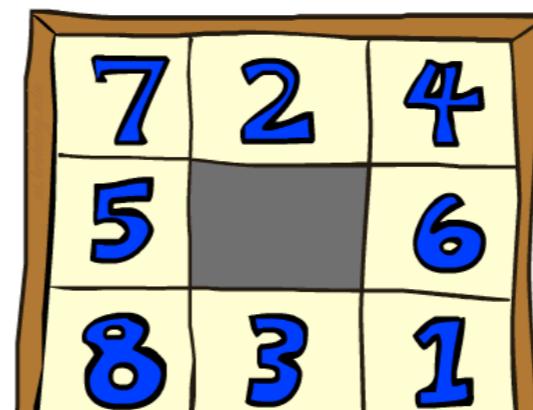
| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

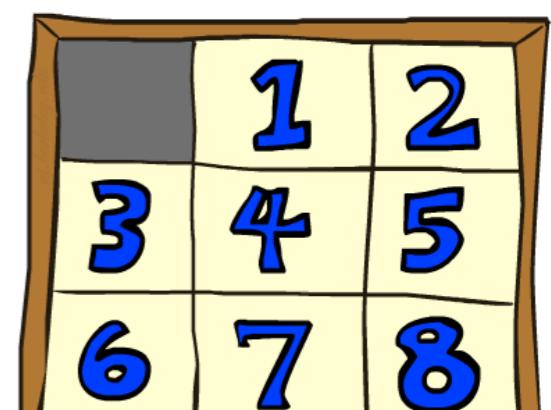
- ❖ What are the states?
- ❖ How many states?
- ❖ What are the actions?
- ❖ How many successors from the start state?
- ❖ What should the costs be?

8-Puzzle I

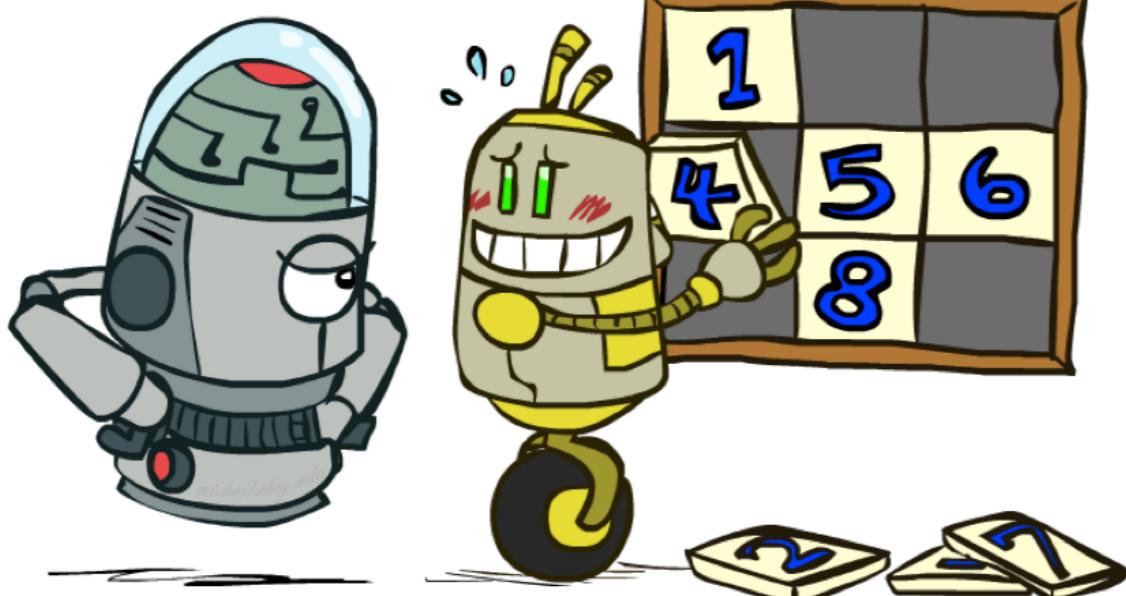
- ❖ Heuristic: Number of tiles misplaced
- ❖ $h(\text{start}) = 8$
- ❖ Why is it admissible?
- ❖ This is a relaxed-problem heuristic



Start State



Goal State



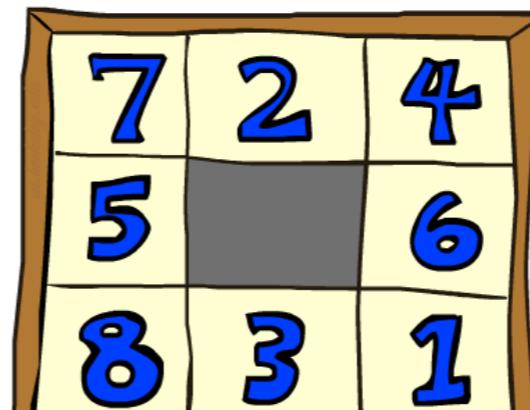
Average nodes expanded
when the optimal path has...

| | ...4 steps | ...8 steps | ...12 steps |
|-------|------------|------------|-------------------|
| UCS | 112 | 6,300 | 3.6×10^6 |
| TILES | 13 | 39 | 227 |

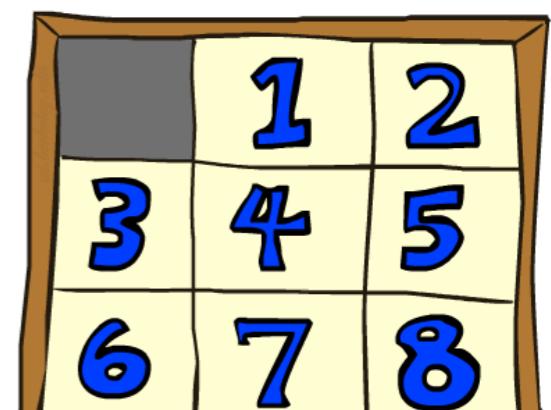
Statistics from Andrew Moore

8-Puzzle II

- ❖ What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?



Start State



Goal State

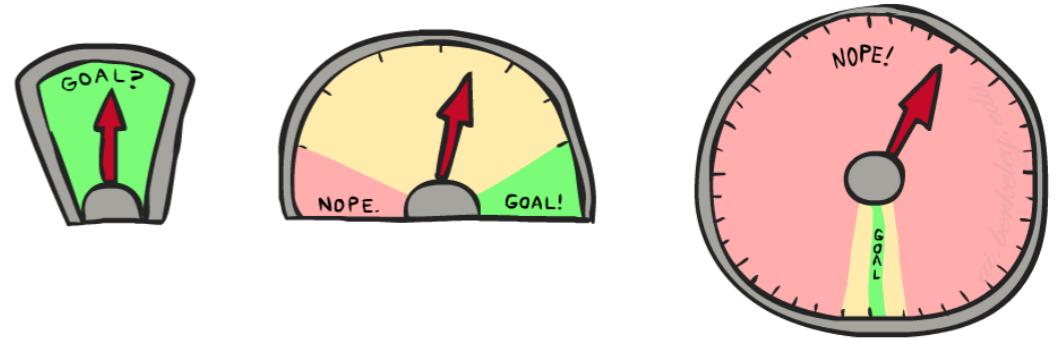
- ❖ Total Manhattan distance
- ❖ Why is it admissible?
- ❖ $h(\text{start}) = 3 + 1 + 2 + \dots = 18$

Average nodes expanded
when the optimal path has...

| | ...4 steps | ...8 steps | ...12 steps |
|-------|------------|------------|-------------|
| TILES | 13 | 39 | 227 |
| MD | 12 | 25 | 73 |

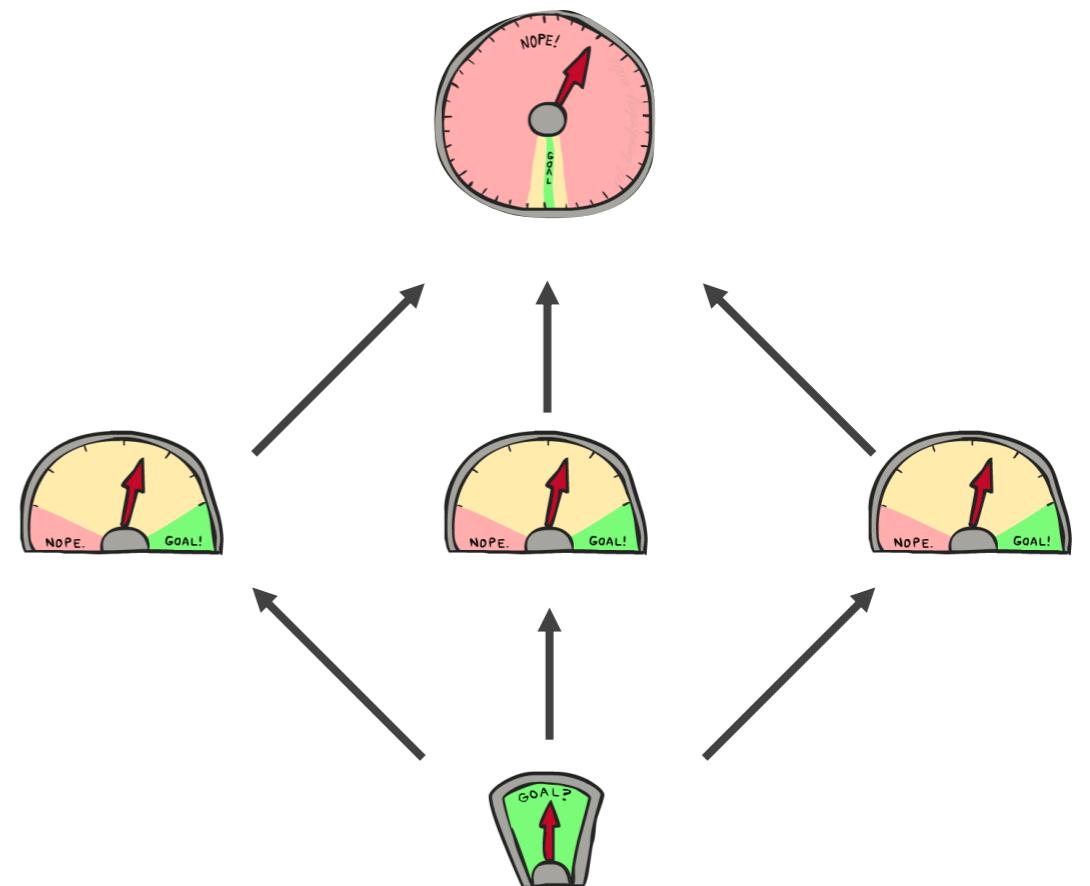
8-Puzzle III

- ❖ How about using the actual cost as a heuristic?
 - ❖ Would it be admissible?
 - ❖ Would we save on nodes expanded?
 - ❖ What's wrong with it?
- ❖ With A*: a trade-off between quality of estimate and work per node
 - ❖ As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself



*Informed Search: A**

Semi-Lattice of Heuristics



Trivial Heuristics, Dominance

- ❖ Dominance: $h_a \geq h_c$ if

$$\forall n : h_a(n) \geq h_c(n)$$

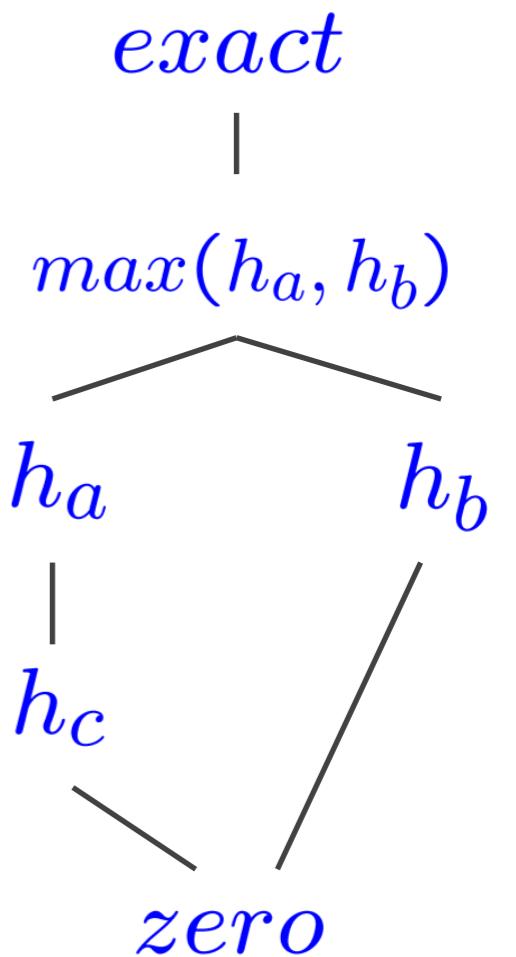
- ❖ Heuristics form a semi-lattice:

- ❖ Max of admissible heuristics is admissible

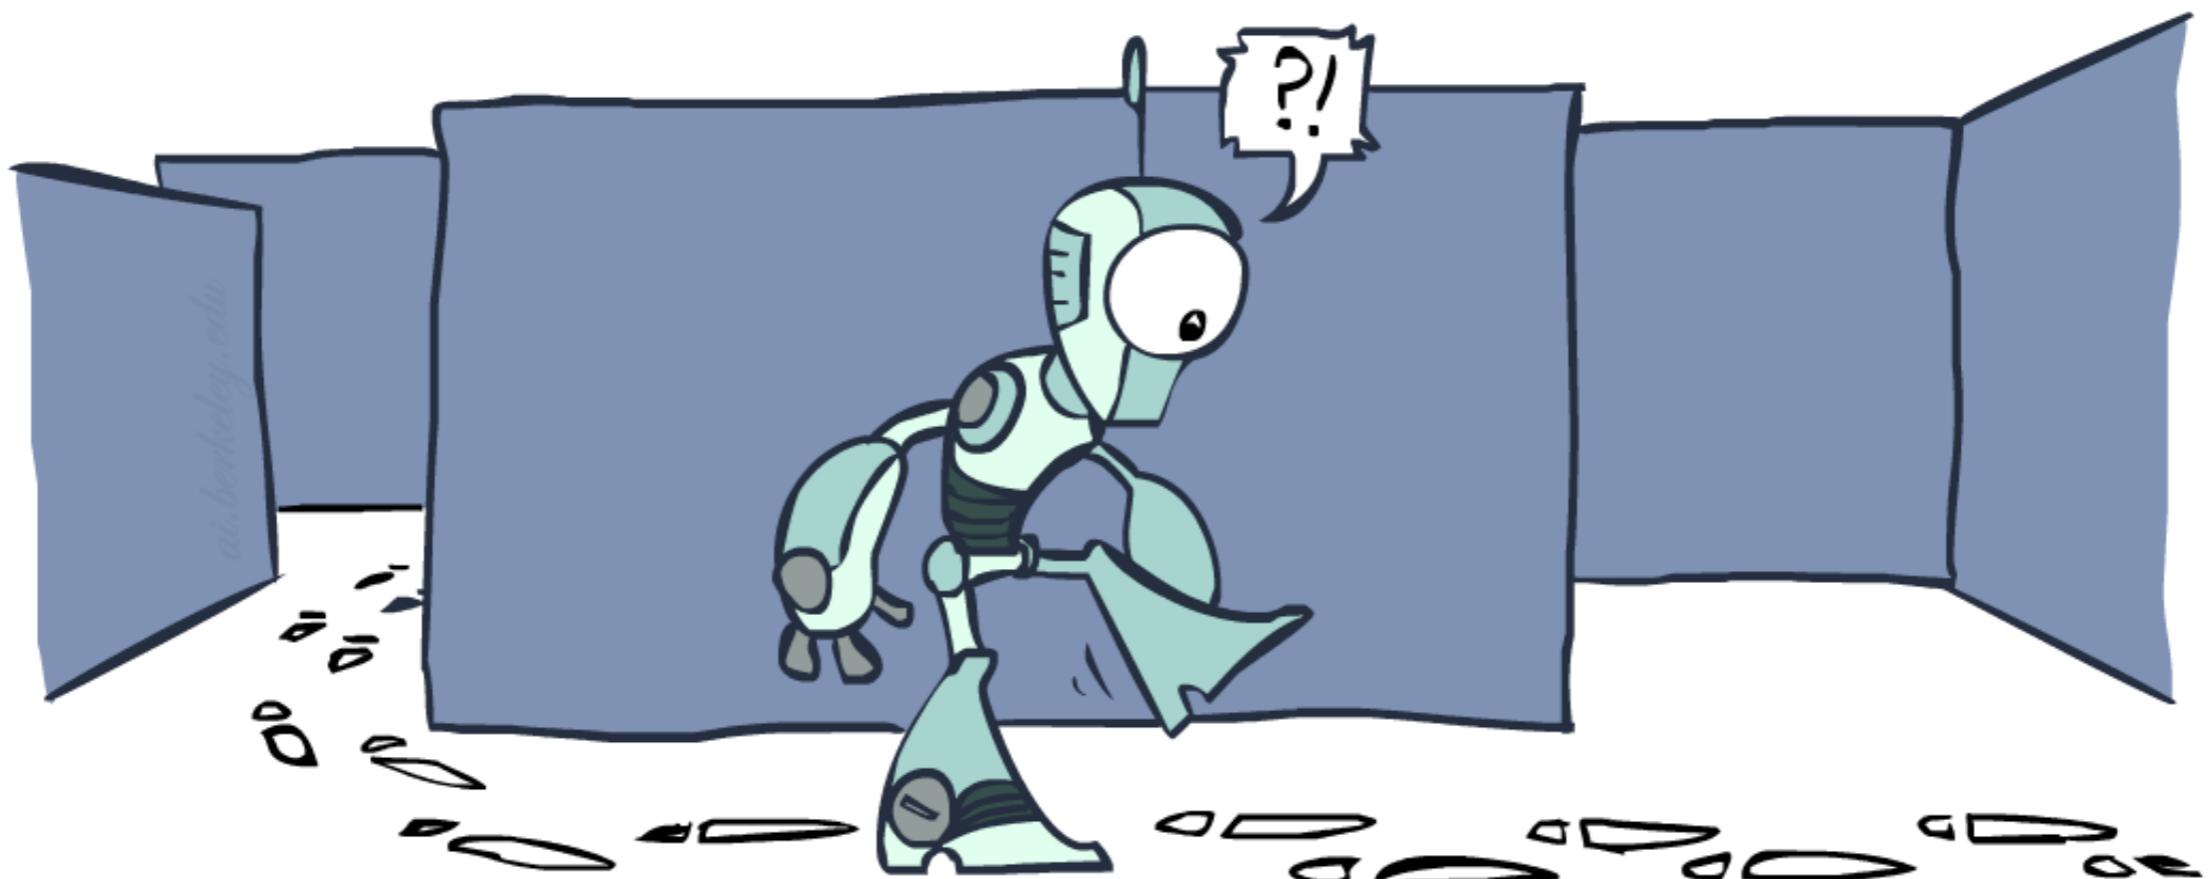
$$h(n) = \max(h_a(n), h_b(n))$$

- ❖ Trivial heuristics:

- ❖ Bottom of lattice is the zero heuristic (what does this give us?)
 - ❖ Top of lattice is the exact heuristic

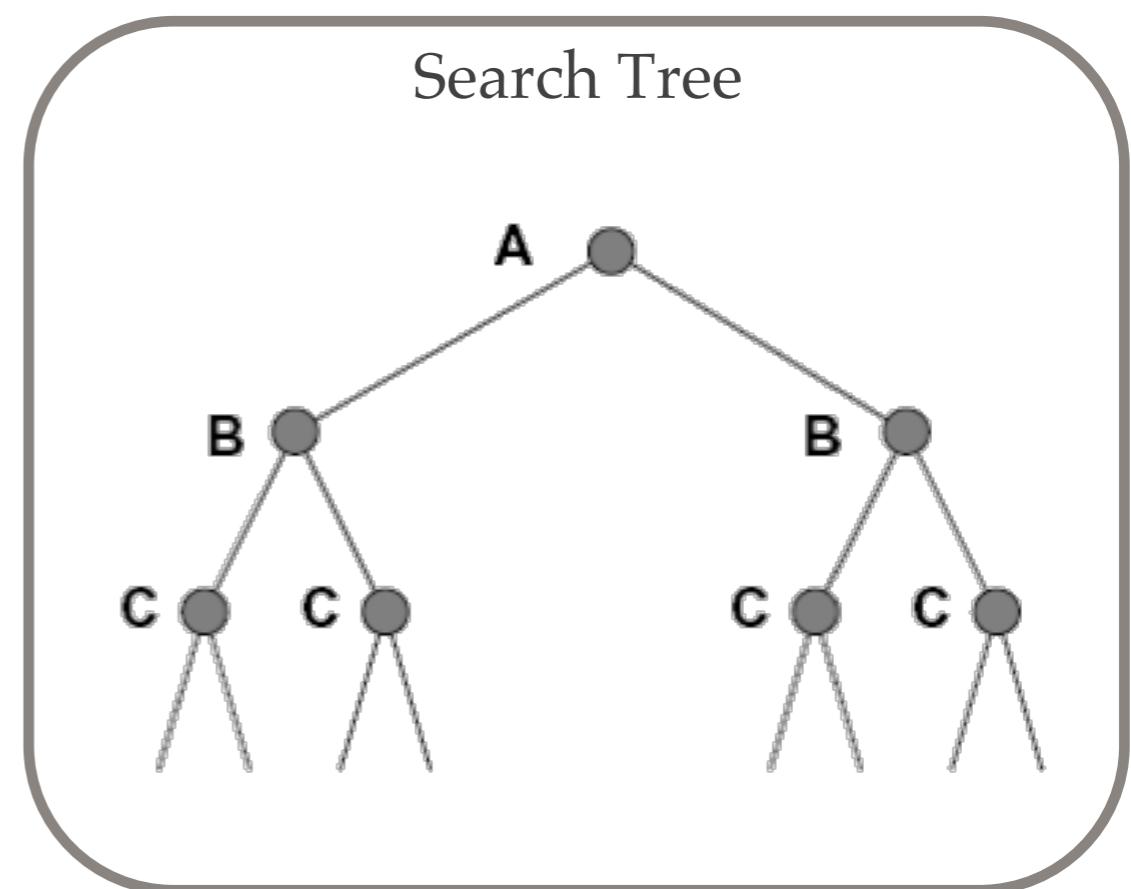
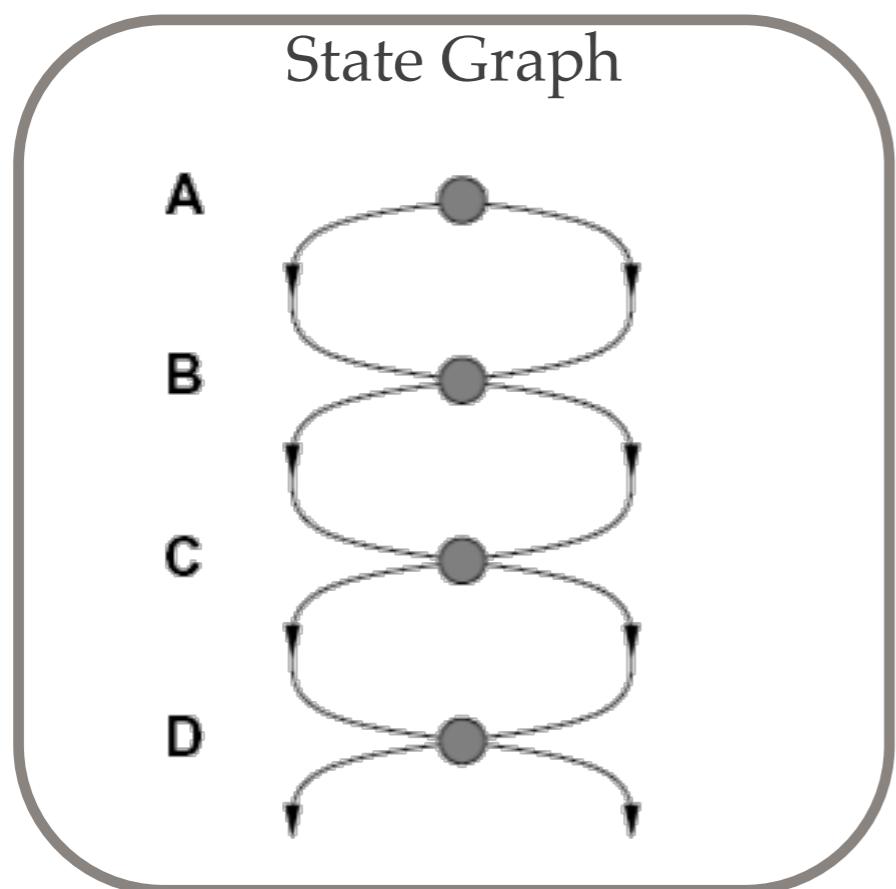


Graph Search



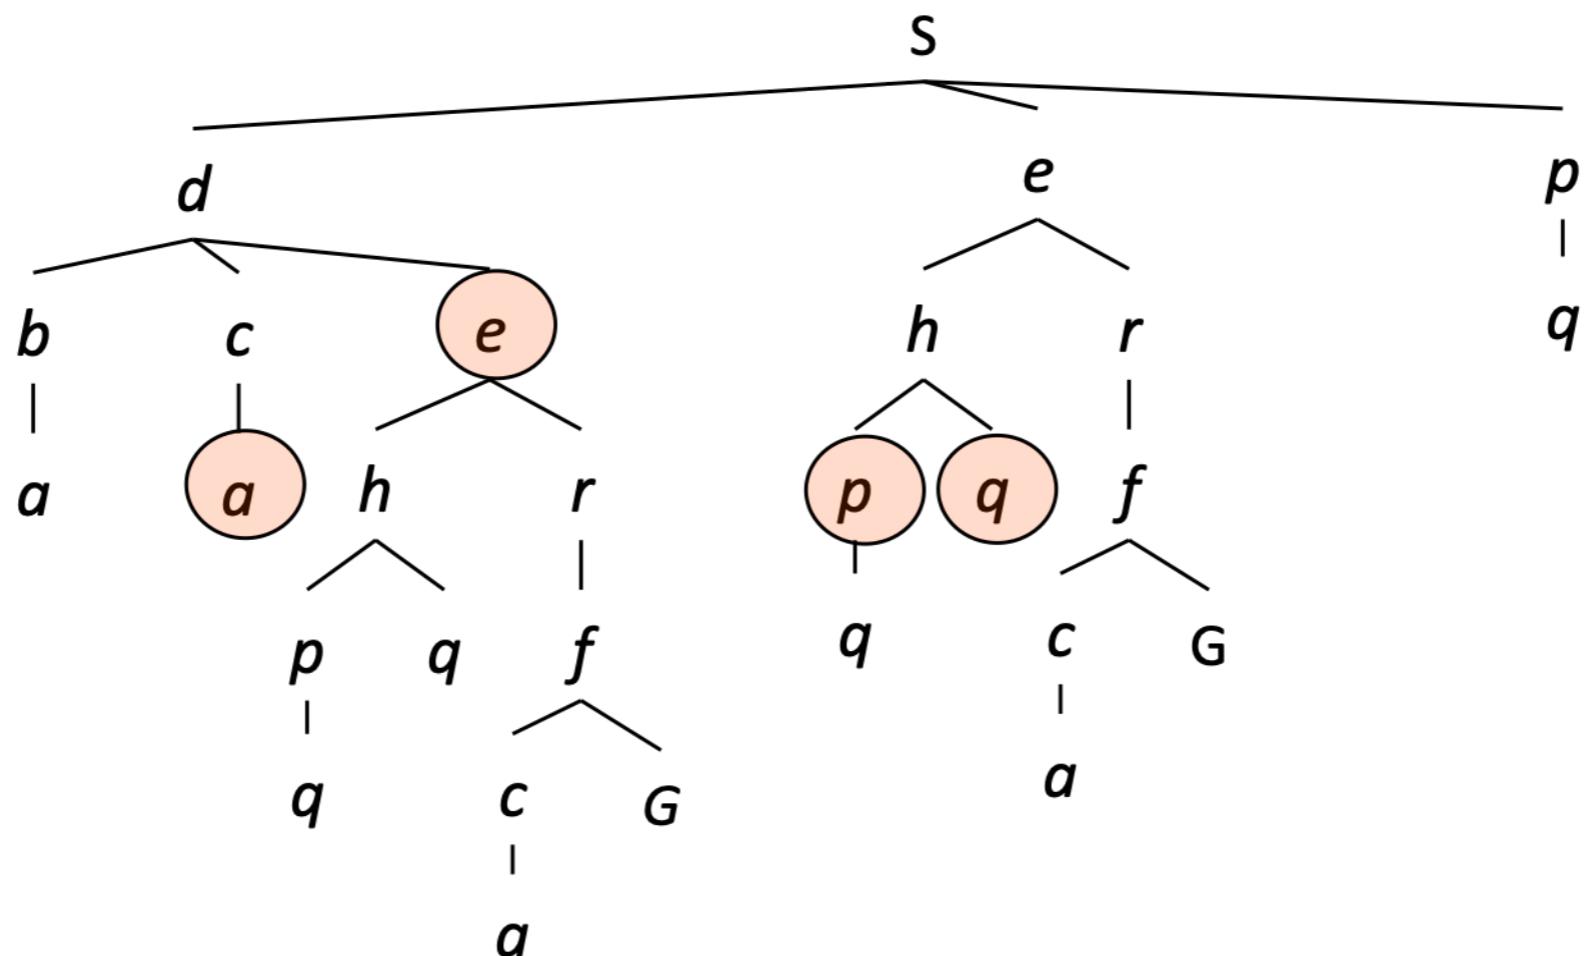
Tree Search: Extra Work!

- ❖ Failure to detect repeated states can cause exponentially more work.



Graph Search

- ❖ In BFS, for example, we shouldn't bother expanding the circled nodes (why?)

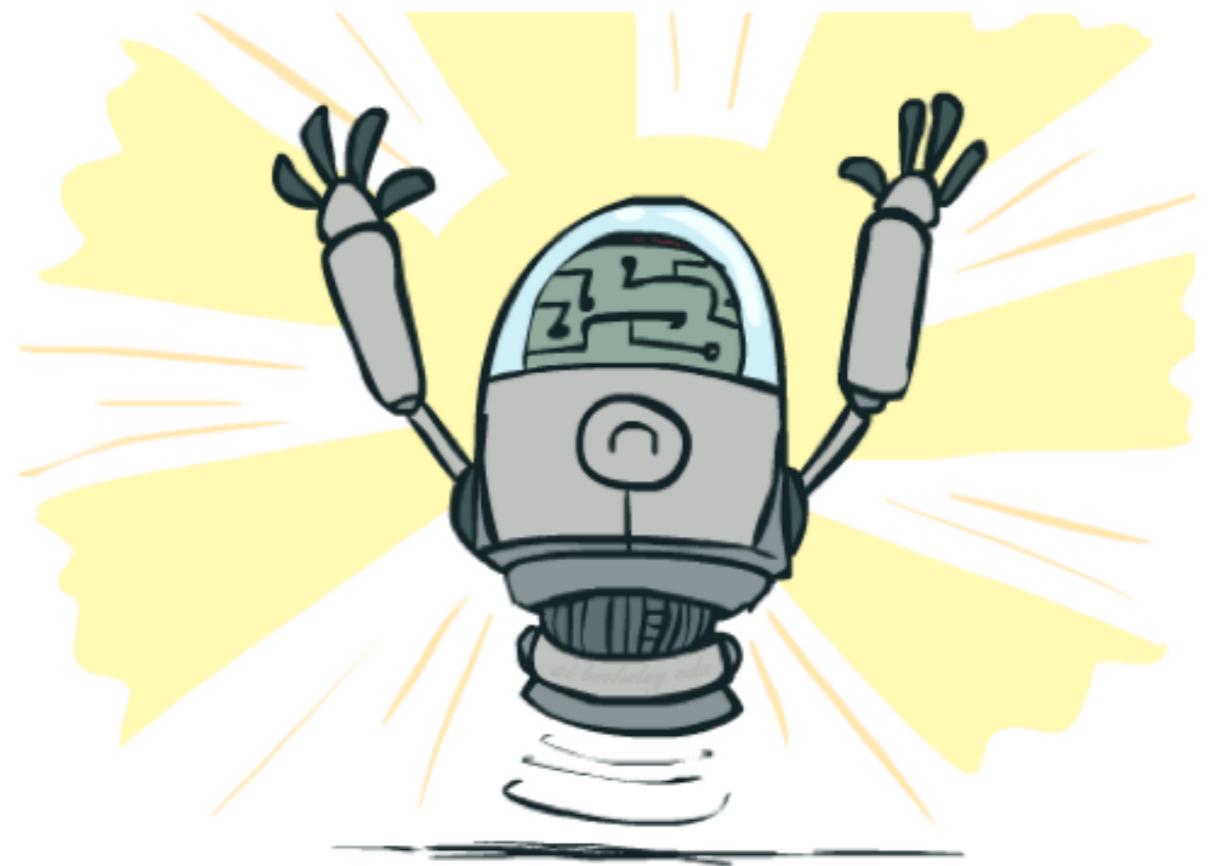


Graph Search

- ❖ Idea: never **expand** a state twice
- ❖ How to implement:
 - ❖ Tree search + set of expanded states (“closed set”)
 - ❖ Expand the search tree node-by-node, but...
 - ❖ Before expanding a node, check to make sure its state has never been expanded before
 - ❖ If not new, skip it, if new add to closed set
- ❖ Important: store the closed set as a set, not a list
- ❖ Can graph search wreck completeness? Why / why not?
- ❖ How about optimality?

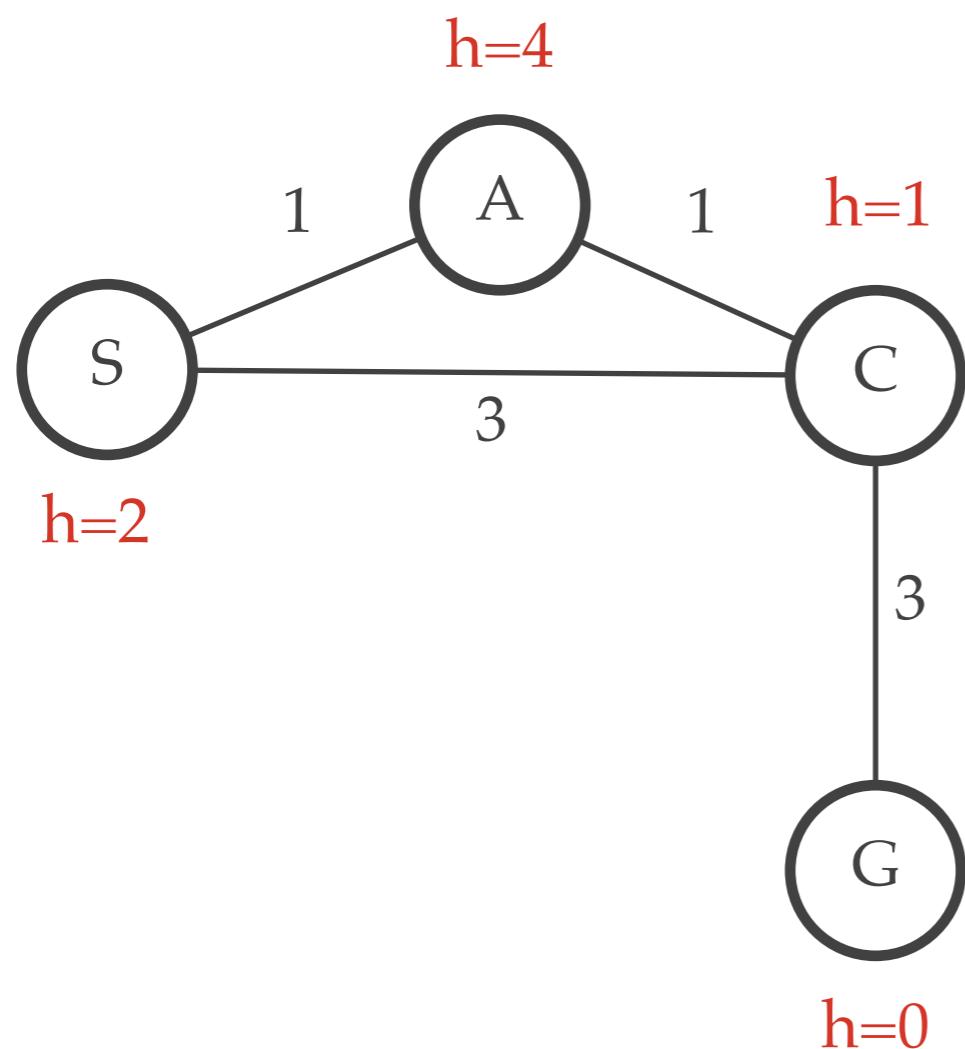
Graph Search

Optimality of A* Graph Search

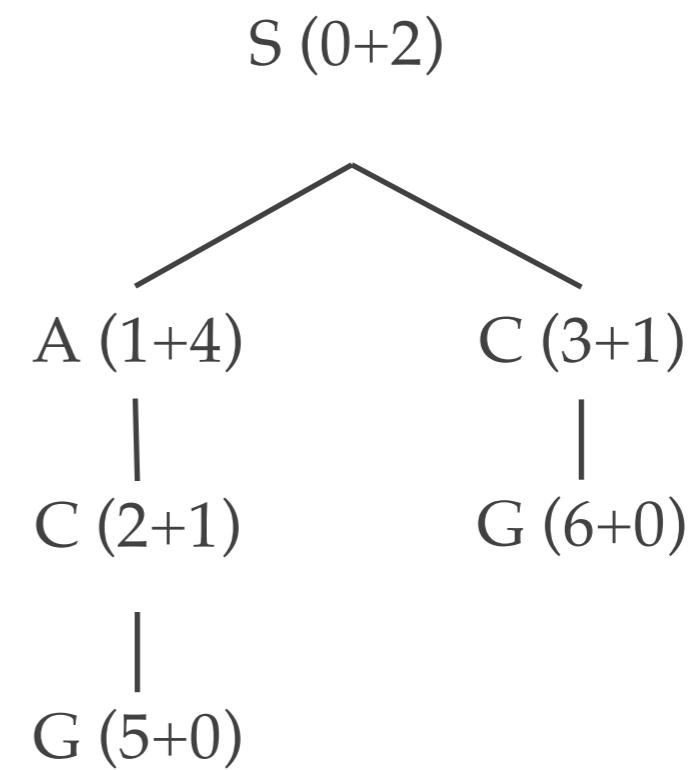


A* Tree Search

State space graph

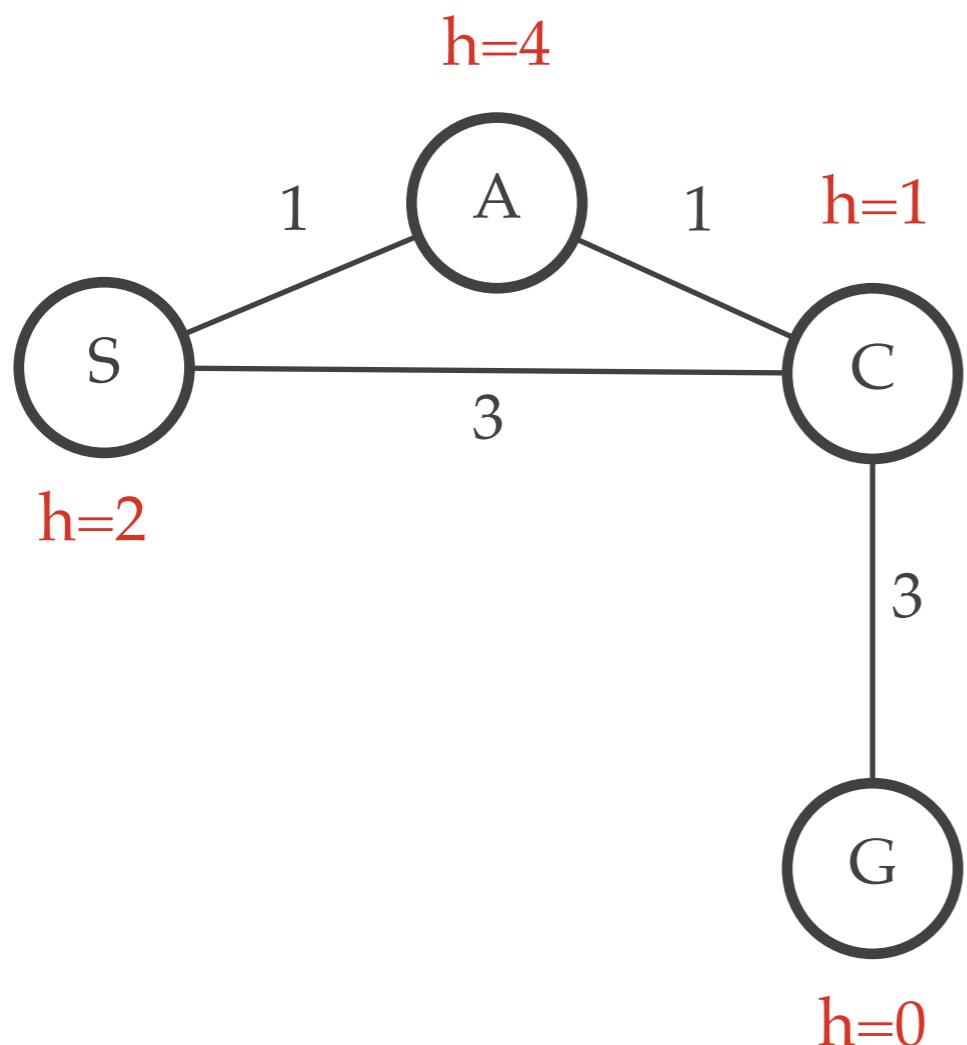


Search tree



Quiz: A* Graph Search

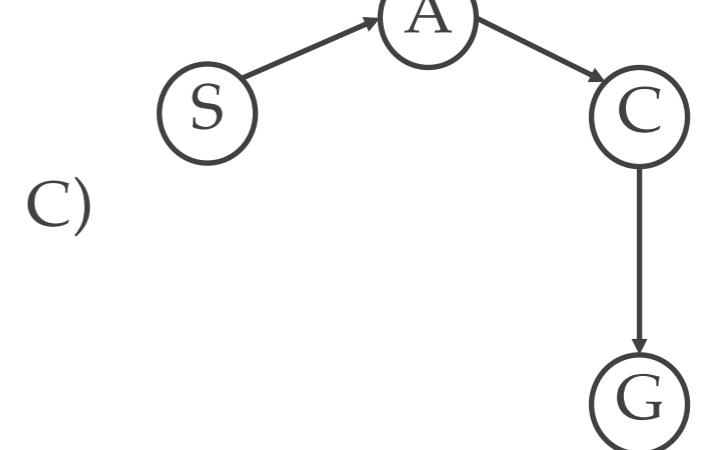
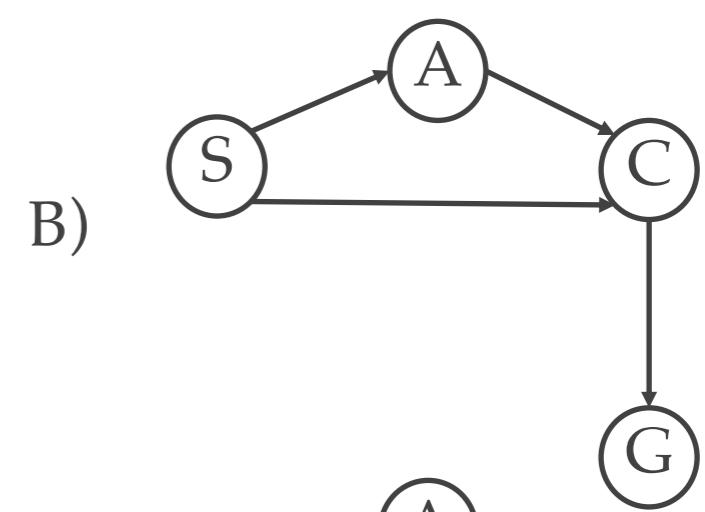
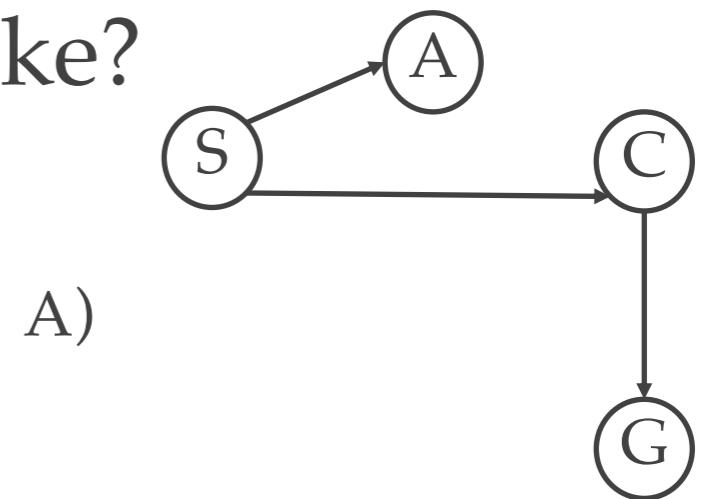
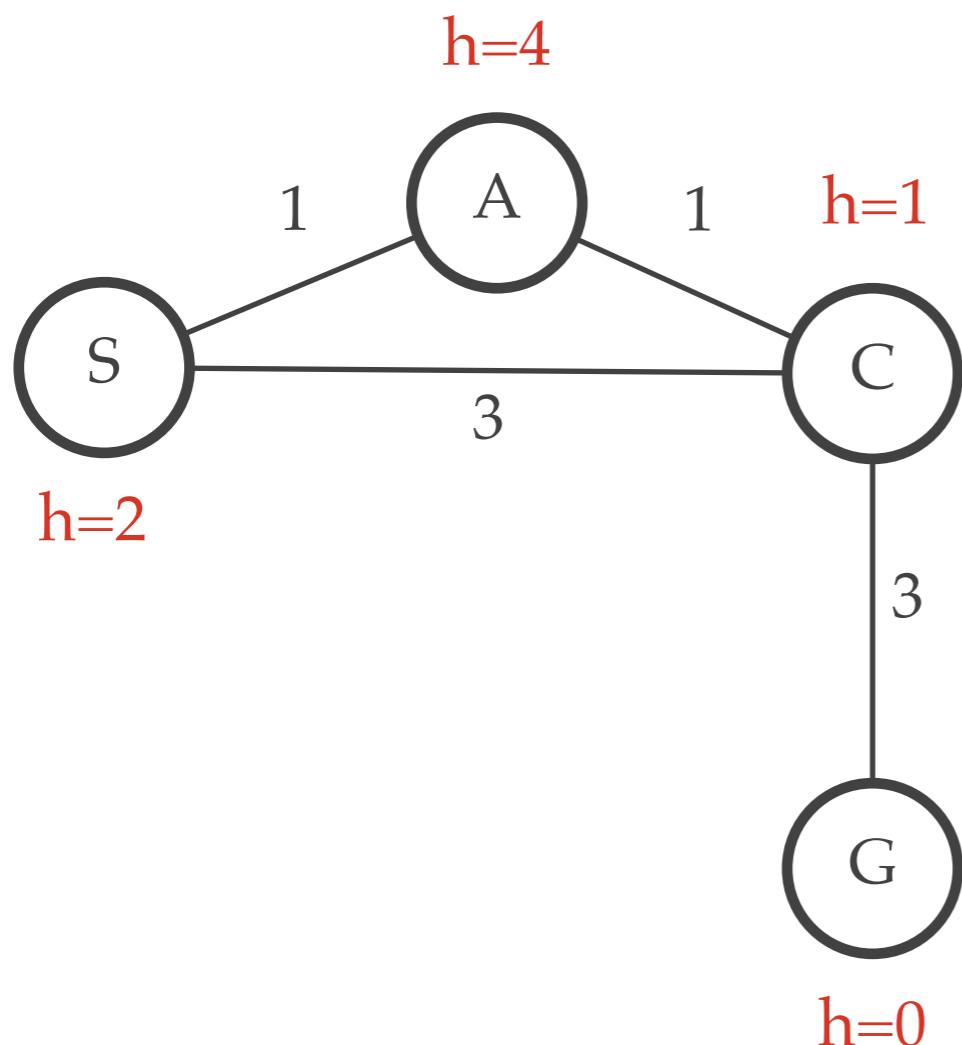
What paths does A* graph search consider during its search?



- A) S, S-A, S-C, S-C-G
- B) S, S-A, S-C, S-A-C, S-C-G
- C) S, S-A, S-A-C, S-A-C-G
- D) S, S-A, S-C, S-A-C, S-A-C-G

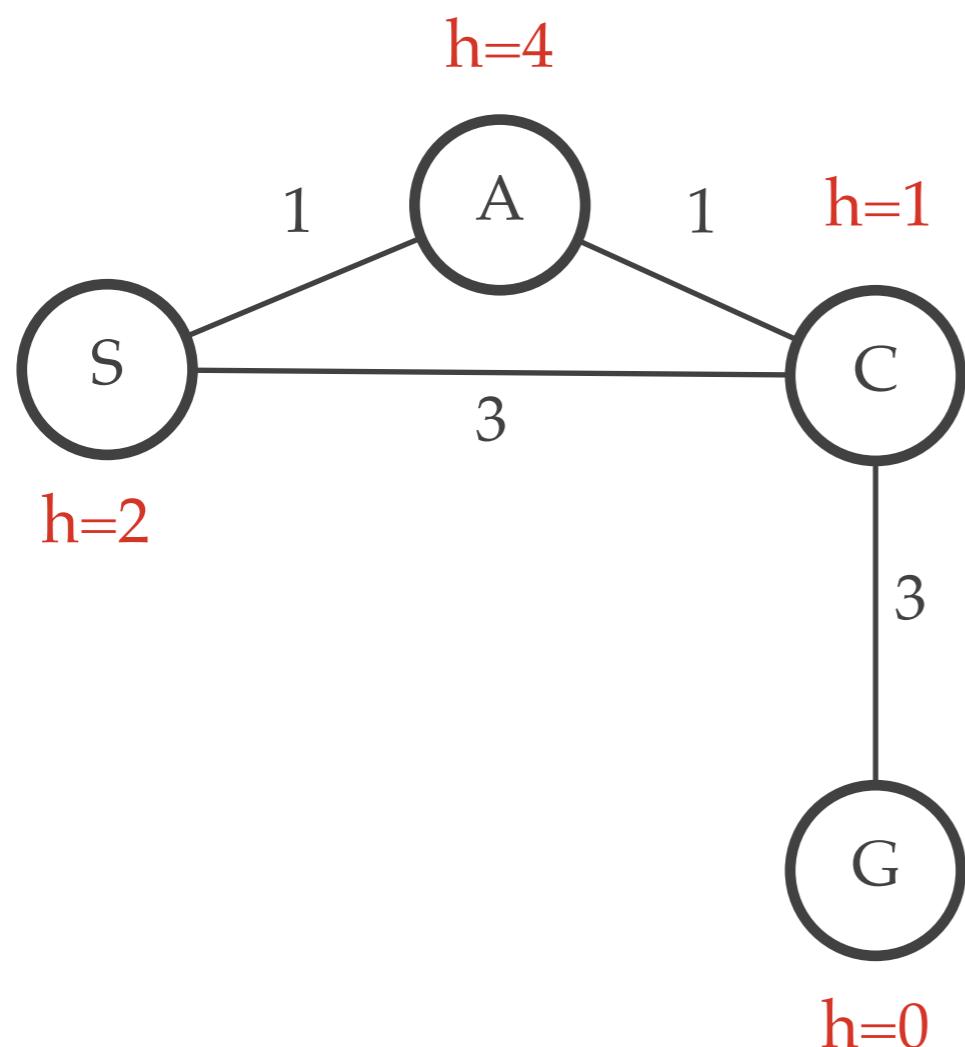
Quiz: A* Graph Search

What does the resulting graph tree look like?

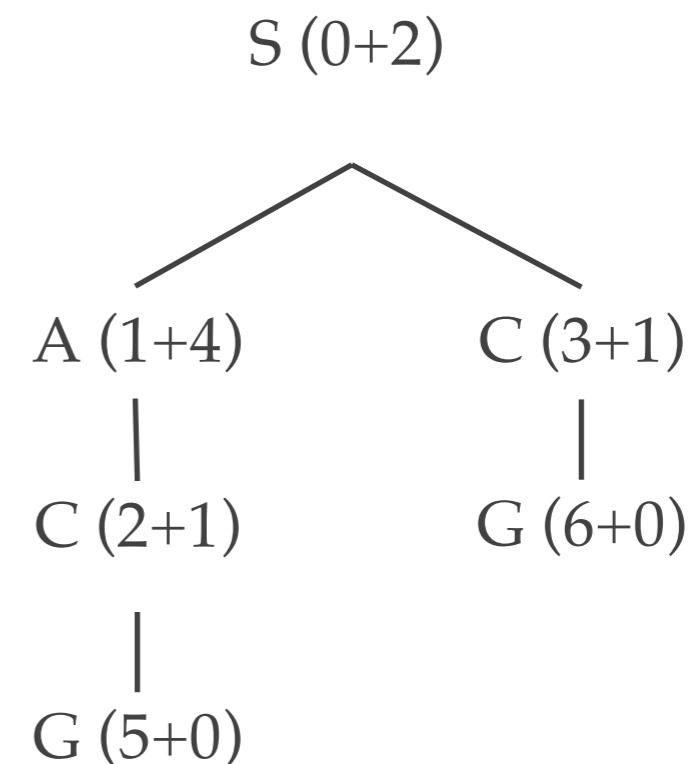


A* Graph Search Gone Wrong?

State space graph



Search tree



- ❖ Check if state already explored
- ❖ Revisit if cheaper but requires recalculating descendants

Consistency of Heuristics

- ❖ Main idea: estimated heuristic costs \leq actual costs

- ❖ Admissibility: heuristic cost \leq actual cost to goal

$h(A) \leq$ actual cost from A to G

- ❖ Consistency:

- ❖ triangular inequality

$h(A) \leq$ cost(A to C) + $h(C)$

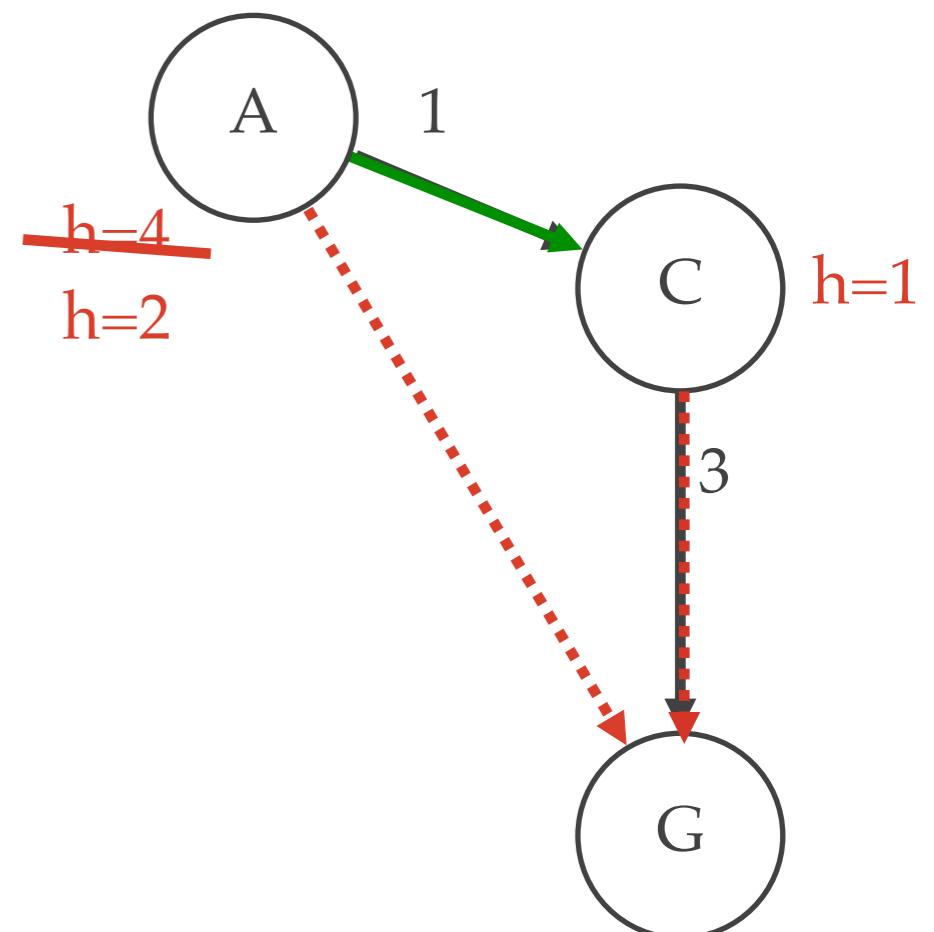
- ❖ heuristic “arc” cost \leq actual cost for each arc

$h(A) - h(C) \leq$ cost(A to C)

- ❖ Consequences of consistency:

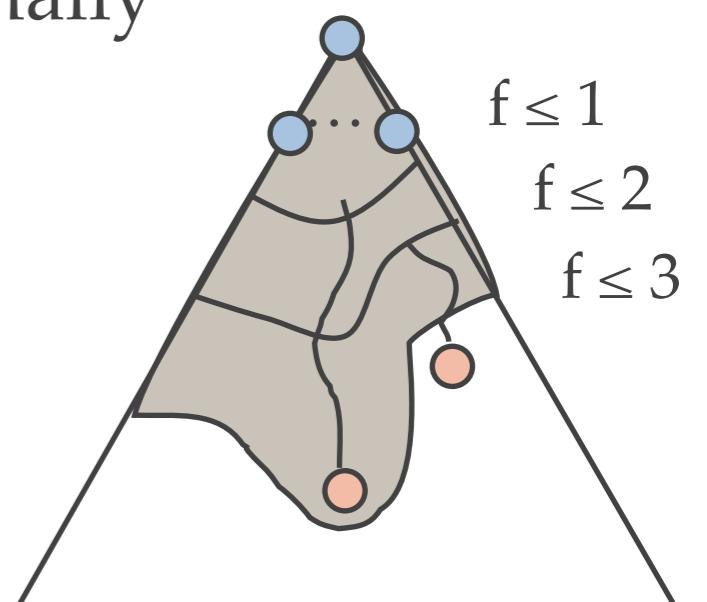
- ❖ The f value along a path never decreases

- ❖ A* graph search is optimal



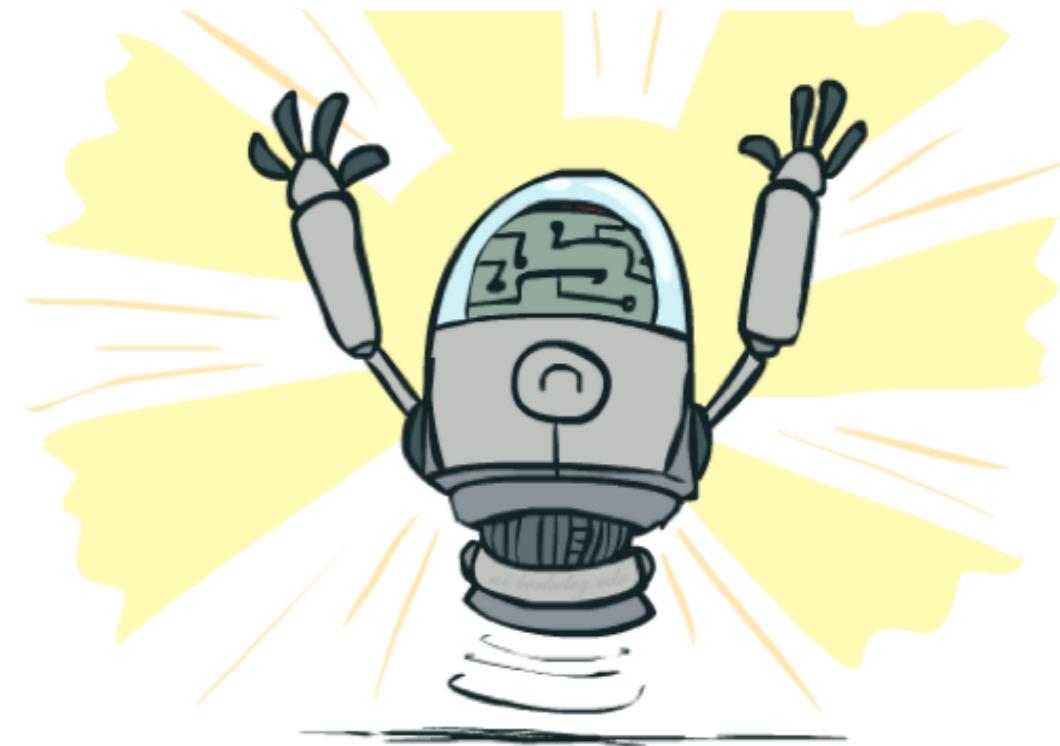
Optimality of A* Graph Search

- ❖ **Sketch:** consider what A* does with a consistent heuristic:
 - ❖ **Fact 1:** In tree search, A* expands nodes in increasing total f value (f-contours)
 - ❖ **Fact 2:** For every state n, nodes that reach n optimally are expanded before nodes that reach n suboptimally
 - ❖ **Result:** A* graph search is optimal



Optimality

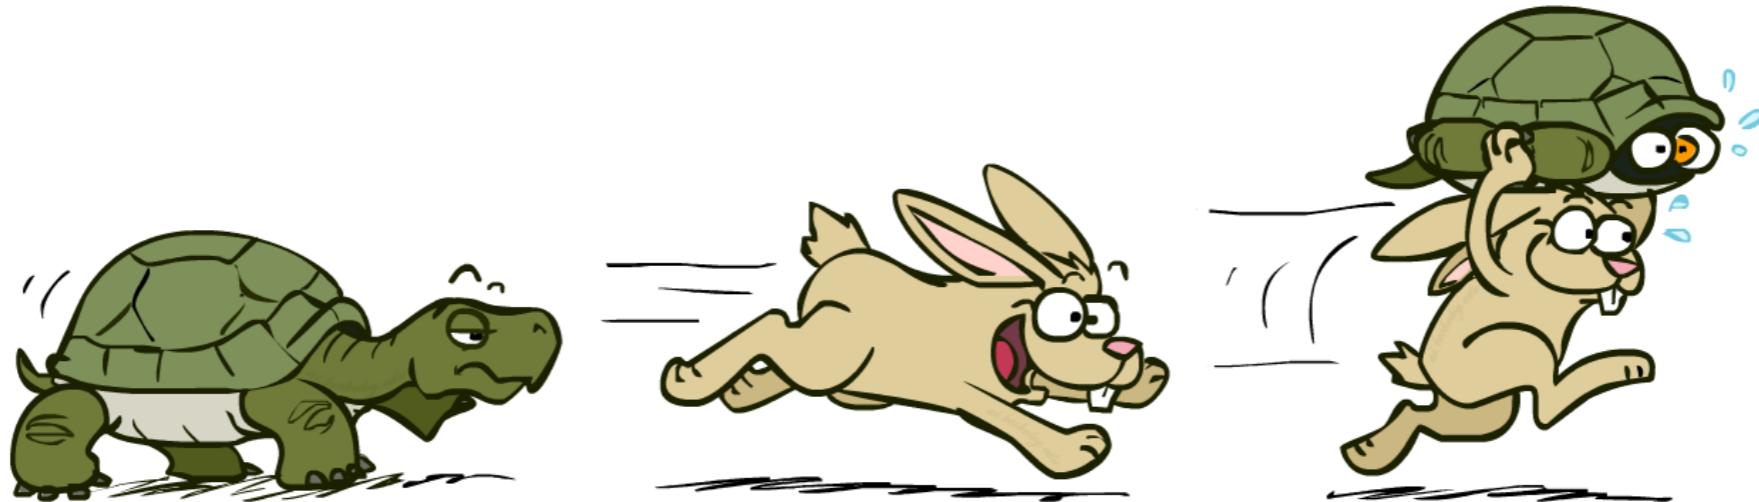
- ❖ Tree search:
 - ❖ A* is optimal if heuristic is admissible
 - ❖ UCS is a special case ($h = 0$)
- ❖ Graph search:
 - ❖ A* optimal if heuristic is consistent
 - ❖ UCS optimal ($h = 0$ is consistent)
- ❖ Consistency implies admissibility
- ❖ In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems



A*: Summary



A*: Summary



- ❖ A* uses both backward costs and (estimates of) forward costs
- ❖ A* is optimal with admissible / consistent heuristics
- ❖ Heuristic design is key: often use relaxed problems

Tree Search Pseudo-Code

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe  $\leftarrow$  INSERT(child-node, fringe)
    end
  end
```

Graph Search Pseudo-Code

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
    end
  end
```