
Announcements

- ❖ Mid-term exam: June 22, 4pm-5:40pm
 - ❖ Open book, open notes
 - ❖ No communication
- ❖ HW5 on CSP
 - ❖ Released today
 - ❖ Due June 24 at 11:59pm
- ❖ P3 on MDP and RL
 - ❖ Early release
 - ❖ Due July 3 at 11:59pm

Ve492: Introduction to Artificial Intelligence

Mid-term Review



Paul Weng

UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

What have we learned so far?

- ❖ Search and planning

- ❖ Define a state space, goal test; Find path from start to goal

- ❖ Game trees

- ❖ Define utilities; Find path from start that maximizes utility

- ❖ Decision theory and game theory

- ❖ Foundation for MEU; Basic concepts in game theory

- ❖ MDPs

- ❖ Define rewards, $utility = (\text{discounted}) \text{ sum of rewards}$
 - ❖ Find policy that maximizes utility

- ❖ Reinforcement learning

- ❖ Just like MDPs, only T and / or R are not known in advance



- ❖ Constraint satisfaction

- ❖ Find solution that satisfies constraints; Not just for finding a sequential plan



High-Level Framework

- ❖ How to build AI system?



Search

- ❖ Environment: single-agent, fully-observable state, deterministic transition, sequential, model known
- ❖ Search problem
 - ❖ States, transition model, goal test, initial state
 - ❖ Search tree
- ❖ Algorithms
 - ❖ Uninformed search
 - ❖ BFS, DFS, UCS
 - ❖ Informed search
 - ❖ Greedy search, A^*
- ❖ Properties
 - ❖ Complete, optimal
 - ❖ Space and computational complexities

Search in Games

- ❖ Environment: multi-agent, fully-observable state, deterministic or stochastic transition, turn-taking , model known
- ❖ Multi-agent search problems as games
 - ❖ States, players, transition model, terminal test/ values, initial state
 - ❖ Game tree
- ❖ Algorithm for adversarial agent (zero-sum game)
 - ❖ Minimax search algorithm
 - ❖ Alpha-beta pruning
 - ❖ Depth-limited search, iterative deepening
- ❖ Algorithm for random agent
 - ❖ Expectimax
- ❖ Algorithm for multi-agent search
 - ❖ Expectiminimax

Decision Theory and Game Theory

❖ Axiomatization of Expected Utility

- ❖ Completeness, Transitivity, Independence, Continuity
- ❖ Unicity of utility function up to positive affine transformation
- ❖ Preference elicitation

❖ Game theory

- ❖ Extensive form vs normal form
- ❖ Best response, dominant/ dominated strategies
- ❖ Nash equilibrium (pure or mixed)
- ❖ Pareto optimal, correlated equilibrium

Markov Decision Process

- ❖ Environment: single-agent, fully-observable state, stochastic transition, sequential, model known
- ❖ Model
 - ❖ States, actions, transition function, reward function
- ❖ Algorithms
 - ❖ Policy evaluation
 - ❖ Policy extraction
 - ❖ Value iteration
 - ❖ Policy iteration

Reinforcement Learning

- ❖ Environment: single-agent, fully-observable state, stochastic transition, sequential, model unknown
- ❖ MDP Model, but unknown!
 - ❖ States, actions, transition function, reward function
- ❖ Algorithms
 - ❖ Policy evaluation with TD learning
 - ❖ Policy learning with Q-learning
 - ❖ Approximate Q-learning
 - ❖ Action selection with ϵ -greedy or exploration function

Constraint Satisfaction

❖ CSP

- ❖ Set of variables, set of domains, set of constraints
- ❖ Find assignments to variables such that all constraints are satisfied

❖ Algorithms

- ❖ Backtracking search
 - ❖ Filtering, forward-checking, arc consistency, k-consistency
 - ❖ Ordering of variables and values
- ❖ Structure of constraint graph
 - ❖ Two-pass algorithm for tree-structured constraint graph
 - ❖ Cutset conditioning
- ❖ Iterative improvement

❖ Local search

Quiz: Search

- ❖ Consider a graph search problem where for every action, the cost is at least ϵ , with $\epsilon > 0$. Assume the used heuristic is consistent.
 - ❖ Greedy graph search is guaranteed to return an optimal solution.
 - ❖ A^* graph search is guaranteed to return an optimal solution.
 - ❖ A^* graph search is guaranteed to expand no more nodes than depth-first graph search.
 - ❖ A^* graph search is guaranteed to expand no more nodes than uniform-cost graph search.

Quiz: A* Heuristics

- ❖ Let H_1 and H_2 both be admissible heuristics.
 - ❖ $\max(H_1, H_2)$ is necessarily admissible
 - ❖ $\min(H_1, H_2)$ is necessarily admissible
 - ❖ $(H_1 + H_2)/2$ is necessarily admissible
 - ❖ $\max(H_1, H_2)$ is necessarily consistent

Quiz: Search under Uncertainty

- ❖ You are given a game tree for which you are the maximizer, and in the nodes in which you don't get to make a decision an action is chosen uniformly at random amongst the available options. Your objective is to maximize the probability you win \$10 or more (rather than the usual objective to maximize your expected value).
- ❖ Running expectimax will result in finding the optimal strategy to maximize the probability of winning \$10 or more.
- ❖ Running minimax, where chance nodes are considered minimizers, will result in finding the optimal strategy to maximize the probability of winning \$10 or more.
- ❖ Running expectimax in a modified game tree where every pay-off of \$10 or more is given a value of 1, and every pay-off lower than \$10 is given a value of 0 will result in finding the optimal strategy to maximize the probability of winning \$10 or more.
- ❖ Running minimax in a modified game tree where every pay-off of \$10 or more is given a value of 1, and every pay-off lower than \$10 is given a value of 0 will result in finding the optimal strategy to maximize the probability of winning \$10 or more.

Quiz: Adversarial Search

- ❖ In the context of adversarial search, α - β pruning
 - ❖ can reduce computation time by pruning portions of the game tree
 - ❖ is generally faster than minimax, but loses the guarantee of optimality
 - ❖ always returns the same value as minimax for the root of the tree
 - ❖ always returns the same value as minimax for all nodes of the tree

Game Theory: Zero-Sum Game

- ❖ Two players choose simultaneously a coin of 10 cents, 50 cents or 1 dollar, which they show to each other.
- ❖ If they chose the same coin, player I wins. Otherwise, player II wins.
- ❖ Write this game in normal form. Is there any pure NE?
- ❖ Express a system of inequalities to find a mixed NE.

Quiz: MDP

- ❖ For Markov Decisions Processes (MDPs), we have that:
 - ❖ A small discount (close to 0) encourages shortsighted, greedy behavior.
 - ❖ A large, negative living reward ($\ll 0$) encourages shortsighted, greedy behavior.
 - ❖ A negative living reward can always be expressed using a discount < 1 .
 - ❖ A discount < 1 can always be expressed as a negative living reward.

Quiz: MDP


- ❖ Value iteration can converge only if the discount factor (γ) satisfies $0 < \gamma < 1$.
- ❖ Policies found by value iteration may be superior to policies found by policy iteration.
- ❖ Policies found by policy iteration may be superior to policies found by value iteration.
- ❖ In some problems, value iteration can converge even though policy iteration may not.

Quiz: Reinforcement Learning

- ❖ Assume that the agent observes the true reward with some Gaussian noise $\mathcal{N}(0,1)$, Q-learning would still converge
- ❖ Q-learning can learn the optimal Q-function Q^* without ever executing the optimal policy.
- ❖ If an MDP has a transition model T that assigns non-zero probability for all triples $T(s, a, s')$ then Q-learning will fail.
- ❖ In Q-learning, we decide to explore every k steps, i.e., if $t = 0 [k]$ we choose a random action with a uniform distribution, otherwise we choose the greedy action. This version would still converge.

Quiz: CSP

- ❖ Assume given a CSP whose constraint graph is given below and that all the variables have the same domain.
- ❖ What is the complexity of solving it with a direct application of backtracking search?
- ❖ Which efficient strategy could you apply to solve it? What would be the complexity?



CSP Problem: Job Scheduling

❖ When can I move in?

Task	Description	Duration	Predecessor
a	Erecting walls	7	none
b	Carpentry for roof	3	a
c	Roof	1	b
d	Installations	8	a
e	Facade painting	2	c & d
f	Windows	1	c & d
g	Garden	1	c & d
h	Ceilings	3	a
i	Painting	2	f & h
j	Moving in	1	i