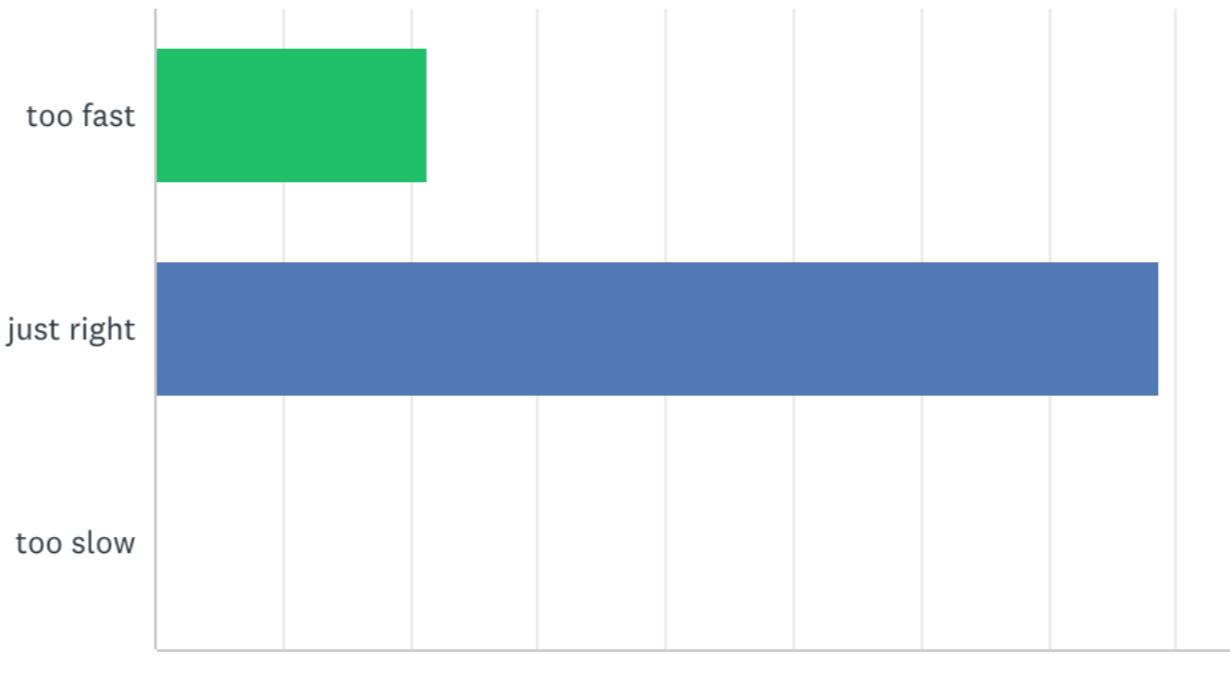


Survey Response

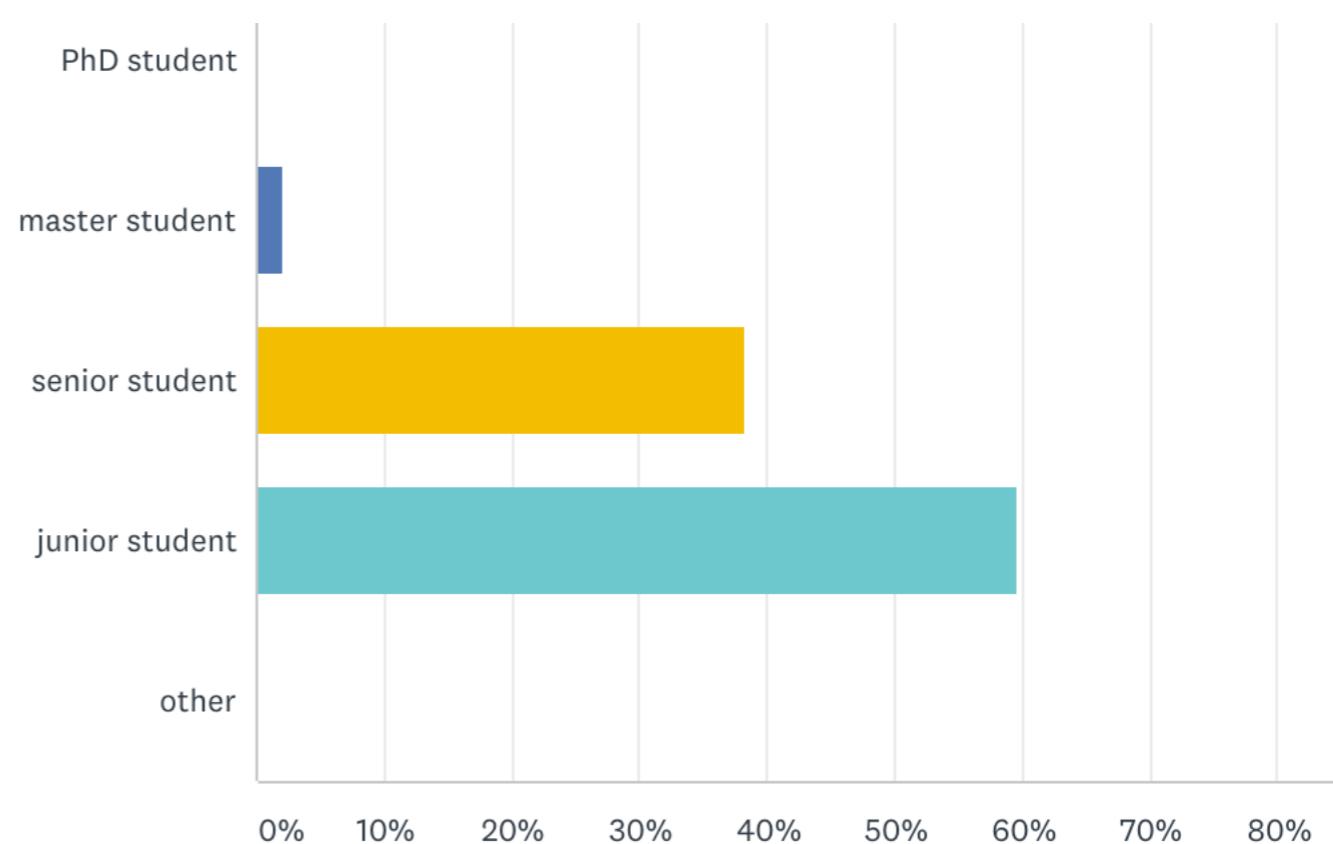
❖ Strongest features

- ❖ Content, homework, projects
- ❖ CS188



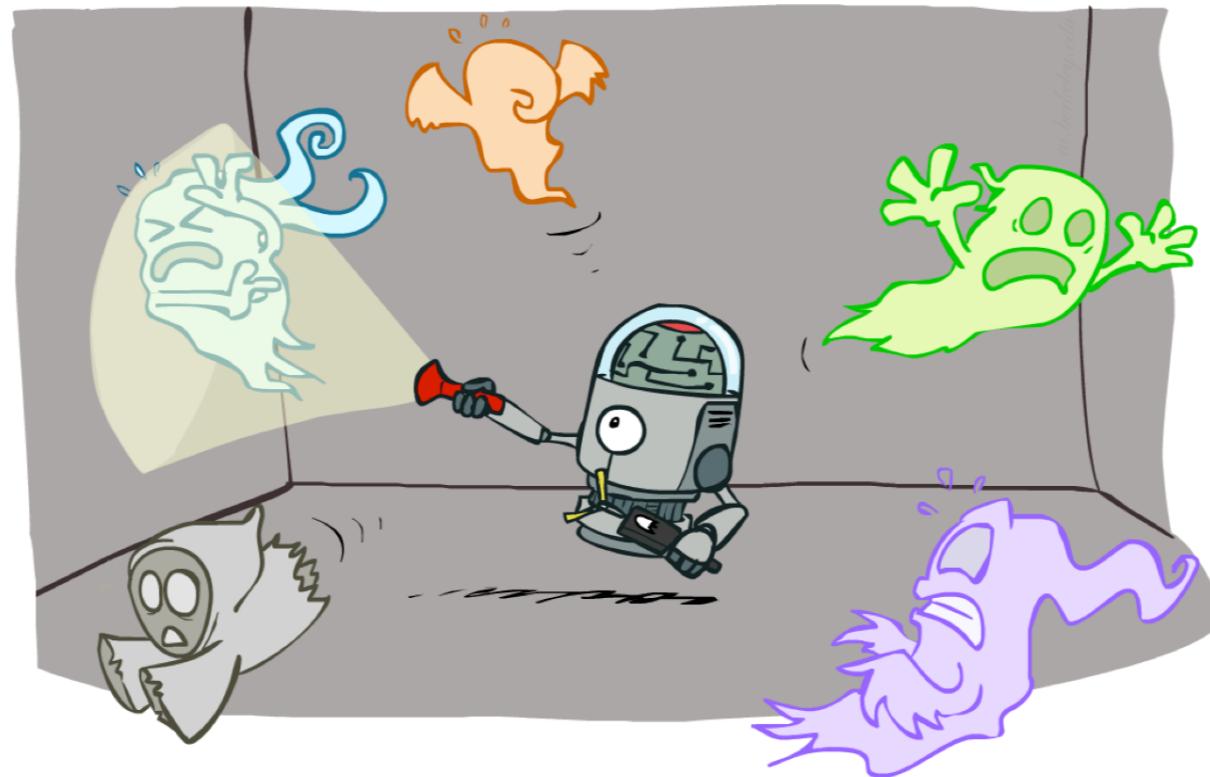
❖ Suggestions

- ❖ Exam
- ❖ Workload
- ❖ More examples, practice



Ve492: Introduction to Artificial Intelligence

Hidden Markov Models II



Paul Weng

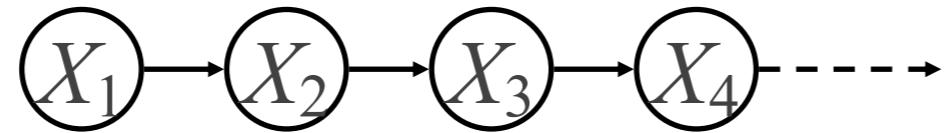
UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

Today

- ❖ Quick Review of (H)MM
- ❖ Particle Filtering
- ❖ Viterbi Algorithm for Most Likely Explanation
- ❖ Dynamic Bayesian Network

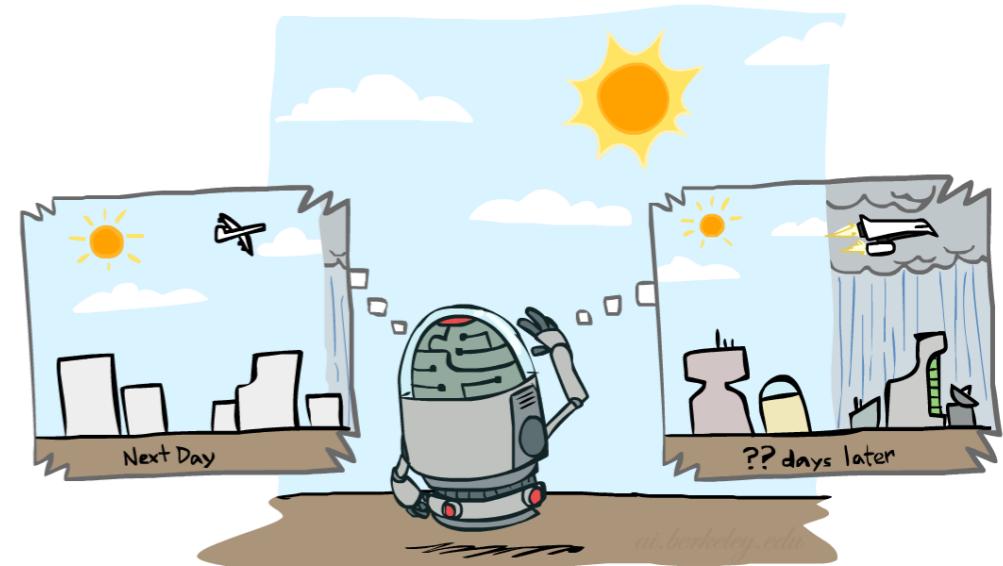
Markov Chain



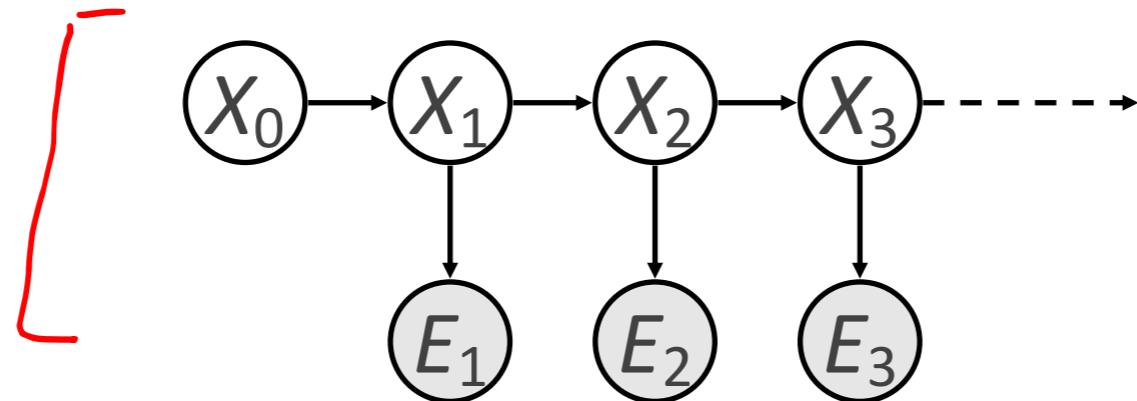
- ❖ You know transition probabilities, $P(X_t \mid X_{t-1})$, and $P(X_4)$.
- ❖ Write an equation to compute $P(X_5)$.

$$\begin{aligned} P(X_5) &= \sum_{x_4} P(x_4, X_5) \\ &= \sum_{x_4} P(X_5 \mid x_4)P(x_4) \end{aligned}$$

- ❖ More generally: $P_{t+1} = T^\top P_t$ where $P_t = P(X_t)$
- ❖ Stationary distribution: $\underline{P_\infty} = T^\top \underline{P_\infty} = (\underline{T^\top})^\infty \underline{P_1}$



HMM, Filtering, Forward Algorithm



Filtering: What is the current state, given all evidence?

$$\underbrace{P(X_{t+1} | e_{1:t+1})}_{\text{belief state}} = \alpha \underbrace{P(e_{t+1} | X_{t+1})}_{\text{Normalize}} \sum_{x_t} P(X_{t+1} | x_t) \underbrace{P(x_t | e_{1:t})}_{\text{Update}}$$

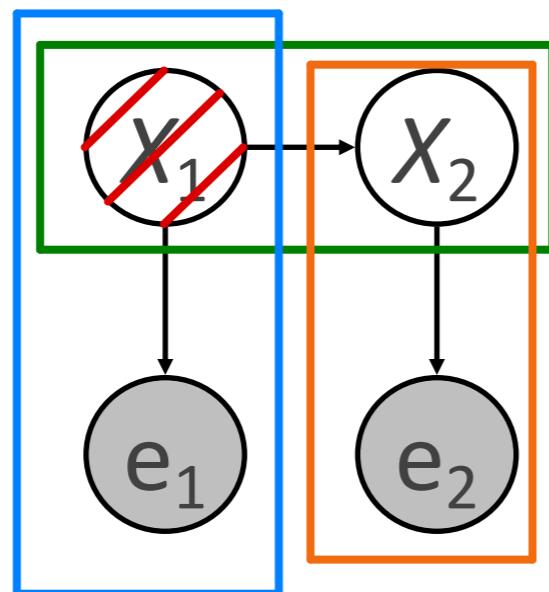
Predict

$$f_{1:t+1} = \text{FORWARD}(f_{1:t}, e_{t+1})$$

Forward Algorithm

Query: What is the current state, given all of the current and past evidence?

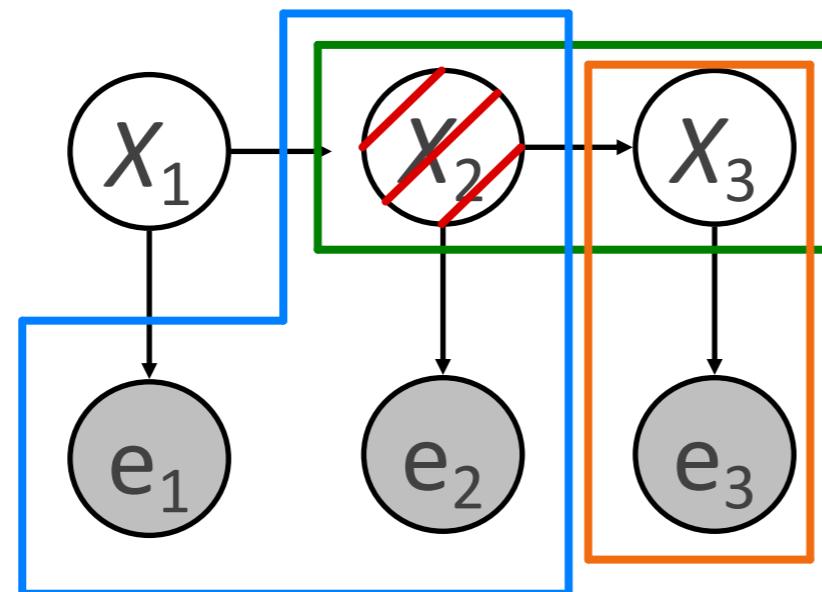
Marching **forward** through the HMM network



Forward Algorithm

Query: What is the current state, given all of the current and past evidence?

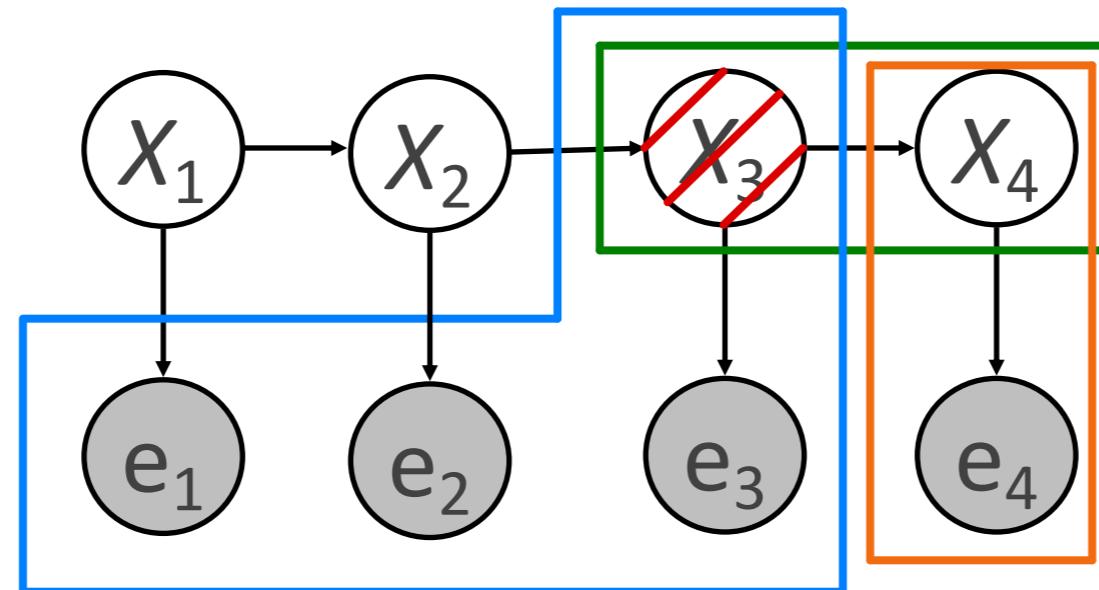
Marching **forward** through the HMM network



Forward Algorithm

Query: What is the current state, given all of the current and past evidence?

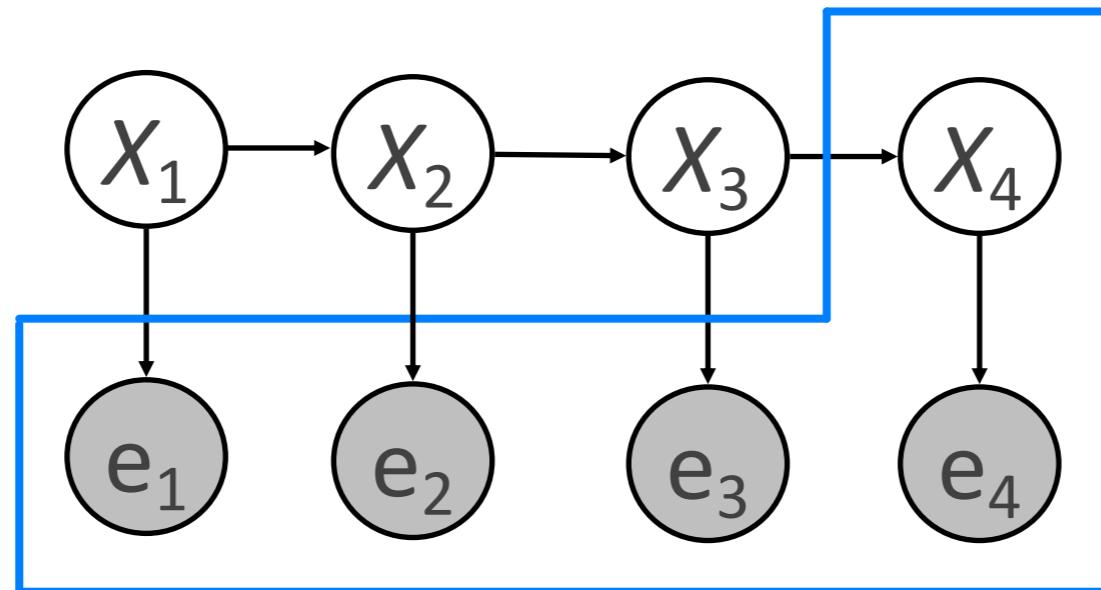
Marching **forward** through the HMM network



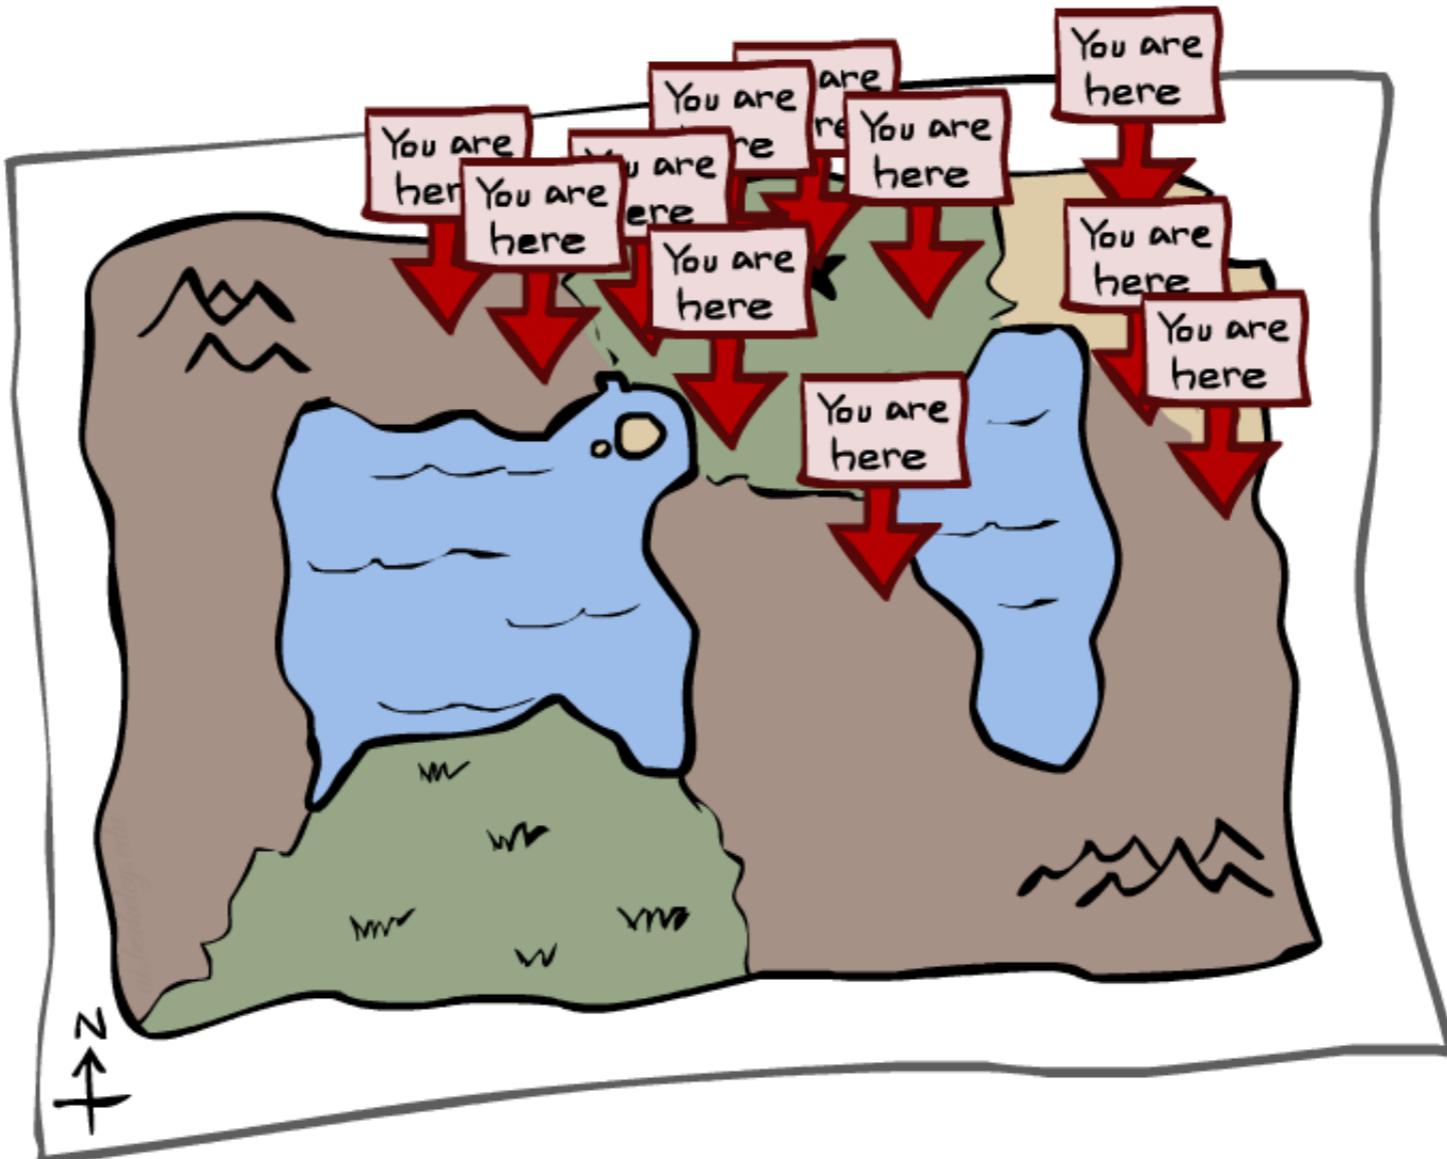
Forward Algorithm

Query: What is the current state, given all of the current and past evidence?

Marching **forward** through the HMM network



Particle Filtering



Quiz: Algorithms for Filtering

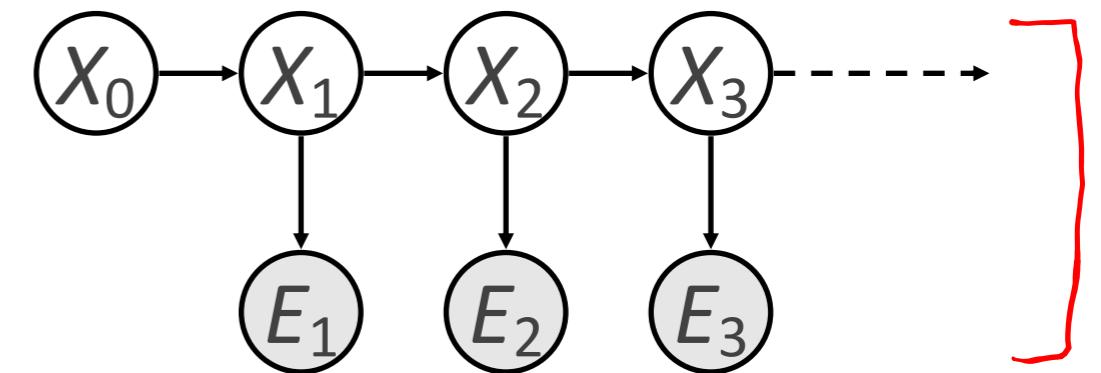
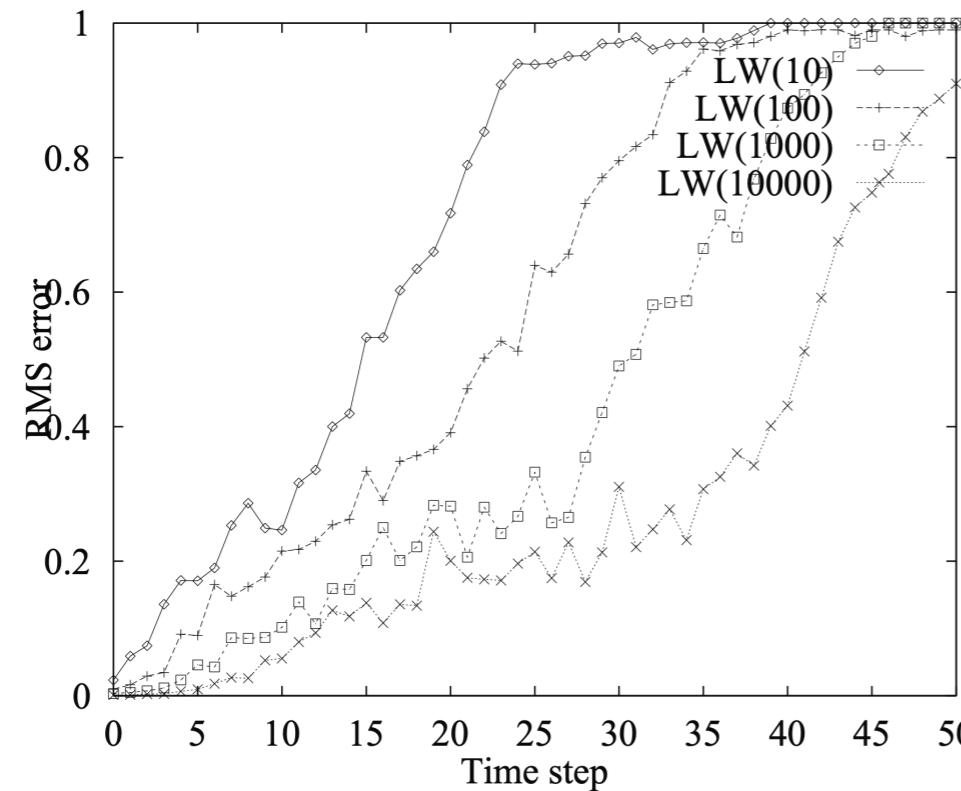
- ❖ Which algorithms can be applied to the filtering task?

- ❖ Variable elimination
- ❖ Prior sampling
- ❖ Rejection sampling
- ❖ Likelihood weighting
- ❖ Gibbs sampling

Inference in HMM

Limitations of Current Algorithms

- ❖ Exact inference infeasible if state space is large or continuous
 - ❖ e.g., when $|X|$ is more than 10^6 or so (tracking 3 ghosts in a 10×20 world)
 - ❖ e.g., robot localization



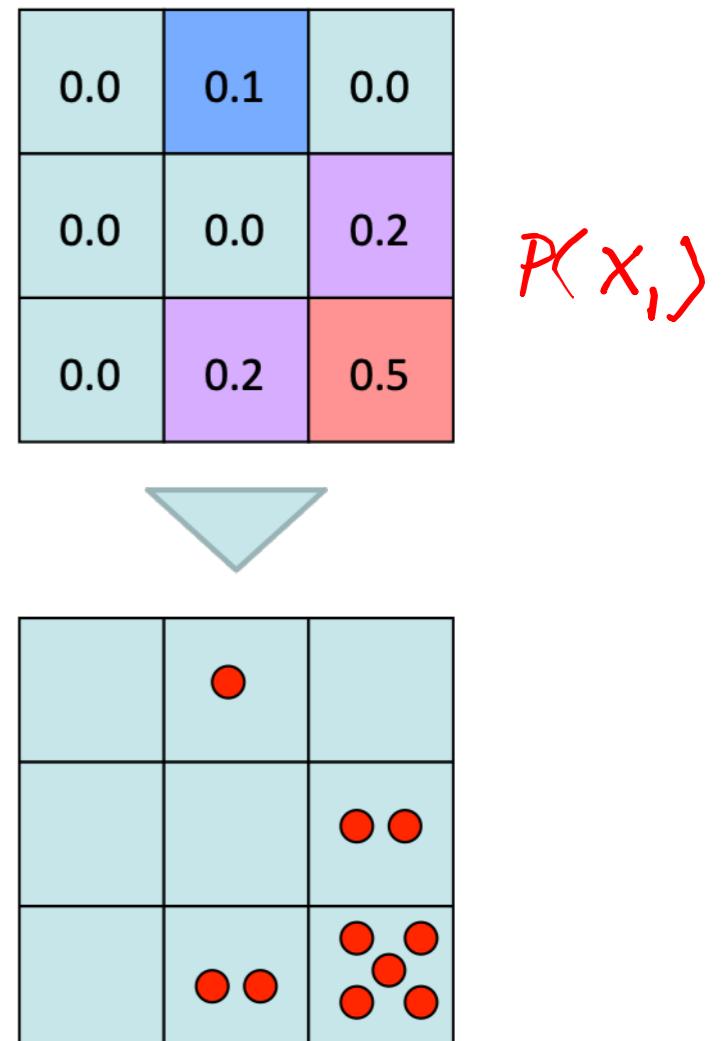
- ❖ Previous sampling methods do not exploit the temporal structure of an HMM

Particle Filtering

- ❖ PF = approximate inference for filtering task

- ❖ Principle of Particle Filtering

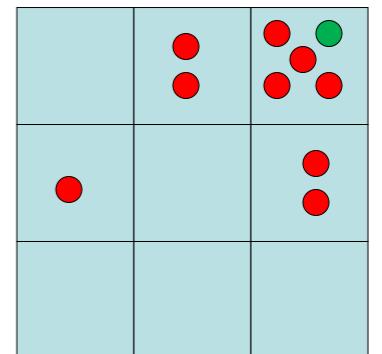
- ❖ Track samples of X_t , not all values $P(X_{t \mid e_{:, t}})$
- ❖ Samples are called particles
- ❖ Time per step is linear in the number of samples
- ❖ But: number needed may be large
- ❖ In memory: list of particles, not states



Representation: Particles

- ❖ Our representation of $P(X)$ is now a list of N particles (samples)

- ❖ Generally, $N \ll |X|$
- ❖ Storing map from X to counts would defeat the point



- ❖ $P(x)$ approximated by number of particles with value x

- ❖ So, many x may have $P(x) = 0$!
- ❖ More particles, more accuracy

Particles:

(3,3) —

(2,3)

(3,3) —

(3,2)

(3,3) —

(3,2)

(1,2)

(3,3) —

(3,3) —

(2,3)

- ❖ For now, all particles have a weight of 1

- ❖ In PF, a set of samples approximates $\underline{f_{1:t}(X_t)} = \underline{P(X_t|e_{1:t})}$

$$\uparrow$$

$$\hat{f}_{1:t}(3,3) = \frac{5}{\# \text{ particles}}$$

Particle Filtering: Elapse Time ✓

- ❖ This first step captures the passage of time
- ❖ Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x)) \quad \checkmark$$

- ❖ This is like prior sampling – samples' frequencies reflect the transition probabilities
- ❖ Here, most samples move clockwise, but some move in another direction or stay in place

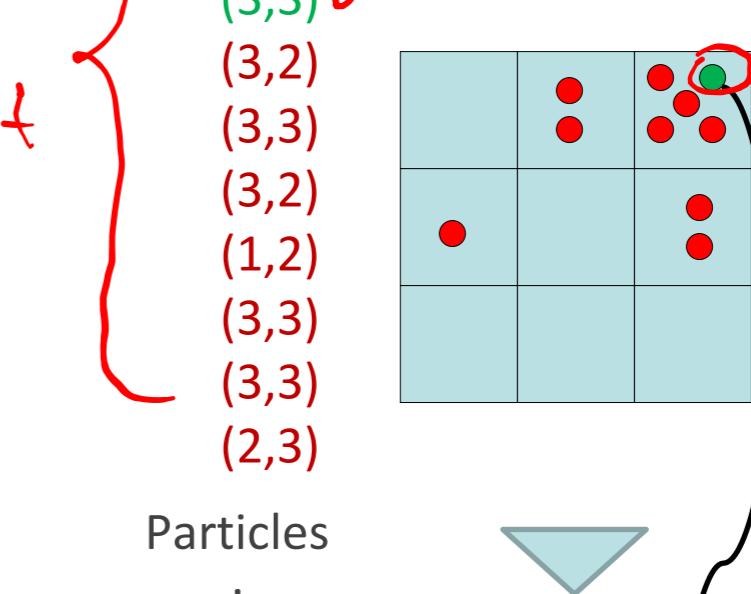
- ❖ New particles approximates:

$$\sum_{x_t} P(X_{t+1}|x_t) f_{1:t}(x_t) = \sum_{x_t} P(X_{t+1}, x_t | e_{1:t})_{t+1}$$

- ❖ Consistency: If enough samples, close to exact values before and after

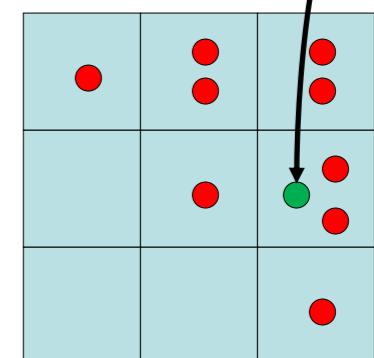
Particles:

(3,3)
(2,3)
(3,3) ✓
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)



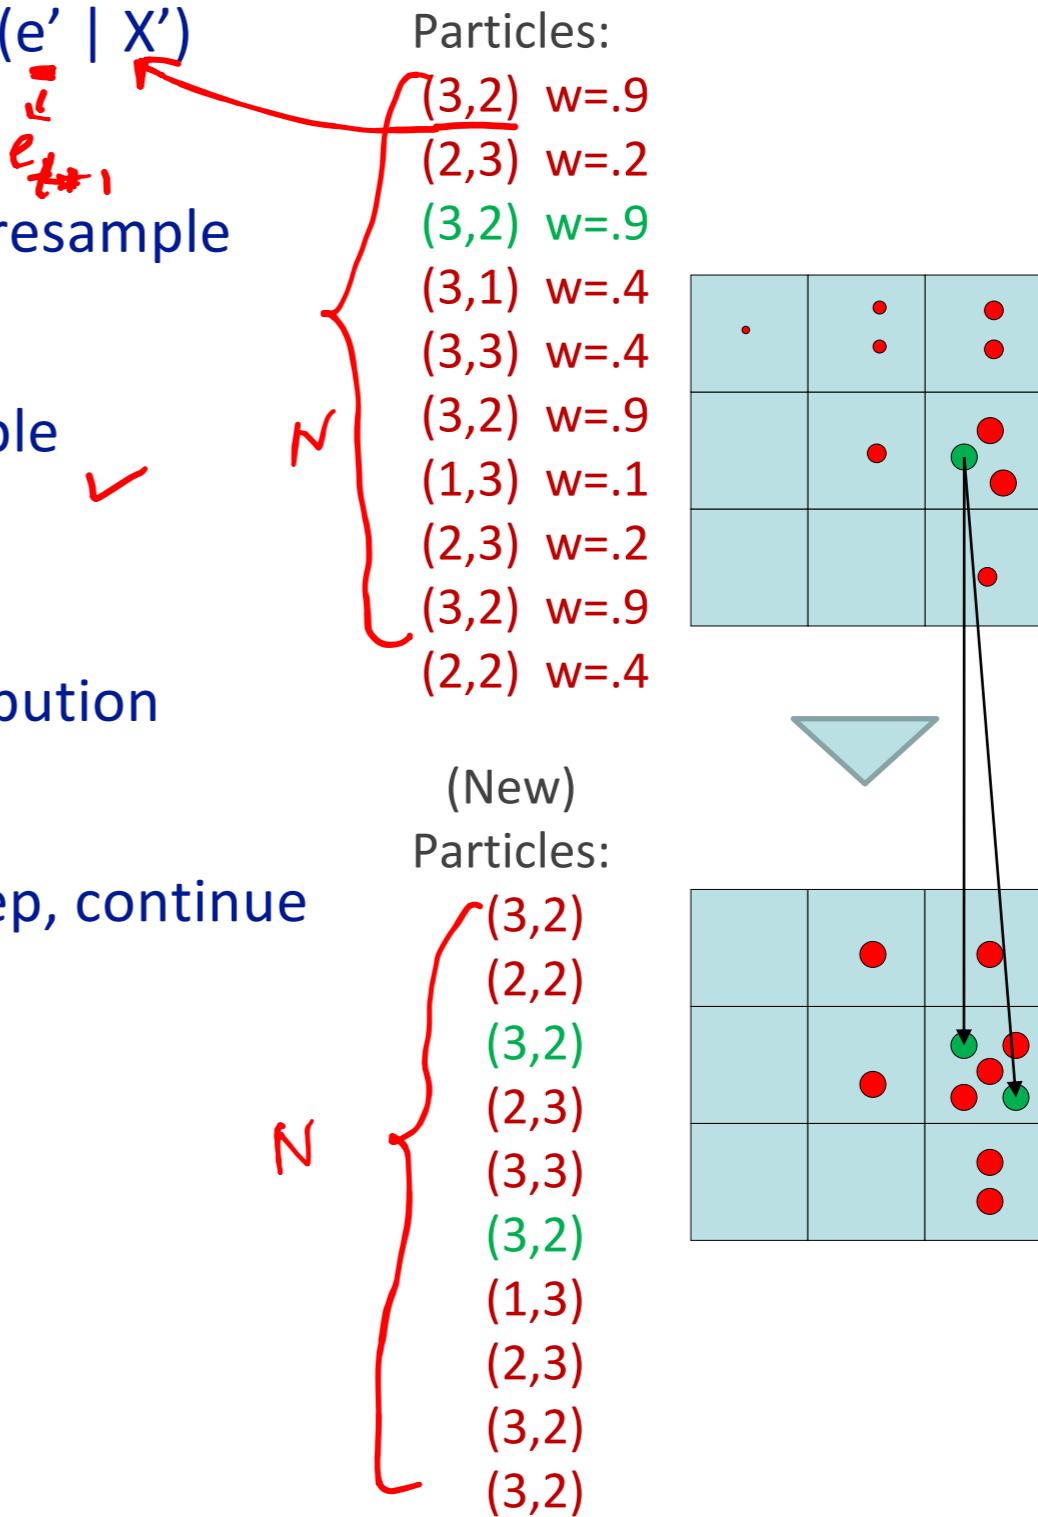
Particles

:
(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



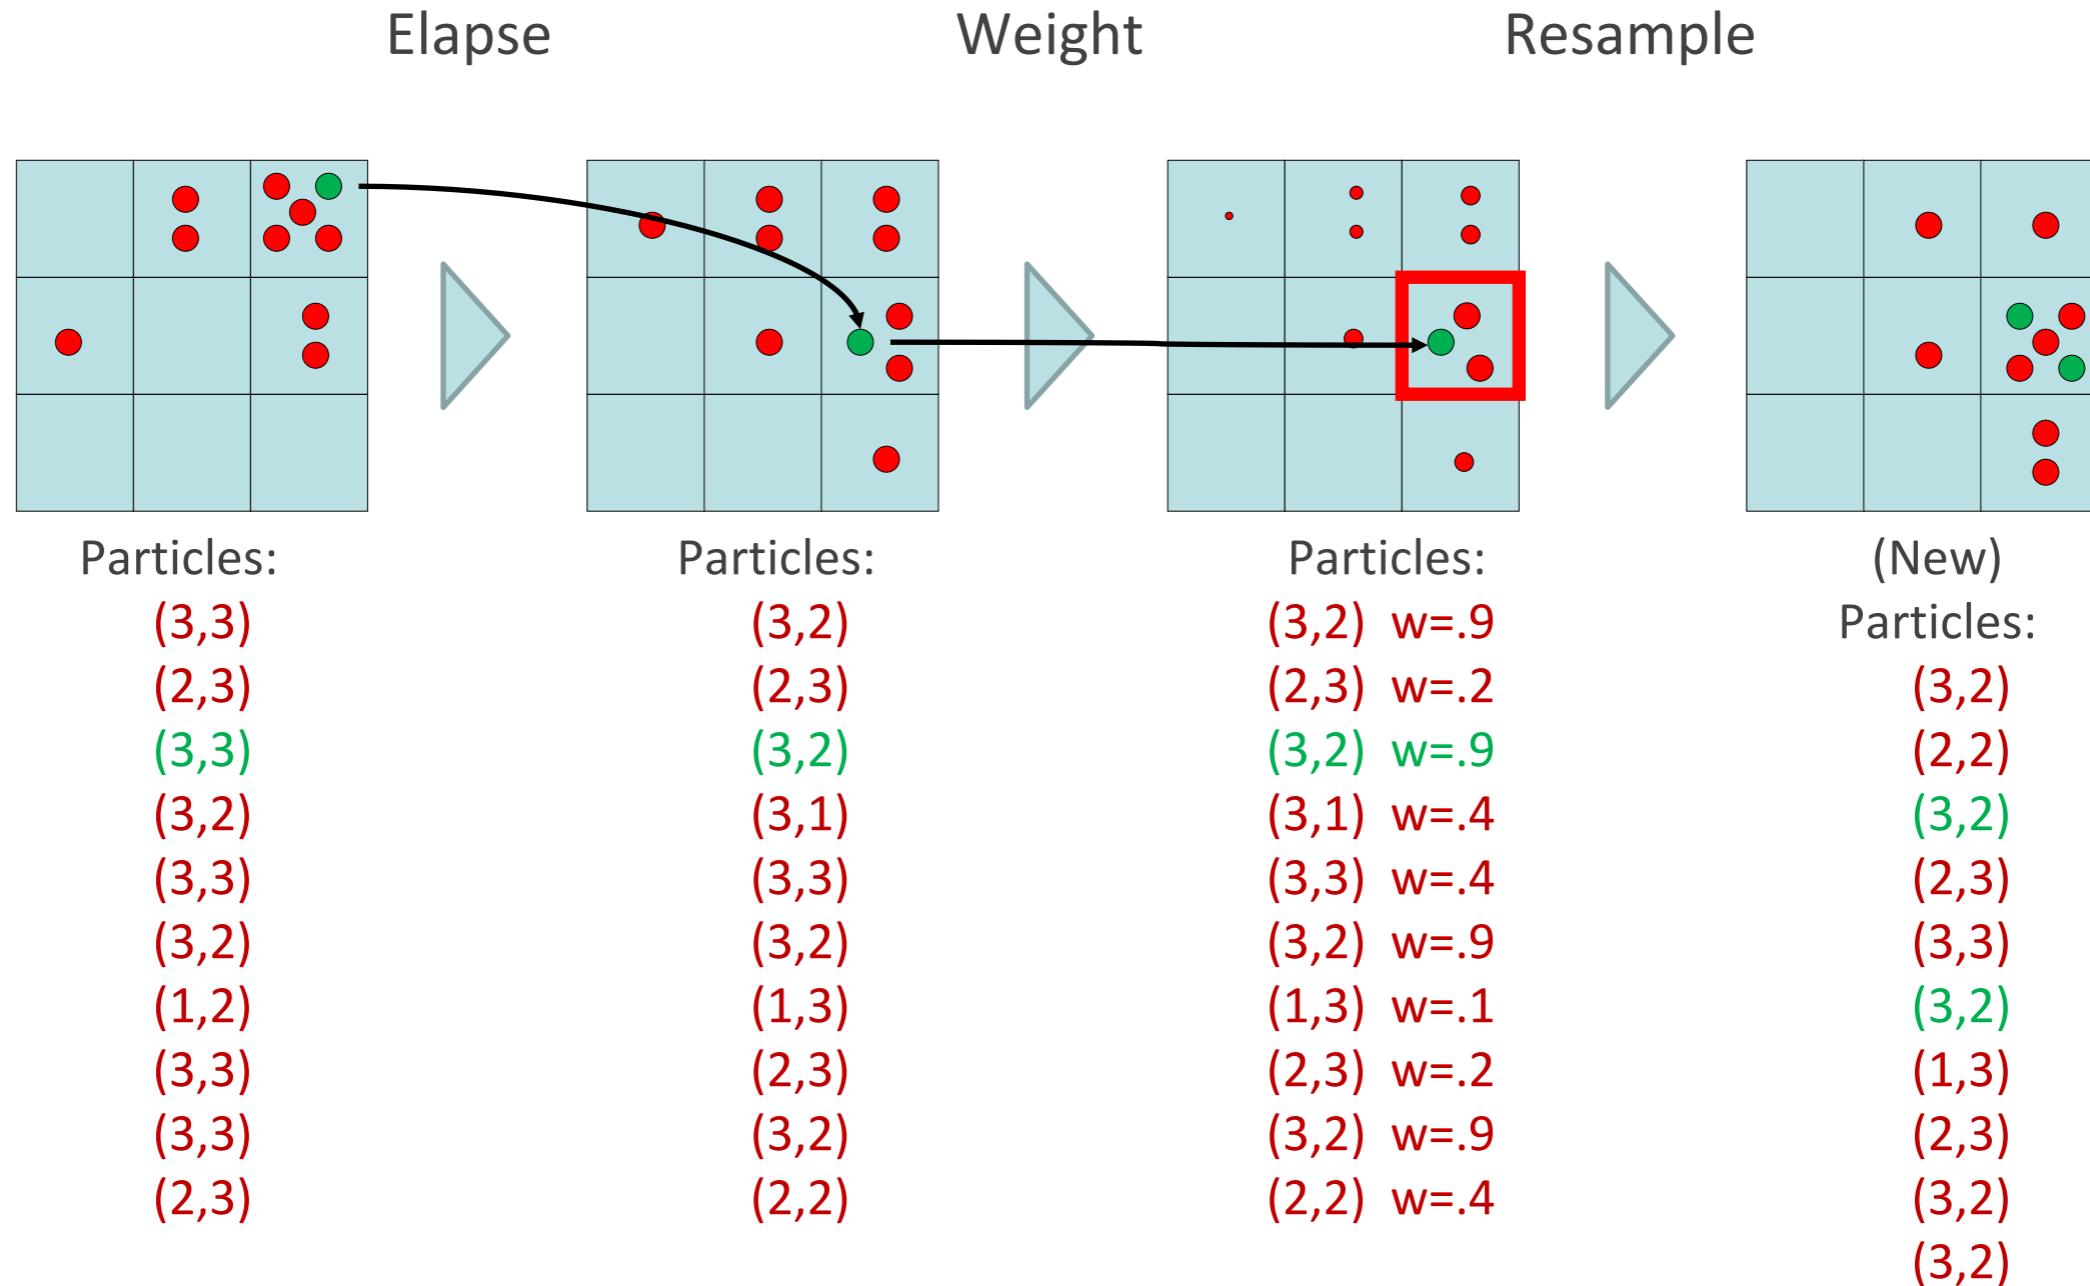
Particle Filtering: Resample

- ❖ Particles receive weights corresponding to $P(e' | X')$
 e_{t+1}
- ❖ Rather than tracking weighted samples, we resample
- ❖ N times, we choose from our weighted sample distribution (i.e. draw with replacement) ✓
- ❖ This is equivalent to renormalizing the distribution
- ❖ Now the update is complete for this time step, continue with the next one

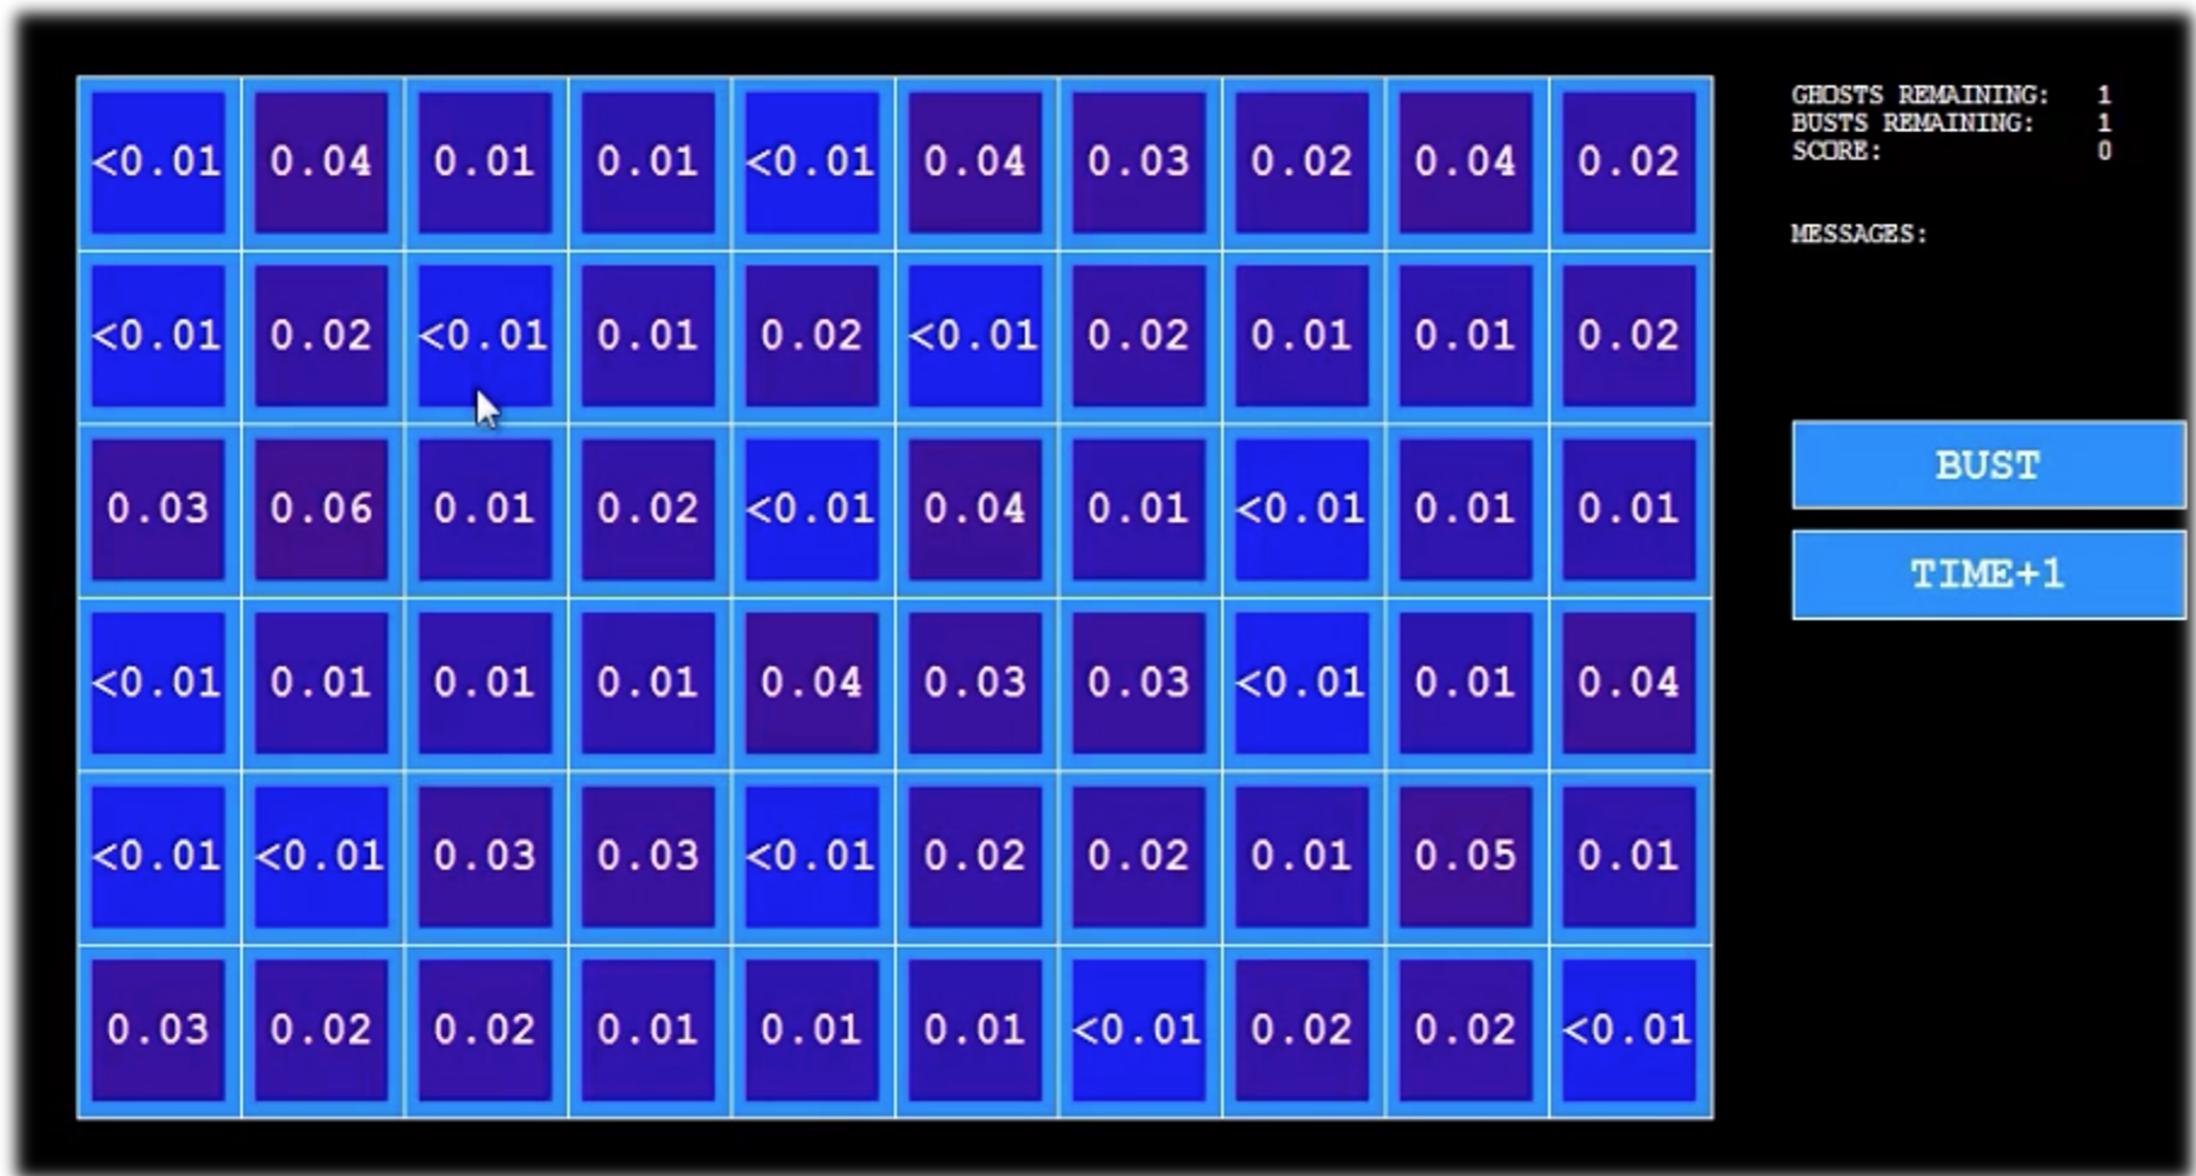


Summary: Particle Filtering

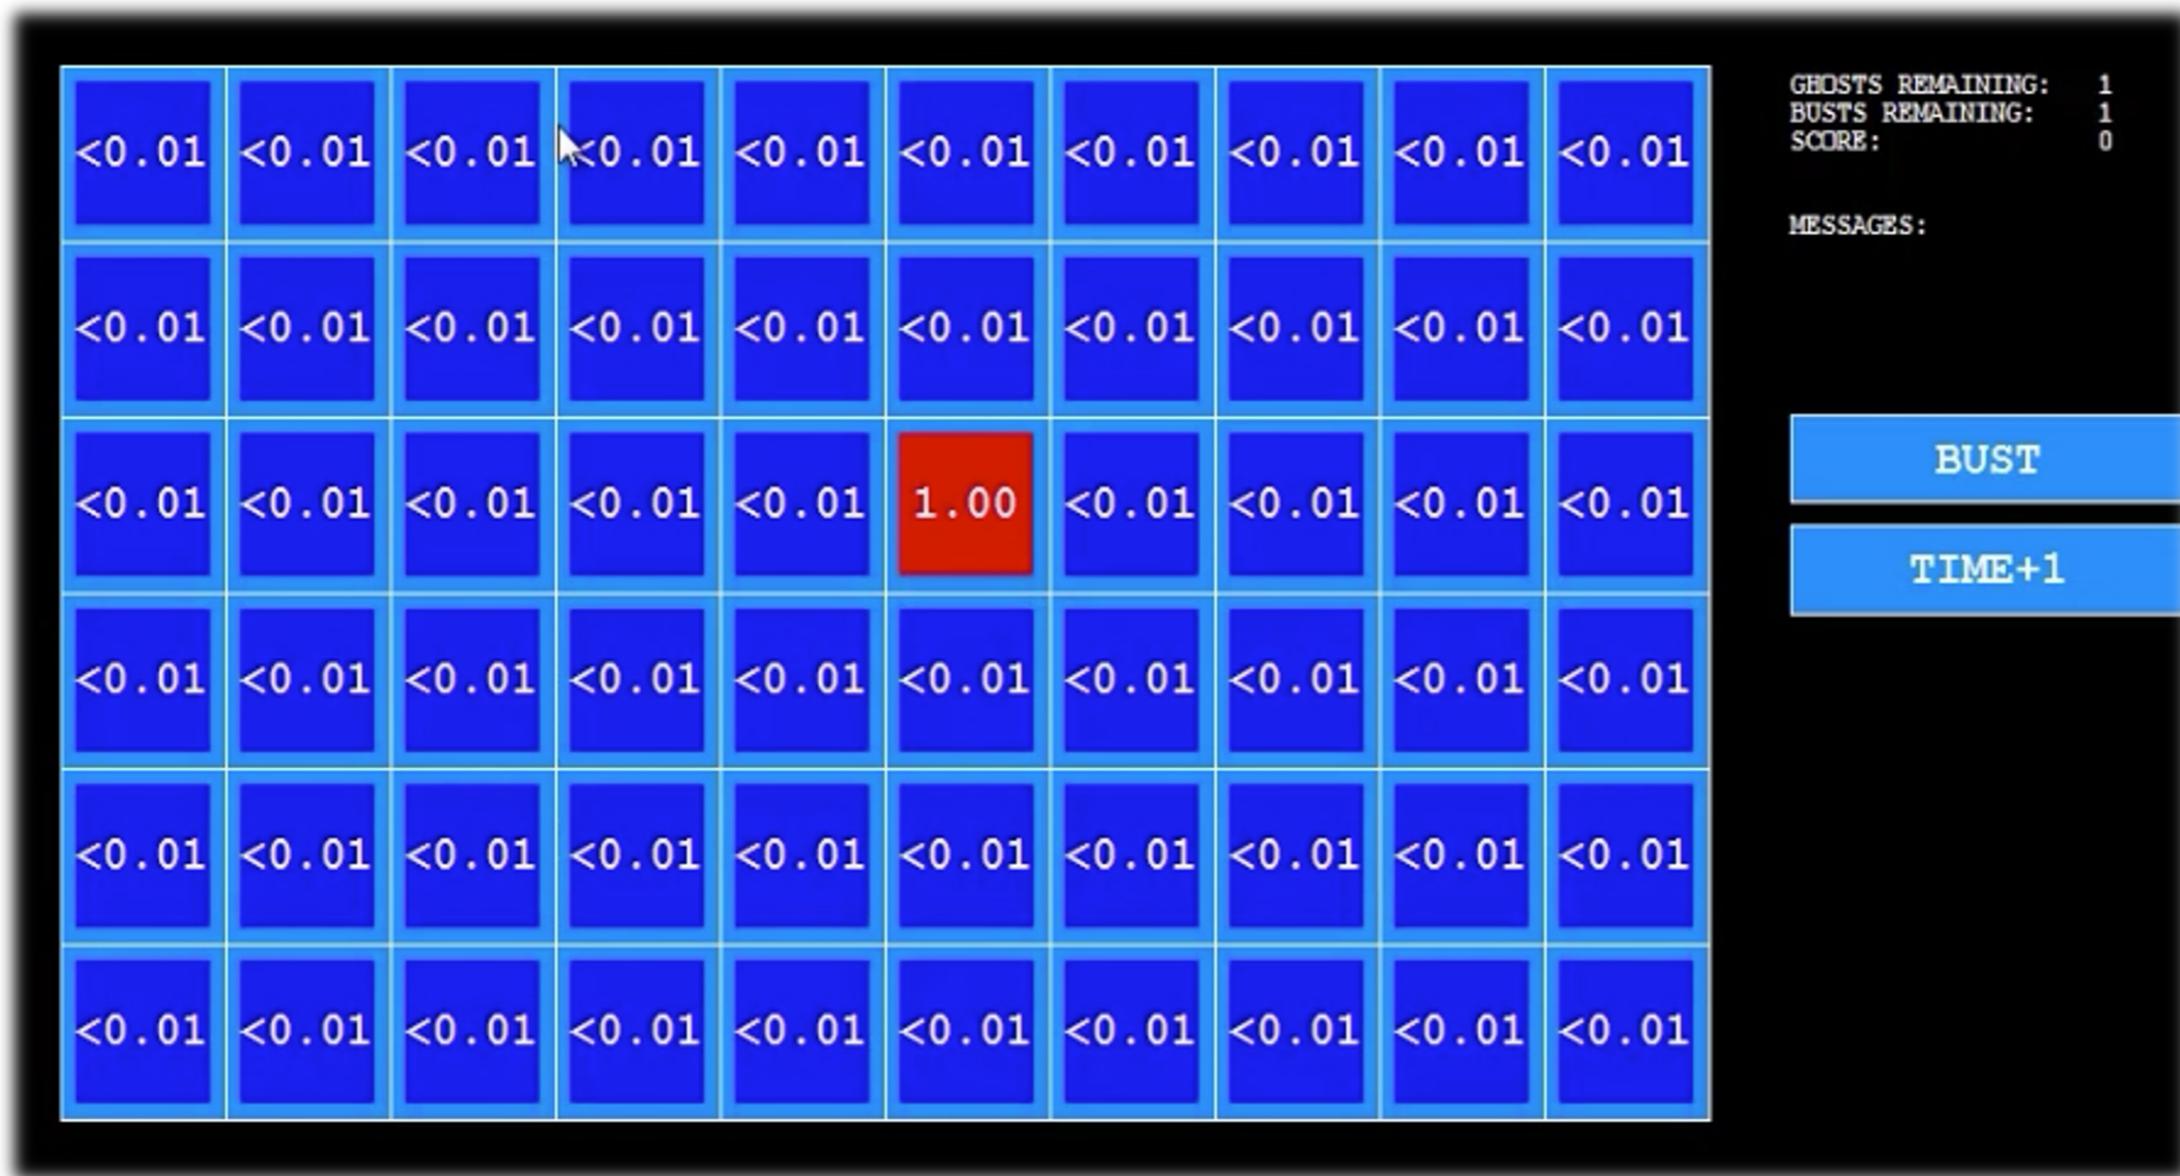
- ❖ Particles: track samples of states rather than an explicit distribution



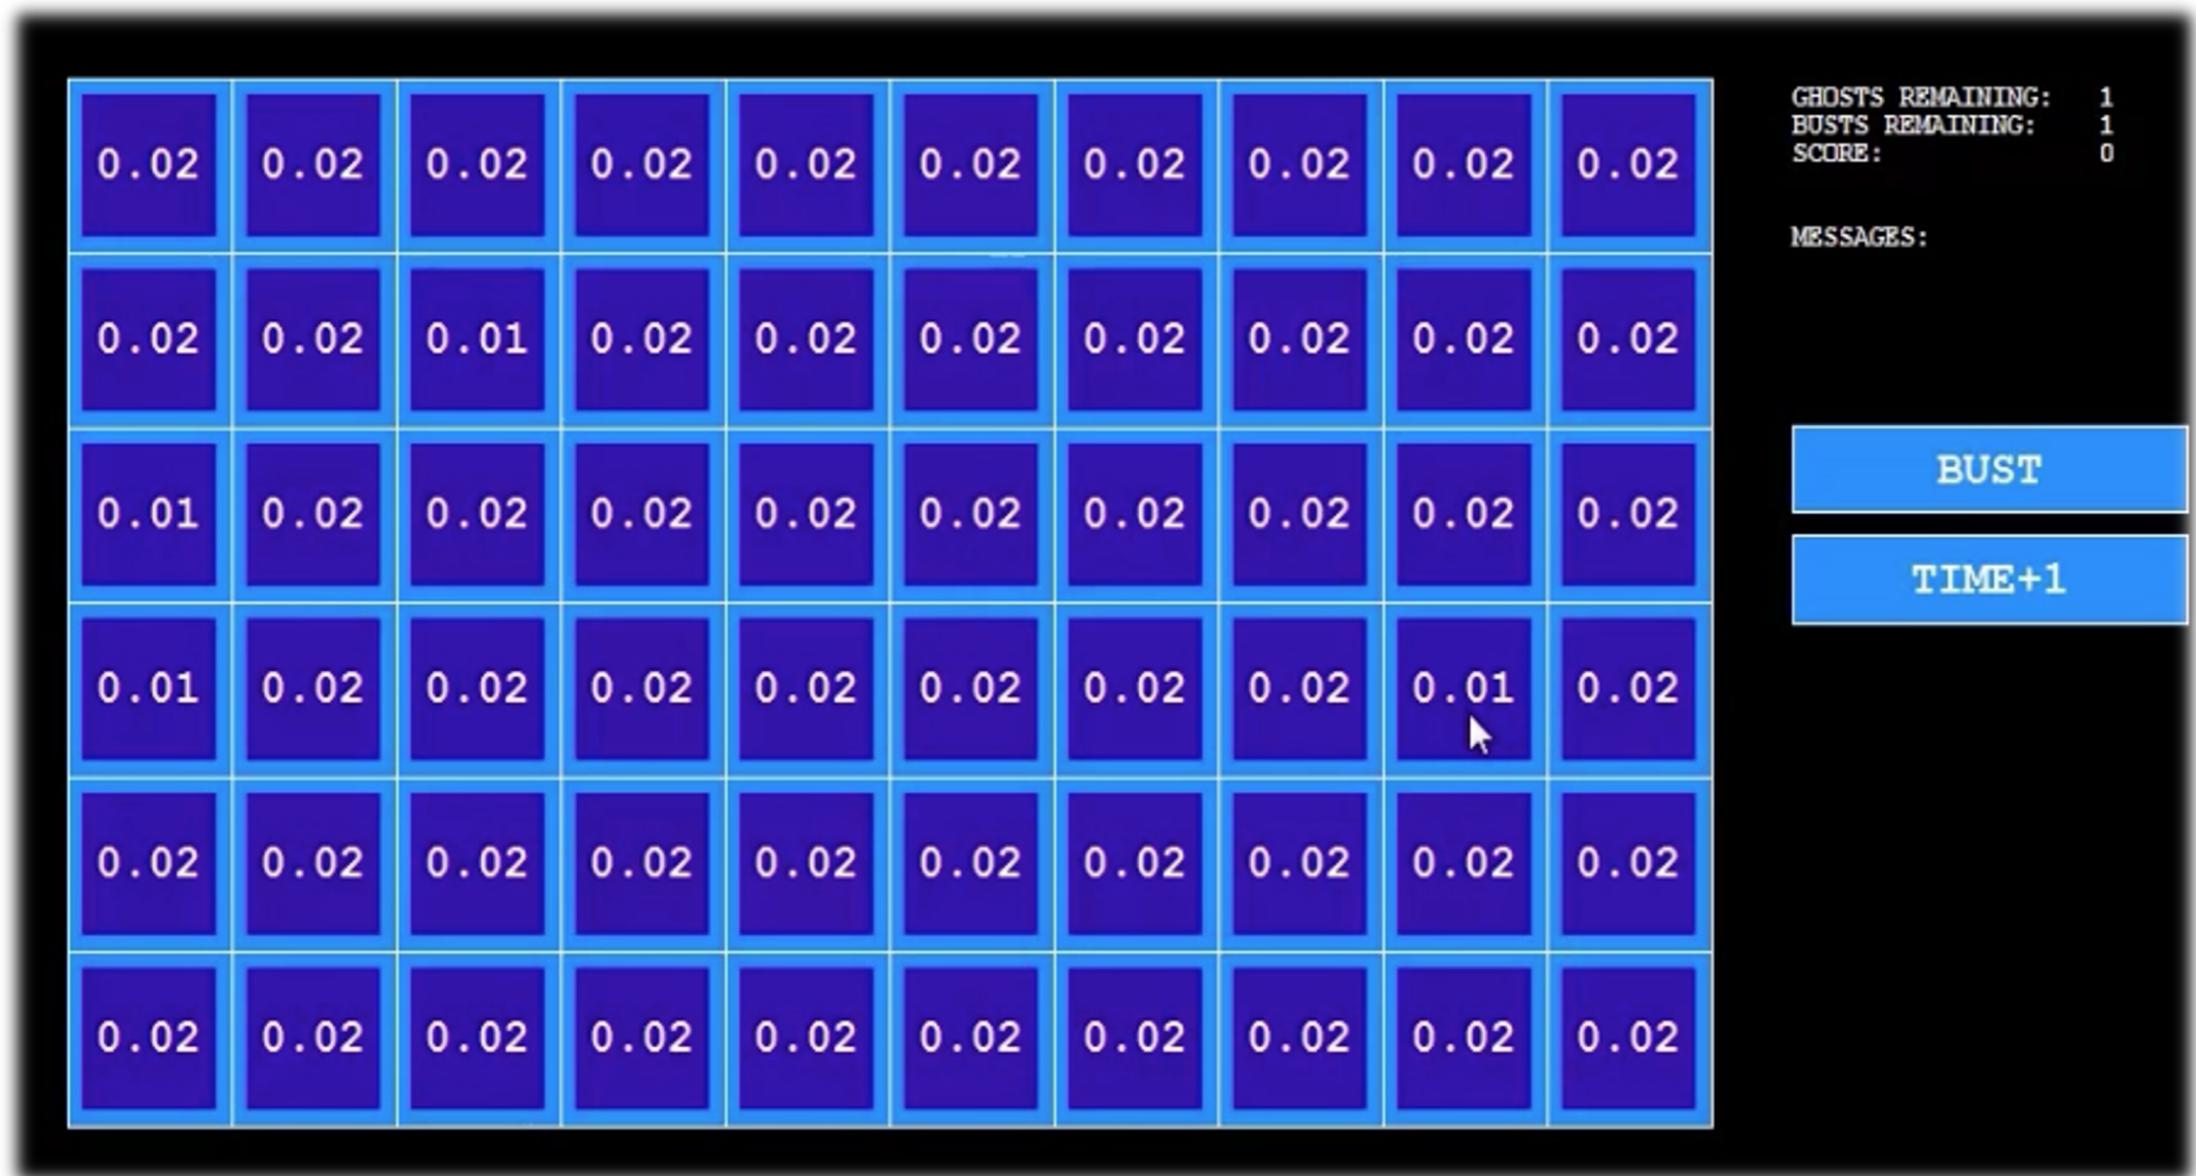
Video of Demo – Moderate Number of Particles



Video of Demo – One Particle



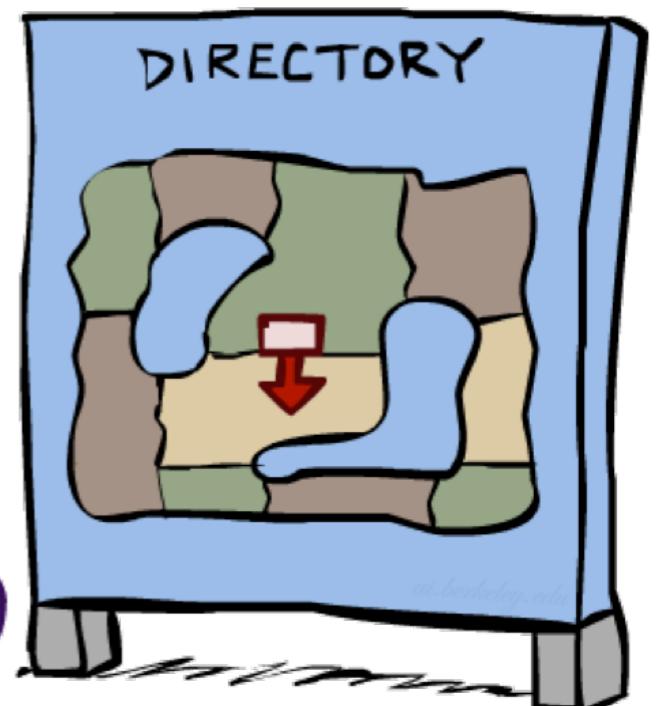
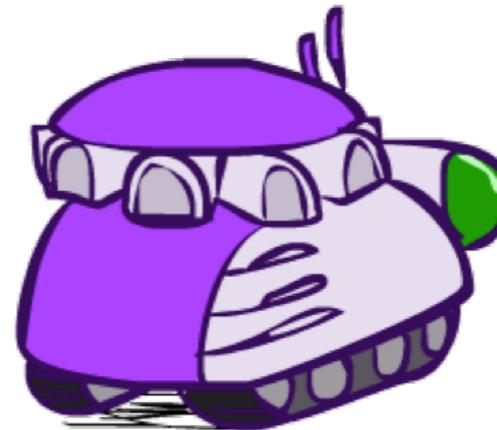
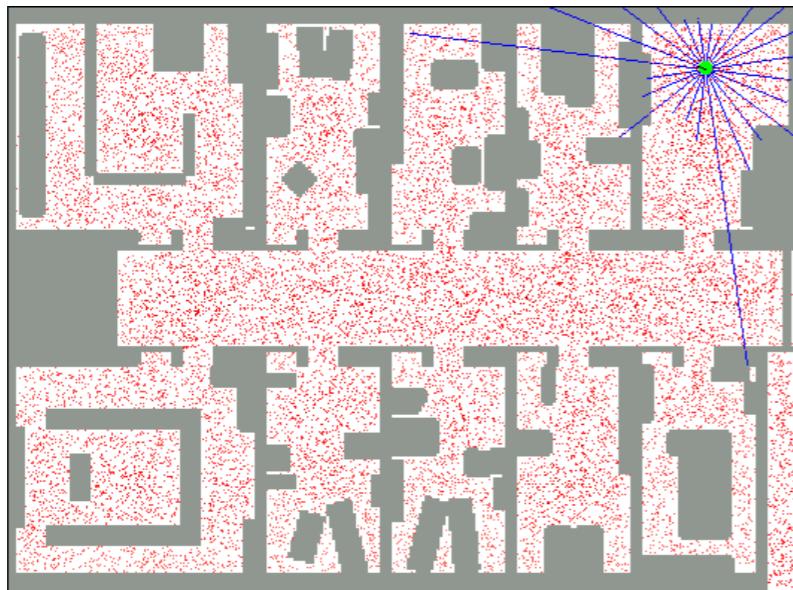
Video of Demo – Huge Number of Particles



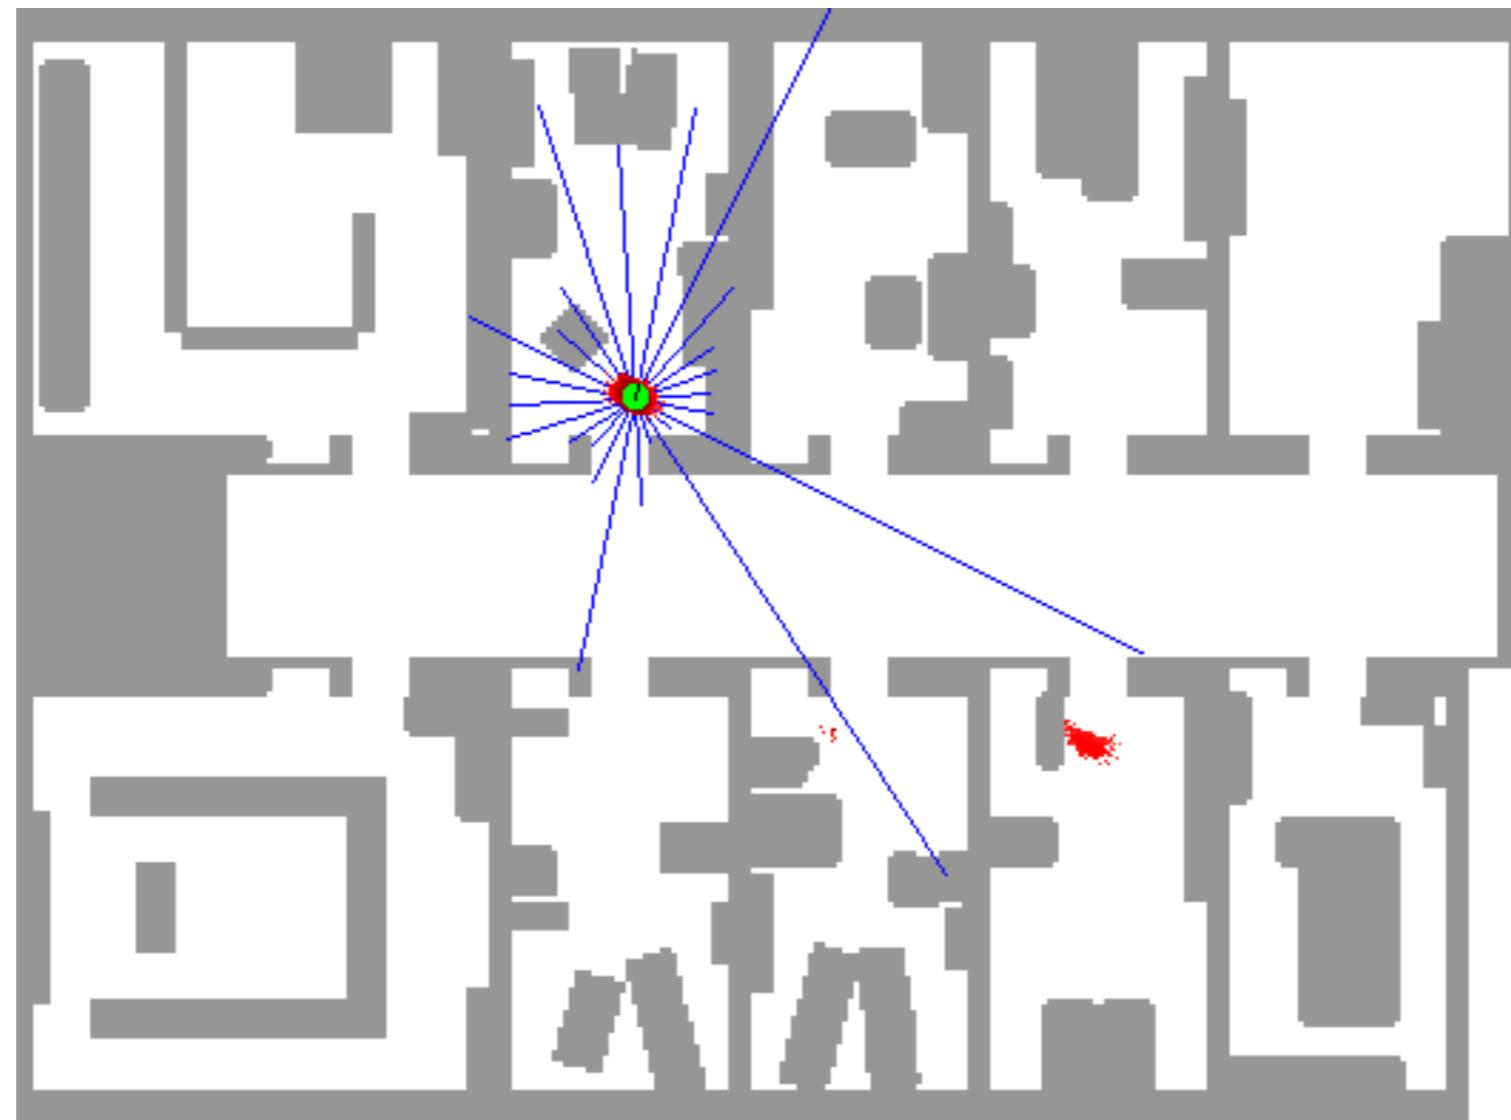
Robot Localization

In robot localization:

- ❖ We know the map, but not the robot's position
- ❖ Observations may be vectors of range finder readings
- ❖ State space and readings are typically continuous
- ❖ Particle filtering is a main technique



Particle Filter Localization (Laser)



[Dieter Fox, et al.]

Most Likely Explanation



Inference tasks

✓ *Filtering*: $P(X_t | e_{1:t})$

- ❖ **belief state**—input to the decision process of a rational agent
- ❖ **Prediction**: $P(X_{t+k} | e_{1:t})$ for $k > 0$
 - ❖ evaluation of possible action sequences; like filtering without the evidence
- ❖ **Smoothing**: $P(X_k | e_{1:t})$ for $0 \leq k < t$
 - ❖ better estimate of past states, essential for learning

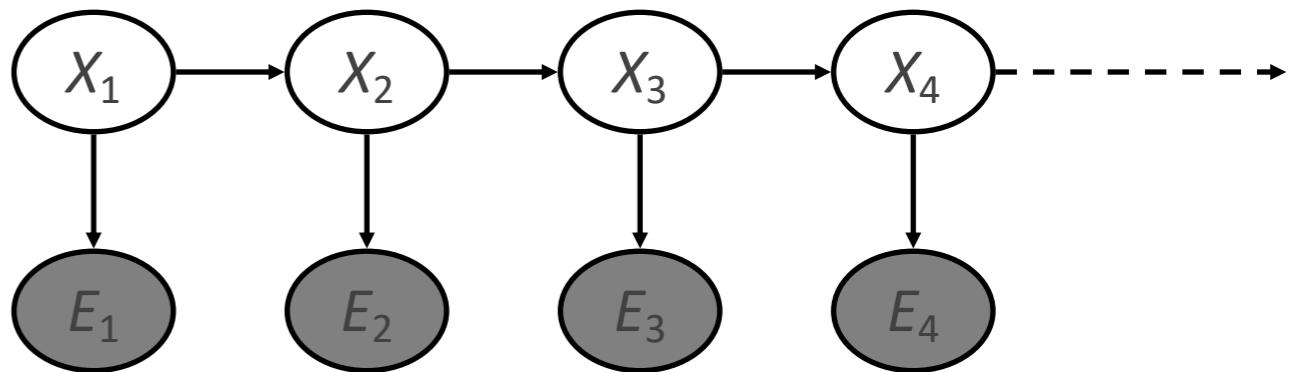
❖ **Most likely explanation**: $\arg \max_{x_{1:t}} P(x_{1:t} | e_{1:t})$

- ❖ speech recognition, decoding with a noisy channel

HMMs: MLE Queries

- ❖ HMMs defined by

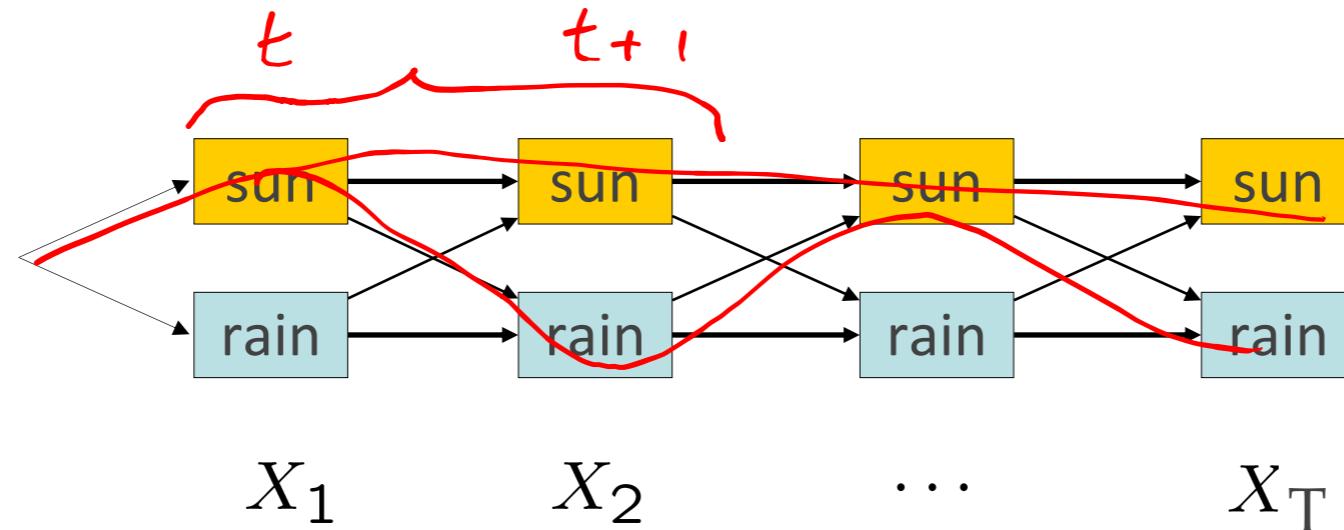
- ❖ States X
- ❖ Observations E
- ❖ Initial distribution: $P(X_1)$
- ❖ Transitions: $P(X_t|X_{t-1})$
- ❖ Emissions: $P(E|X)$



- ❖ New query: most likely explanation: $\text{argmax}_{x_{1:t}} P(x_{1:t}|e_{1:t})$
- ❖ New method: the Viterbi algorithm

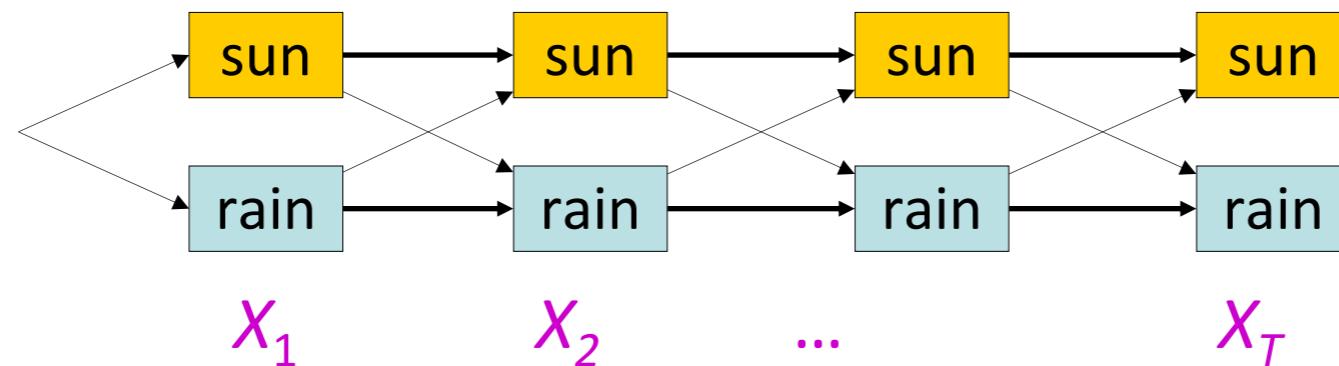
State Trellis

- ❖ State trellis: graph of states and transitions over time



- ❖ Each arc represents some transition $x_{t-1} \rightarrow x_t$
- ❖ Each arc has weight $P(x_t|x_{t-1})P(\underline{e_t|x_t})$
- ❖ Each path is a sequence of states
- ❖ Product of weights on a path = sequence's probability along with the evidence
- ❖ Forward algorithm computes sums of paths, Viterbi computes best paths

Forward / Viterbi algorithms



Forward Algorithm (sum)

For each state at time t , keep track of the **total probability of all paths** to it

$$\begin{aligned} \mathbf{f}_{1:t+1}^{\leftarrow} &= \text{FORWARD}(\mathbf{f}_{1:t}^{\leftarrow}, e_{t+1}) \\ &= \alpha P(e_{t+1}|X_{t+1}) \sum_{x_t} P(X_{t+1}|x_t) \mathbf{f}_{1:t}^{\leftarrow} \end{aligned}$$

update

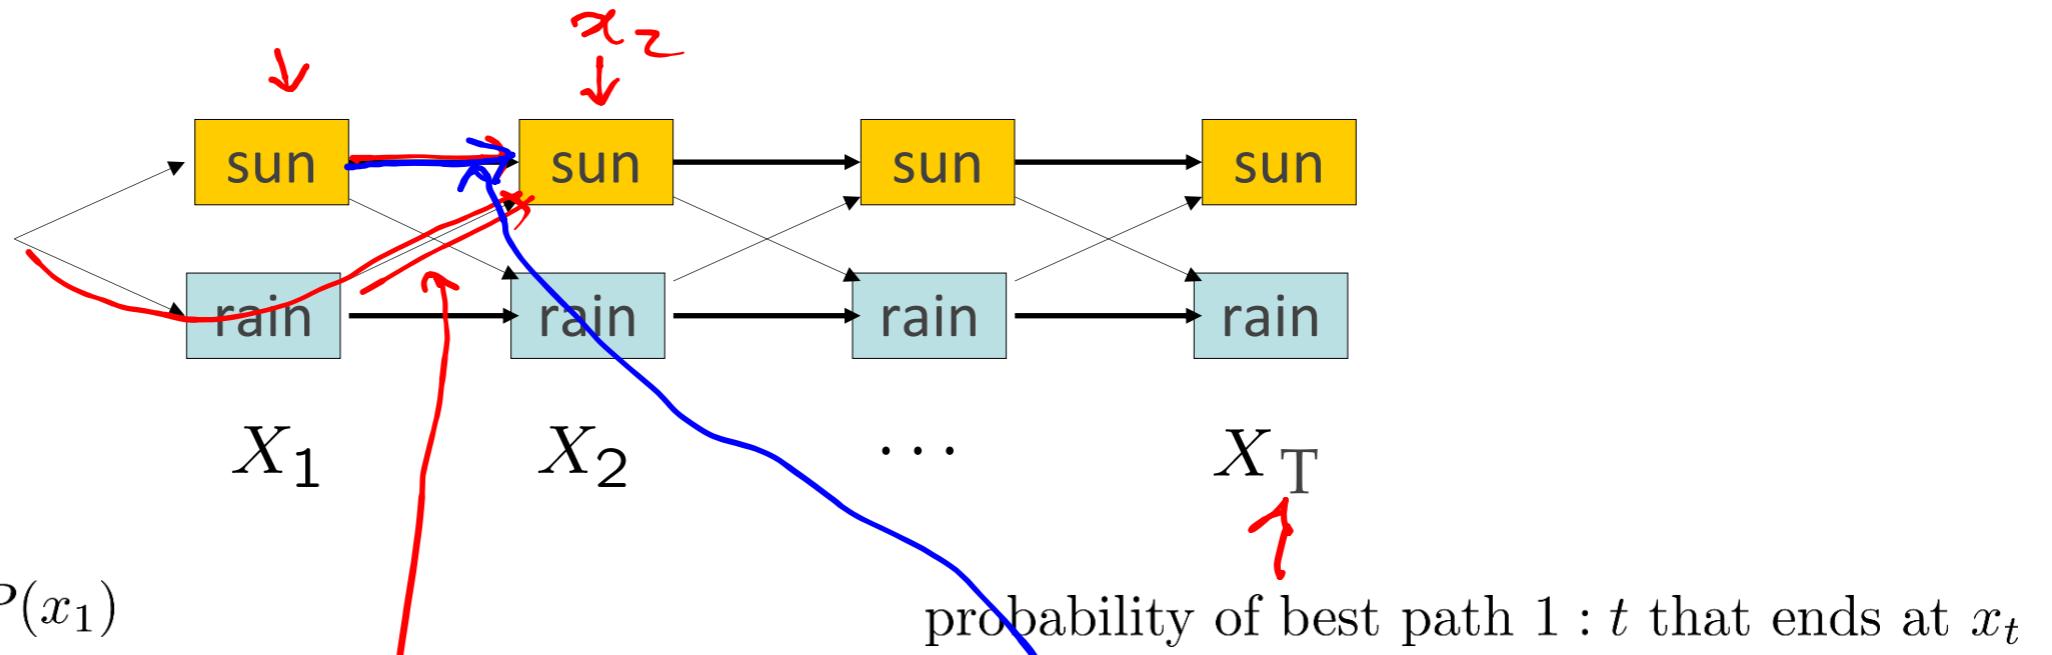
Viterbi Algorithm (max)

For each state at time t , keep track of the **maximum probability of any path** to it

$$\begin{aligned} \mathbf{m}_{1:t+1}^{\leftarrow} &= \text{VITERBI}(\mathbf{m}_{1:t}^{\leftarrow}, e_{t+1}) \\ &= P(e_{t+1}|X_{t+1}) \max_{x_t} P(X_{t+1}|x_t) \mathbf{m}_{1:t}^{\leftarrow} \end{aligned}$$

update move forward

Why is This True?



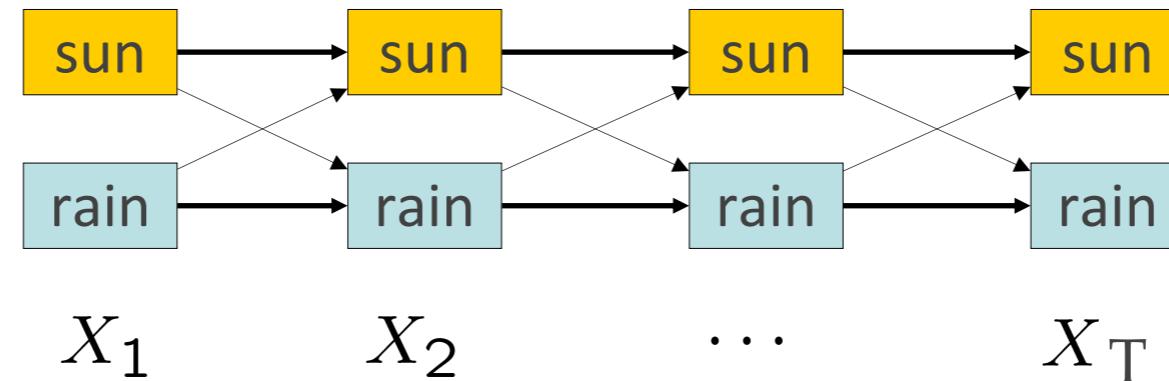
$$m_1[x_1] = P(e_1|x_1)P(x_1)$$

$$\begin{aligned} m_2[x_2] &= \max\{P(e_2|x_2)P(x_2|x_1 = \text{rain})m_1[x_1 = \text{rain}], P(e_2|x_2)P(x_2|x_1 = \text{sun})m_1[x_1 = \text{sun}]\} \\ &= \max_{x_1} P(e_2|x_2)P(x_2|x_1)m_1[x_1] \end{aligned}$$

$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

$$= P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1}) \underline{m_{t-1}[x_{t-1}]}$$

How About the MLE Path?



$m_N[x_N]$ ✓

probability of best path $1 : N$ that ends at x_N

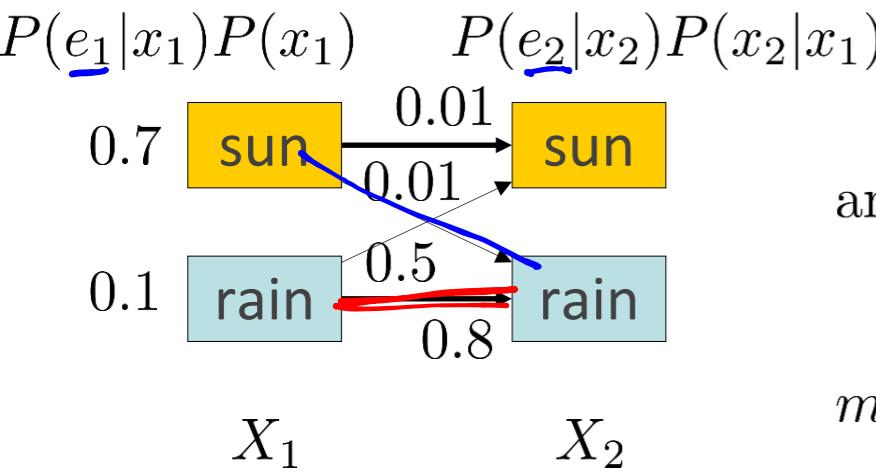
what is the last state on the most likely path?

$$\arg \max_{x_N} m_N[x_N]$$

what is the *second to last* state on the most likely path?

$$\arg \max_{x_{N-1}} m_{N-1}[x_{N-1}] ?$$

A Tricky Counter-Example



$$\arg \max_{x_1} m_1[x_1] = ?$$

sun! ✗

$$m_2[x_2] = \max_{x_1} P(e_2|x_2)P(x_2|x_1)m_1[x_1]$$

$$m_2[x_2 = \text{sun}] = \max\{0.7 \times 0.01, \underline{0.1 \times 0.5}\} = 0.05$$

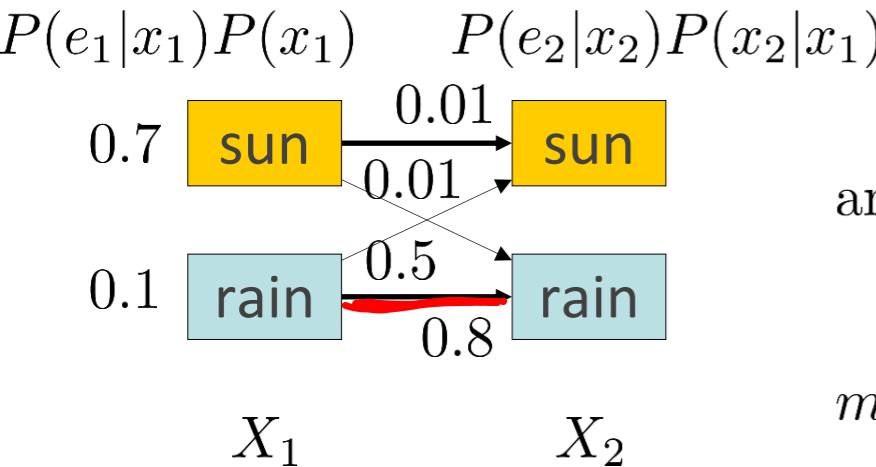
$$m_2[x_2 = \text{rain}] = \max\{0.7 \times 0.01, \underline{0.1 \times 0.8}\} = 0.08$$

arg $\max_{x_2} m_2[x_2] = \text{rain}$ ✓

$$P(\underline{x_1 = \text{sun}}, \underline{x_2 = \text{rain}}, e_1, e_2) = 0.7 \times 0.01 = 0.007$$

best path $1:t$ that ends at $x_t \neq$ best path $1:N$ that goes through x_t

How to Recover the MLE Path?



$$\arg \max_{x_1} m_1[x_1] = ?$$

sun!

$$m_2[x_2] = \max_{x_1} P(e_2|x_2)P(x_2|x_1)m_1[x_1]$$

this path starts at rain

$$m_2[x_2 = \text{sun}] = \max\{0.7 \times 0.01, 0.1 \times 0.5\} = 0.05$$

$$m_2[x_2 = \text{rain}] = \max\{0.7 \times 0.01, 0.1 \times 0.8\} = 0.08$$

$$\arg \max_{x_2} m_2[x_2] = \underline{\text{rain}}$$

rain

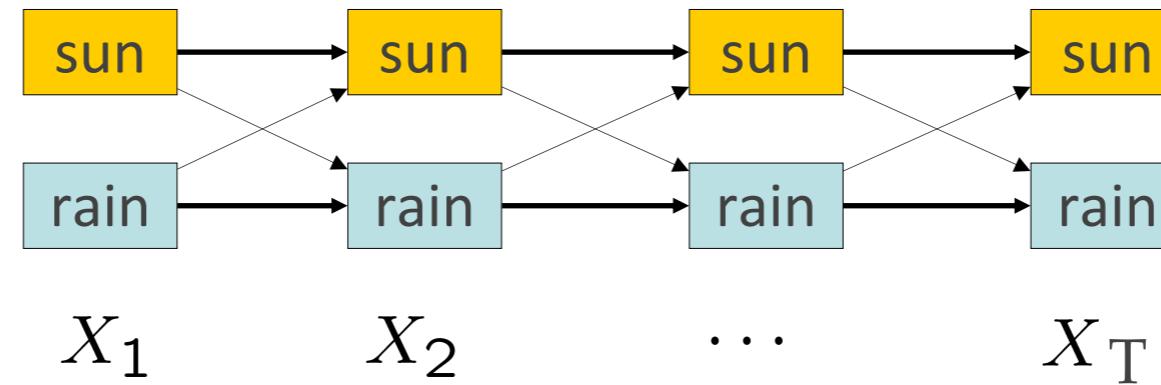
this path starts at rain

idea: what if we also save *where* the best path came from?

$$(m_t[x_t] = \max_{x_{t-1}} \underbrace{P(e_t|x_t)P(x_t|x_{t-1})}_{\text{path}} m_{t-1}[x_{t-1}])$$

$$(a_t[x_t] = \arg \max_{x_{t-1}} \underbrace{P(e_t|x_t)P(x_t|x_{t-1})}_{\text{path}} m_{t-1}[x_{t-1}])$$

Follow the Breadcrumbs...



for $t = 1$ to N :

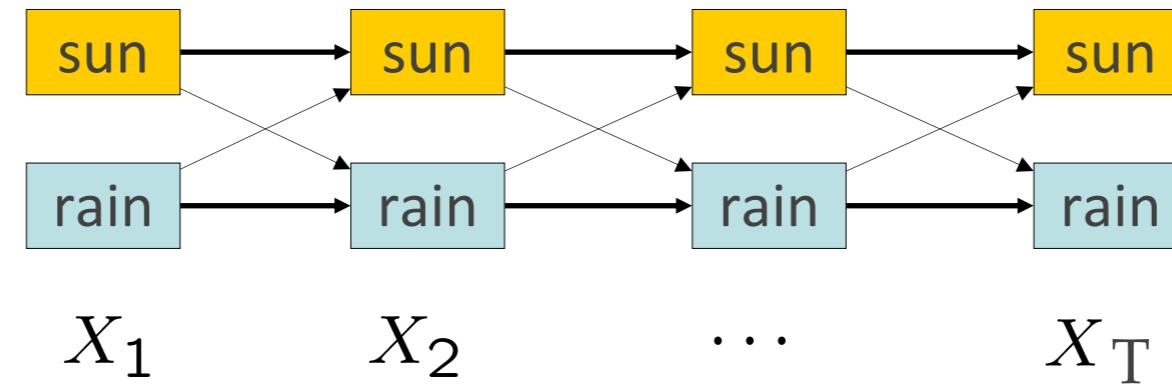
$$\begin{cases} m_t[x_t] = \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}] \\ a_t[x_t] = \arg \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}] \end{cases}$$

last state on most likely path: $x_N^* = \arg \max_{x_N} \underline{m_N}[x_N]$

second to last state on most likely path: $x_{N-1}^* = a_N[x_N^*]$ ↙

third to last state on most likely path: $x_{N-2}^* = a_{N-1}[x_{N-1}^*]$ ↙

Follow the Breadcrumbs...



for $t = 1$ to N :

$$m_t[x_t] = \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

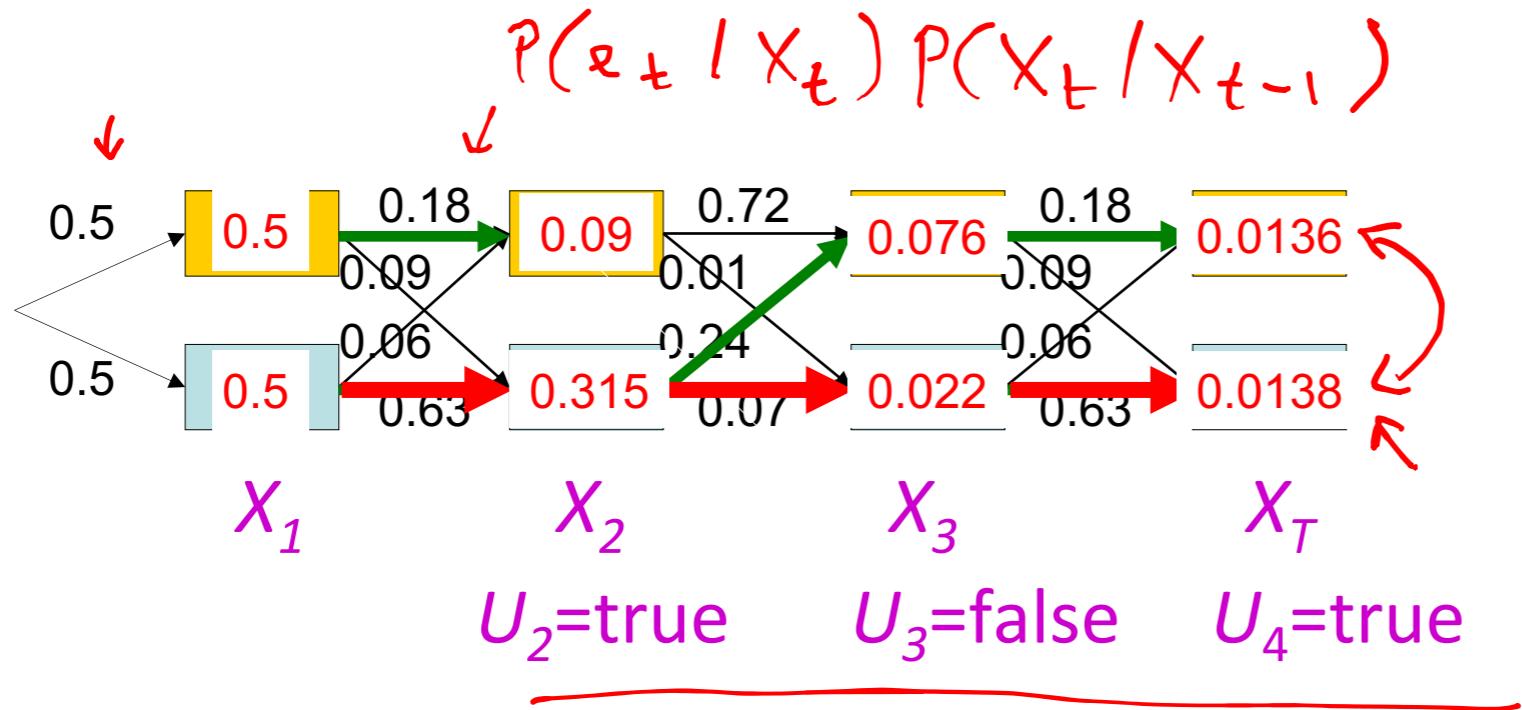
$$a_t[x_t] = \arg \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

$$x_N^* = \arg \max_{x_N} m_N[x_N] \quad \checkmark$$

for $t = N$ to 2:

$$x_{t-1}^* = a_t[x_t^*] \quad \checkmark$$

Viterbi Algorithm



W_{t-1}	$P(W_t W_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7

W_t	$P(U_t W_t)$	
	true	false
sun	0.2	0.8
rain	0.9	0.1

Time complexity?

$O(|X|^2 T)$

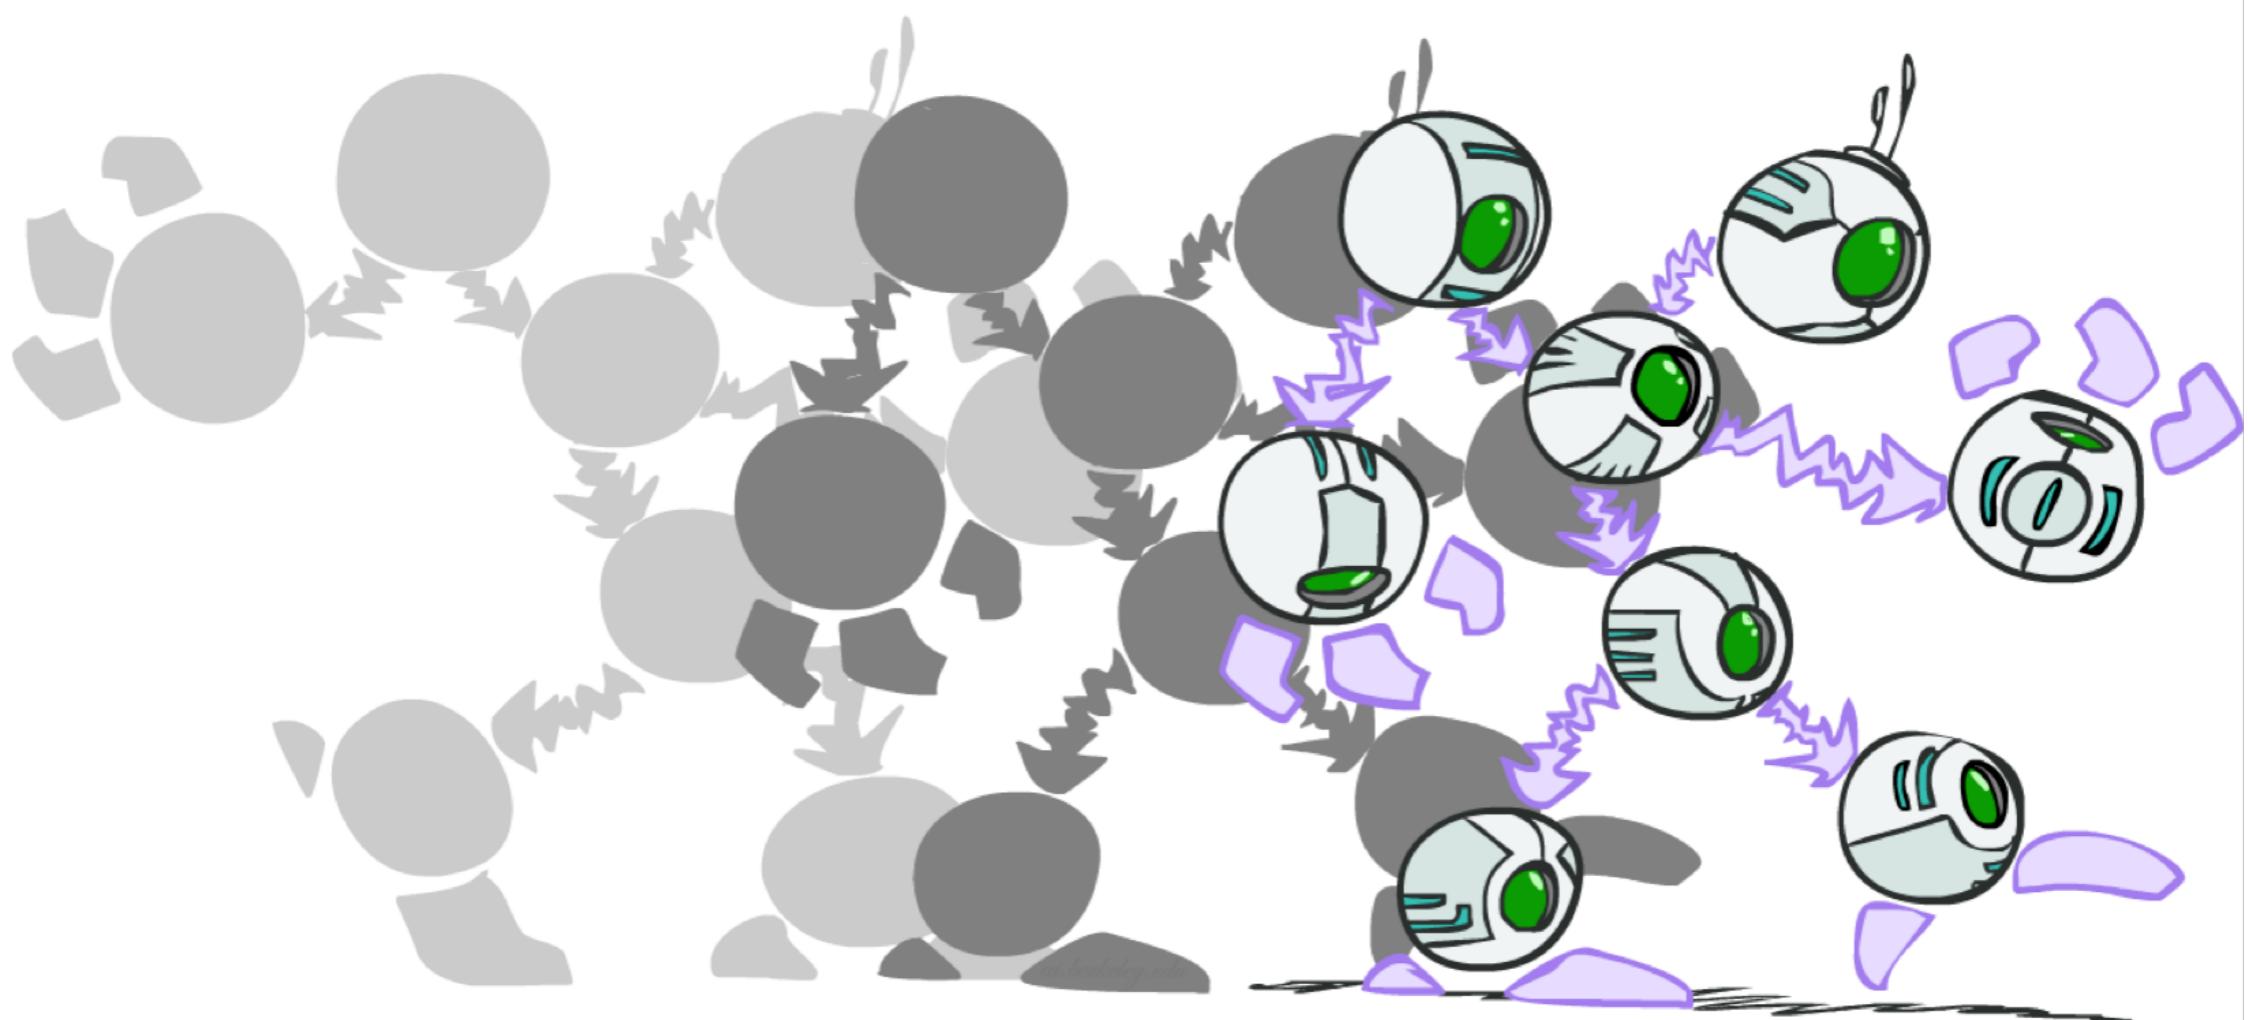
Space complexity?

$O(|X| T)$

Number of paths?

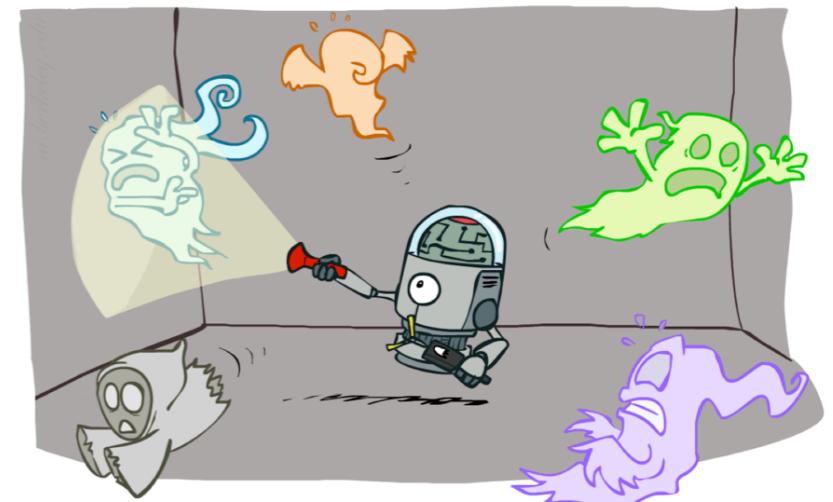
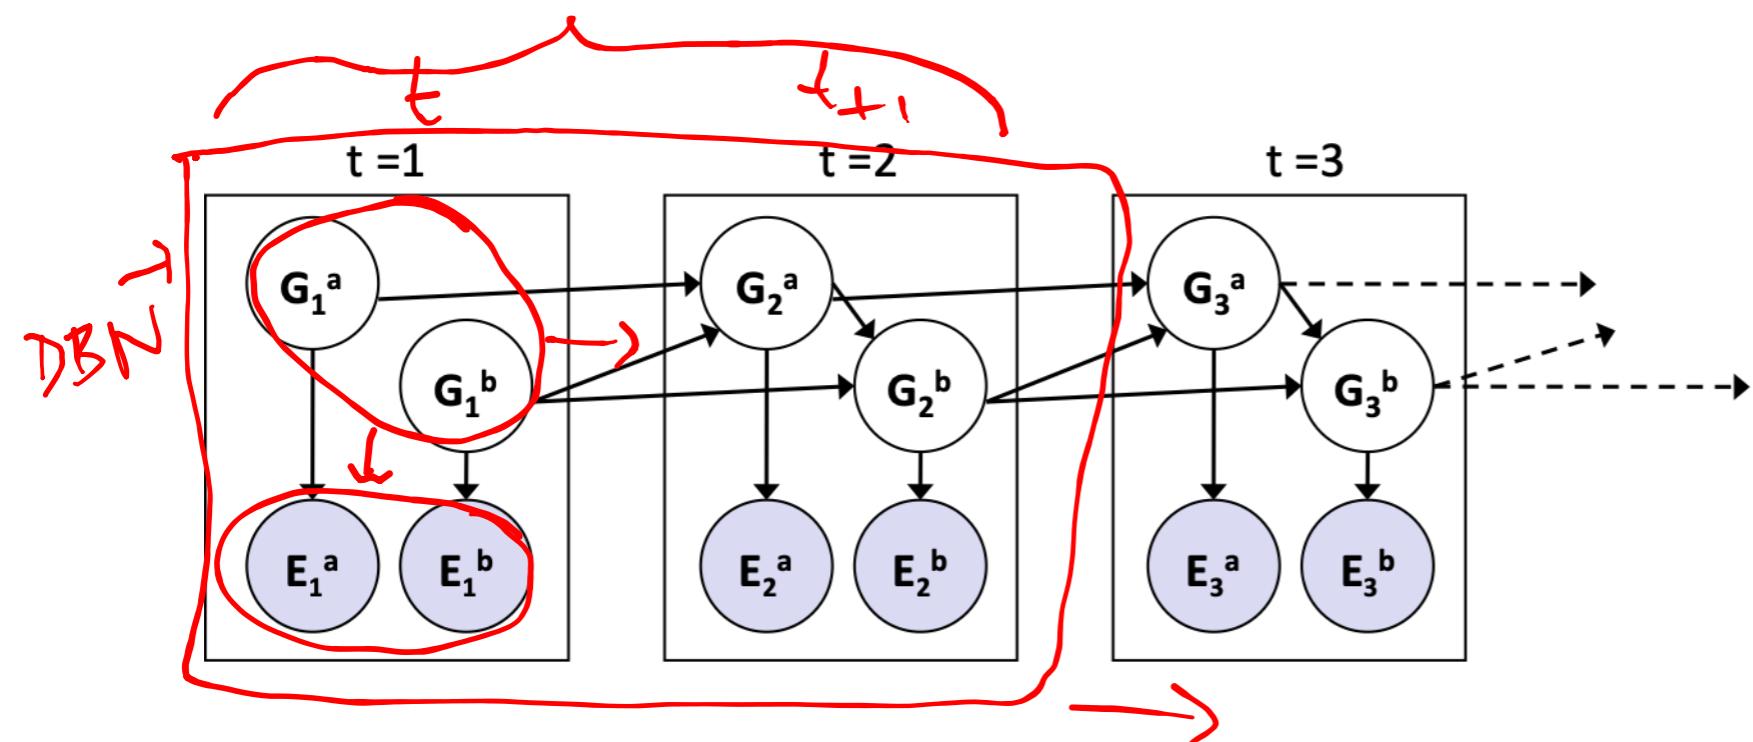
$O(|X|^T)$

Dynamic Bayes Nets



Dynamic Bayes Nets (DBNs)

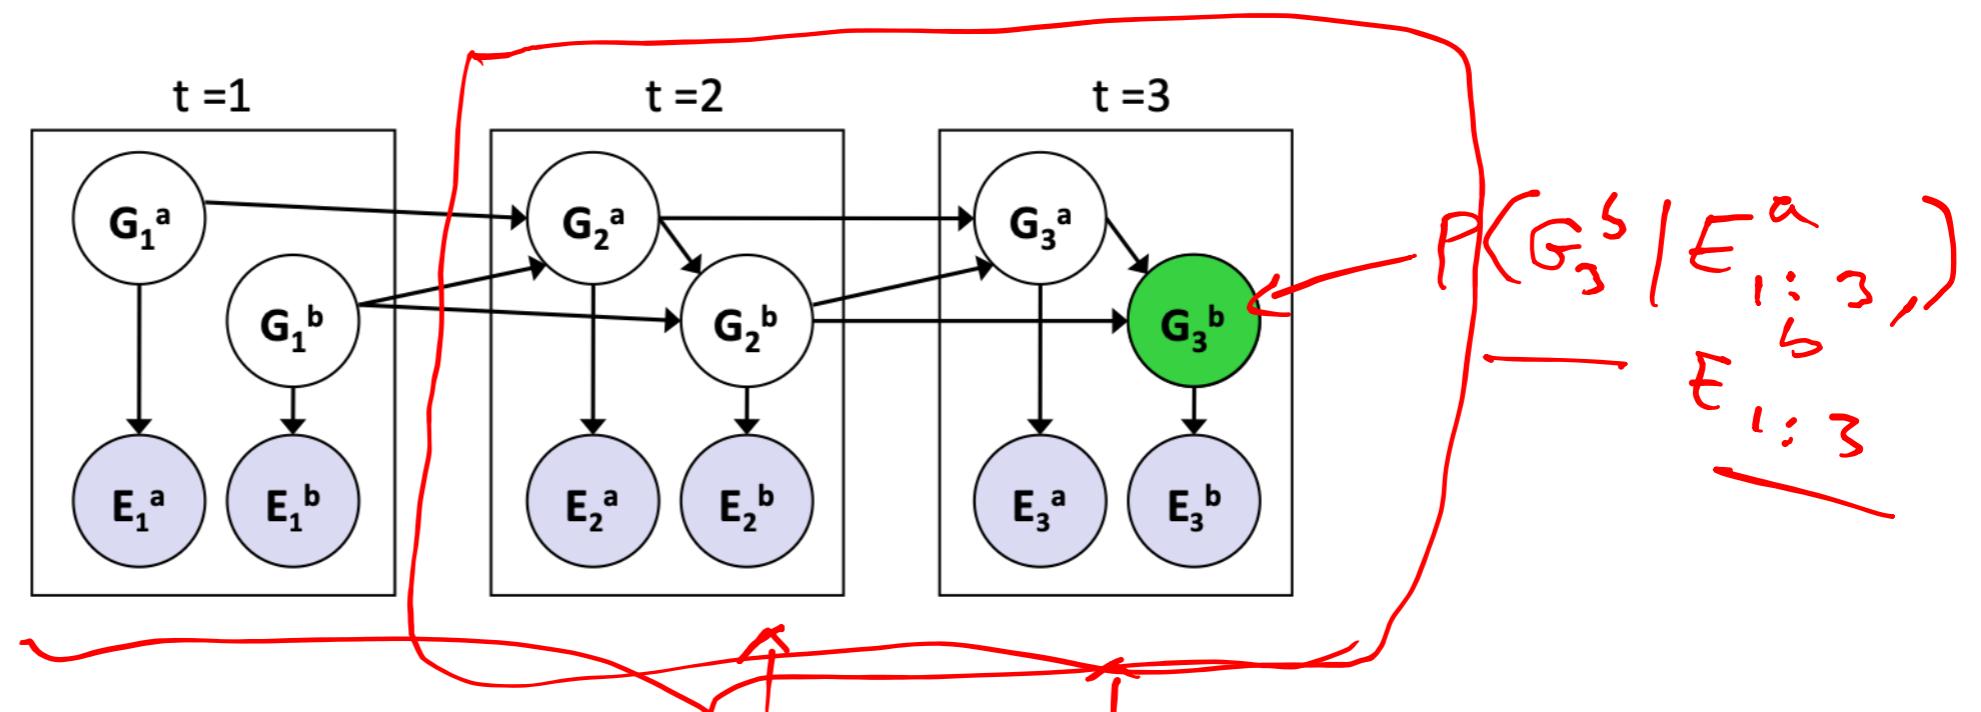
- ❖ We want to track multiple variables over time, using multiple sources of evidence
- ❖ Idea: Repeat a fixed Bayes net structure at each time
- ❖ Variables from time t can condition on those from $t-1$



- ❖ Dynamic Bayes nets are a generalization of HMMs

Exact Inference in DBNs

- ❖ Variable elimination applies to dynamic Bayes nets
- ❖ Procedure: “unroll” the network for T time steps, then eliminate variables until $P(X_T | e_{1:T})$ is computed



- ❖ Online belief updates: Eliminate all variables from the previous time step; store factors for current time only

$$P_L(G_2^a, G_2^b | E_{1:2}^{a,b})$$

DBN Particle Filters

- ❖ A particle is a complete sample for a time step
- ❖ **Initialize:** Generate prior samples for the t=1 Bayes net
 - ❖ Example particle: $\underline{\mathbf{G}_1^a} = (3,3) \underline{\mathbf{G}_1^b} = (5,3)$
- ❖ **Elapse time:** Sample a successor for each particle
 - ❖ Example successor: $\mathbf{G}_2^a = (2,3) \mathbf{G}_2^b = (6,3)$
- ❖ **Observe:** Weight each entire sample by the likelihood of the evidence conditioned on the sample
 - ❖ Likelihood: $P(\underline{\mathbf{E}_1^a} | \mathbf{G}_1^a) * P(\mathbf{E}_1^b | \mathbf{G}_1^b)$
- ❖ **Resample:** Select prior samples (tuples of values) in proportion to their likelihood